

TCP/IP Network, Transport and Application Layers



Cabrillo College

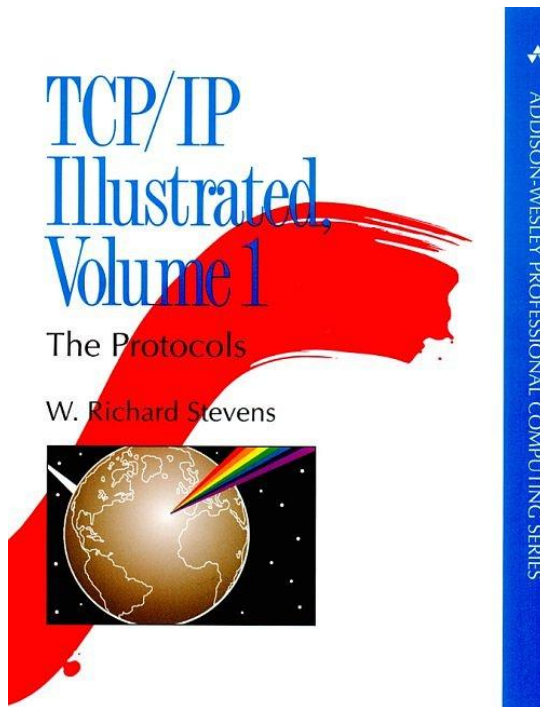
CIS 81 and CST 311

Rick Graziani

Spring 2006

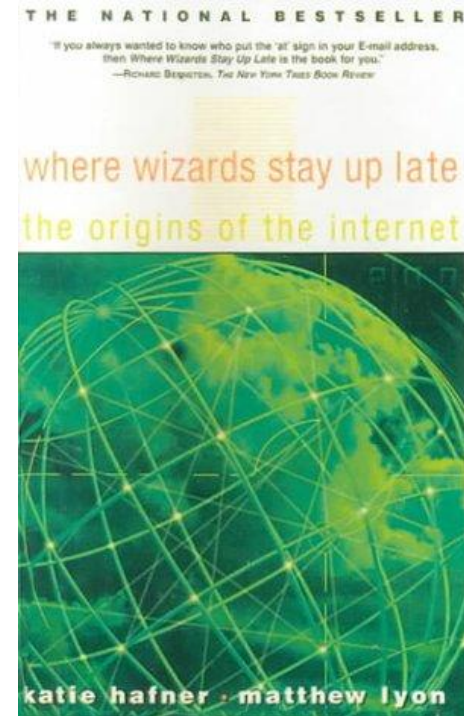
- It is important for networking professionals to have a very good understanding of TCP/IP.
- Various devices communicate using the multiple protocols of the TCP/IP protocol suite.
- A networking professional needs to know how these protocols function and interact with each other in order to properly understand, analyze and troubleshoot networking issues.
- This chapter is only an introduction to this information.
- I strongly suggest taking a separate course in the TCP/IP protocol suite, in addition to system administration courses such as those for Microsoft Windows (MCSE/MCSA) or Unix/Linux.
- The majority of this presentation is taken directly from the on-line curriculum (present and past) – however there are a few mistakes or misconceptions in the on-line curriculum which is addressed in this presentation.
- Many of the concepts in the presentation are missing some important details to keep the amount of information to a reasonable limit – Again I suggest taking a course on TCP/IP protocol suite.
- Also, two other presentations are included on my web site:
 - ARP
 - ICMP – Understanding ping and trace

Important and Interesting Reading



TCP/IP Illustrated, Vol. 1
W. Richard Stevens
Addison-Wesley Pub Co
ISBN: 0201633469

- Although, published in 1994, written by the late Richard Stevens, it is still regarded as the definitive book on TCP/IP.



Where Wizards Stay Up Late
Katie Hafner and Matthew Lyon
ISBN 0613181530

- Very enjoyable reading and you do not have to be a networking geek to enjoy it!
- National Bestseller

Topics

Layer 3, Network Layer Concepts

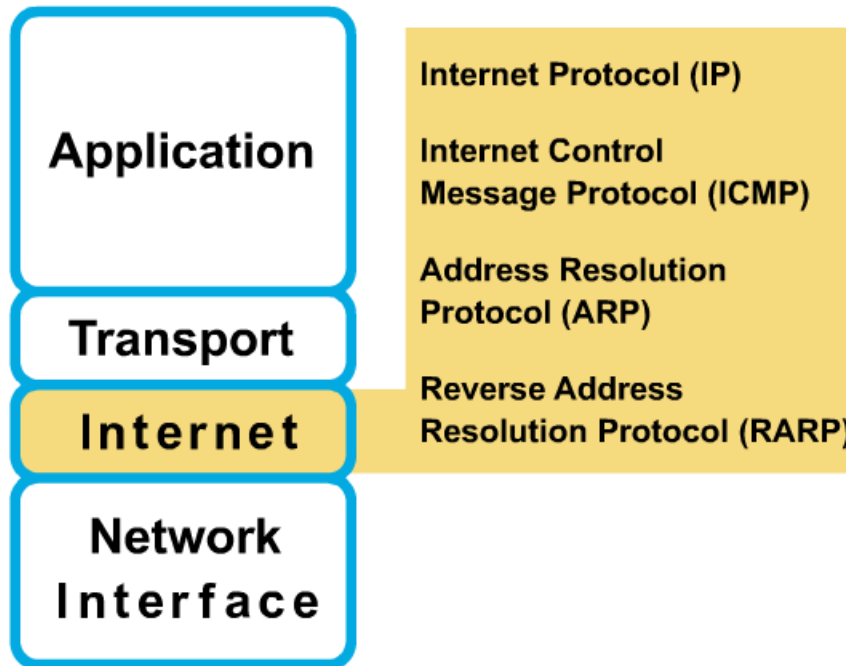
- TCP/IP and the Internet Layer
- Diagram the IP datagram
- Internet Control Message Protocol (ICMP)

TCP/IP protocol stack and the transport layer

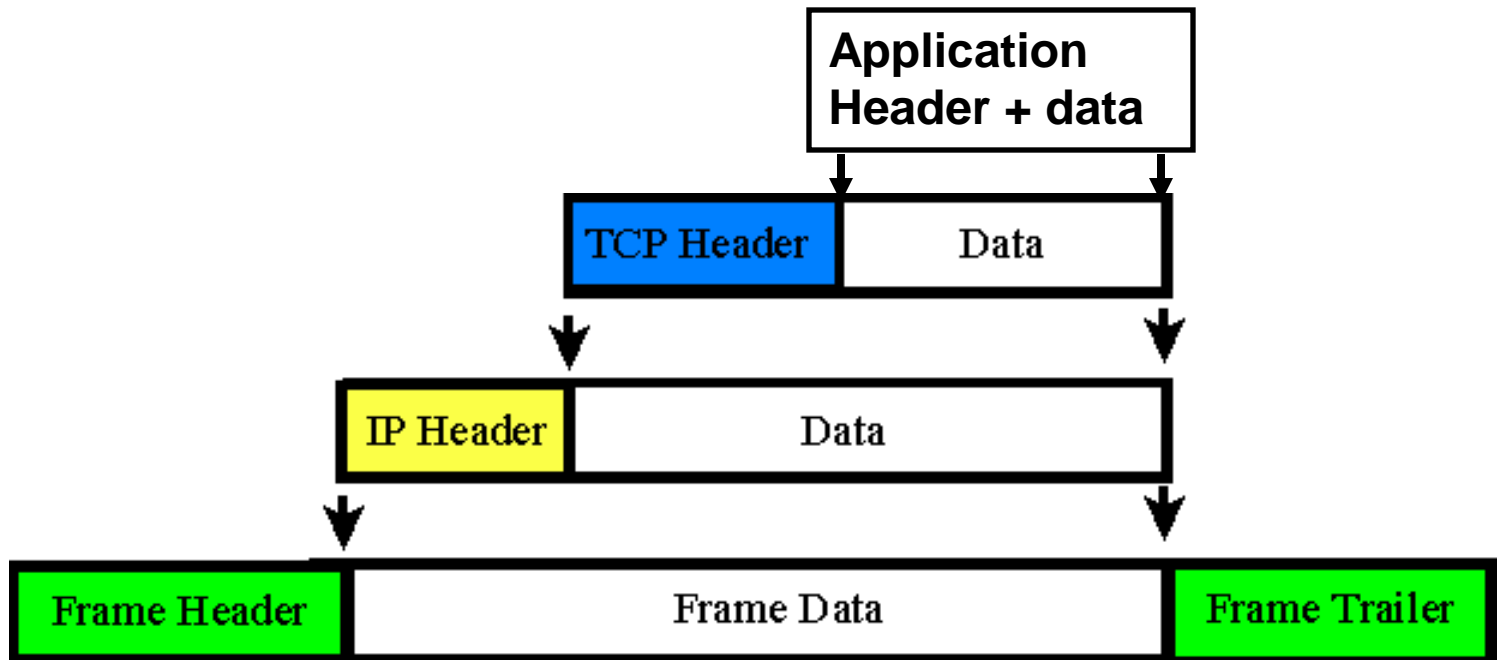
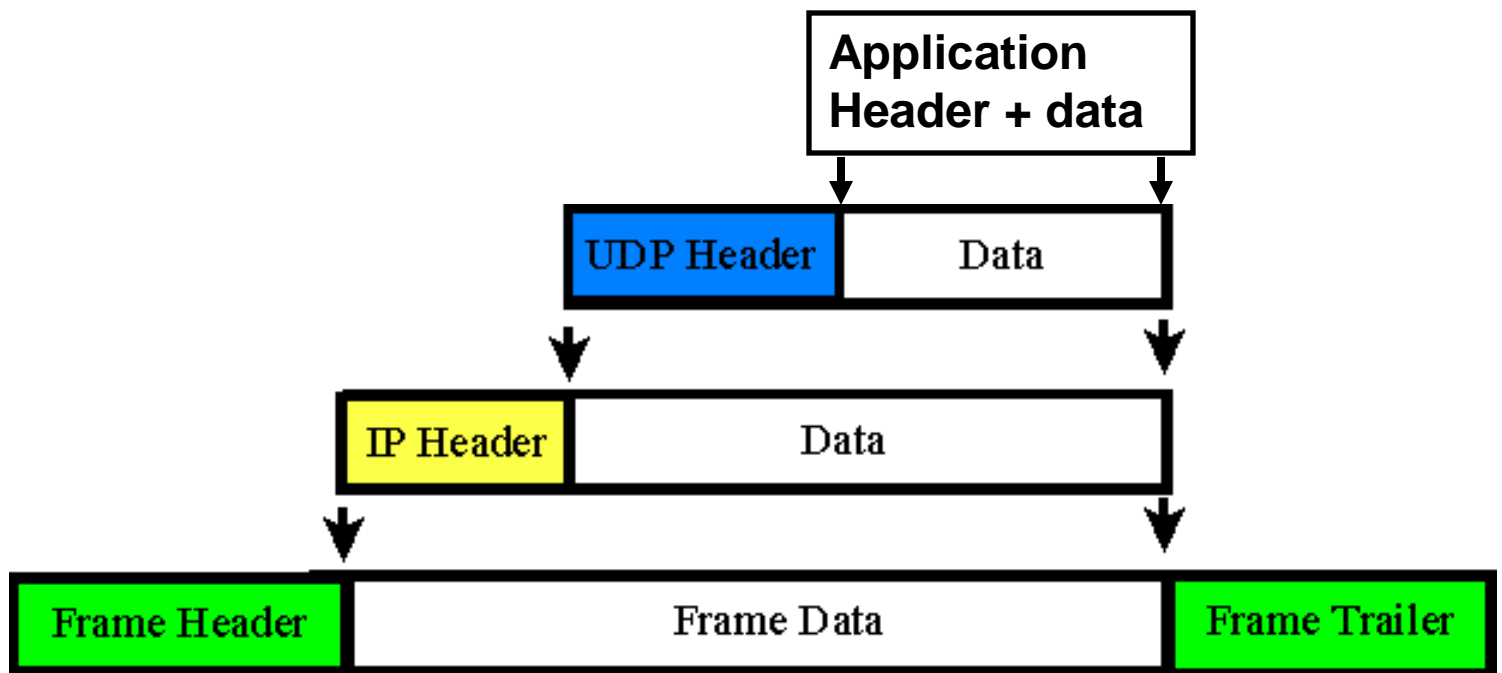
- TCP and UDP segment format
- TCP and UDP port numbers
- TCP three-way handshake/open connection
- TCP simple acknowledgment and windowing

Layer 3: TCP/IP Network Layer

Network Layer Overview

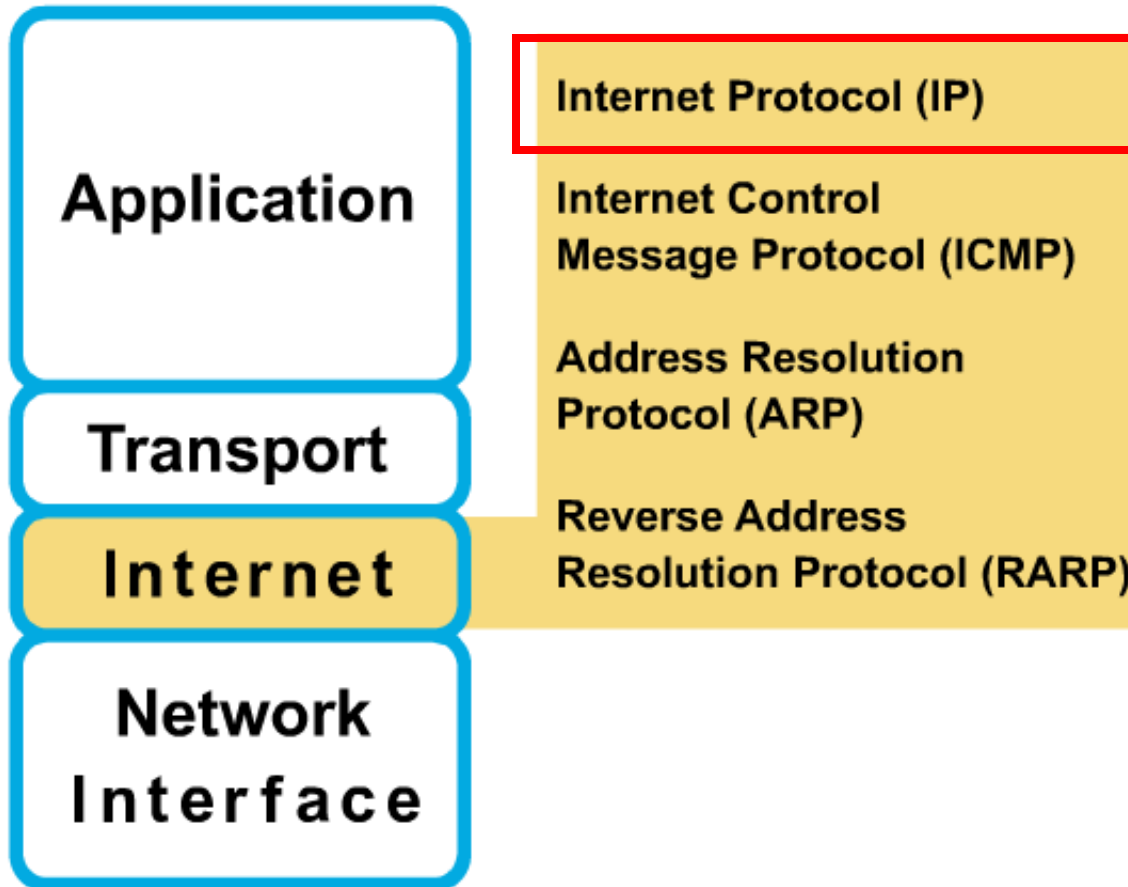


- The Internet layer of the TCP/IP stack corresponds to the network layer of the OSI model.
- Each layer is responsible for getting packets through a network using software addressing.



IP – Internet Protocol

Network Layer Overview



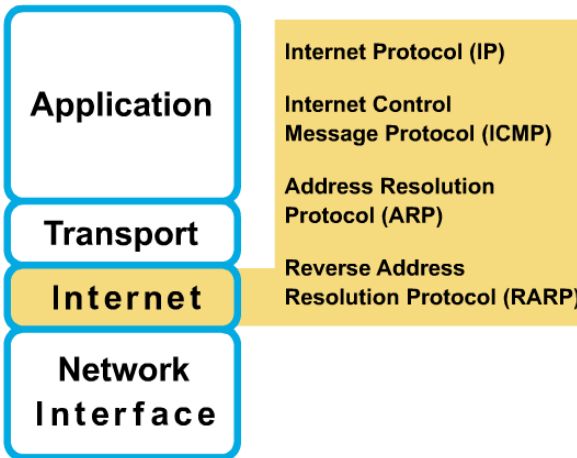
IP Packet (Data Gram) Header

0	15	16	31	
4-bit Version	4-bit Header Length	8-bit Type Of Service (TOS)	16-bit Total Length (in bytes)	
16-bit Identification			3-bit Flags	13-bit Fragment Offset
8 bit Time To Live TTL		8-bit Protocol	16-bit Header Checksum	
32-bit Source IP Address				
32-bit Destination IP Address				

Options (if any)

Data

Network Layer Overview

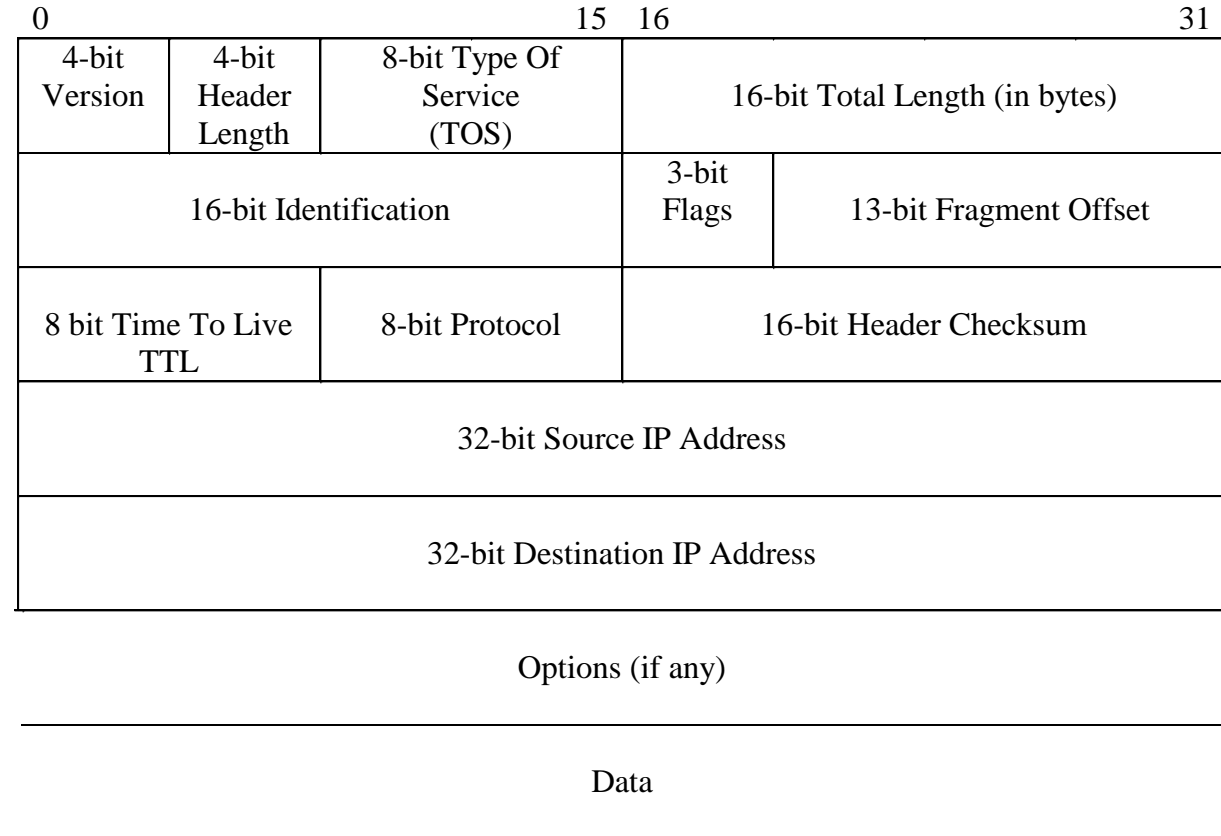
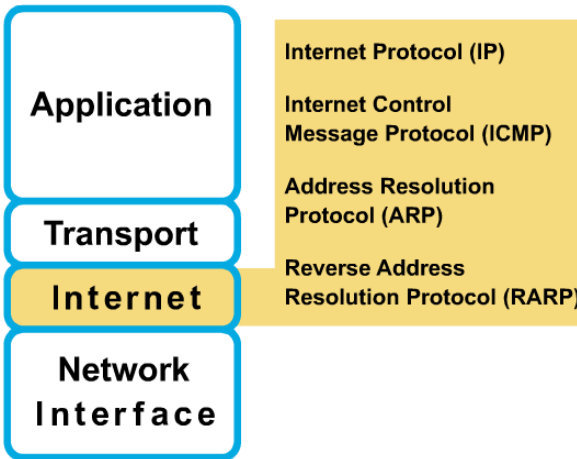


0	15	16	31	
4-bit Version	4-bit Header Length	8-bit Type Of Service (TOS)	16-bit Total Length (in bytes)	
16-bit Identification			3-bit Flags	13-bit Fragment Offset
8 bit Time To Live TTL		8-bit Protocol	16-bit Header Checksum	
32-bit Source IP Address				
32-bit Destination IP Address				
Options (if any)				
Data				

IP Packet (Data Gram) Header

- **VERS** -- version number
- **HLEN** -- header length, in 32-bit words
- **type of service** -- how the datagram should be handled
- **total length** -- total length (header + data)
- **identification, flags, flag offset** -- provides fragmentation of datagrams to allow differing MTUs in the internetwork

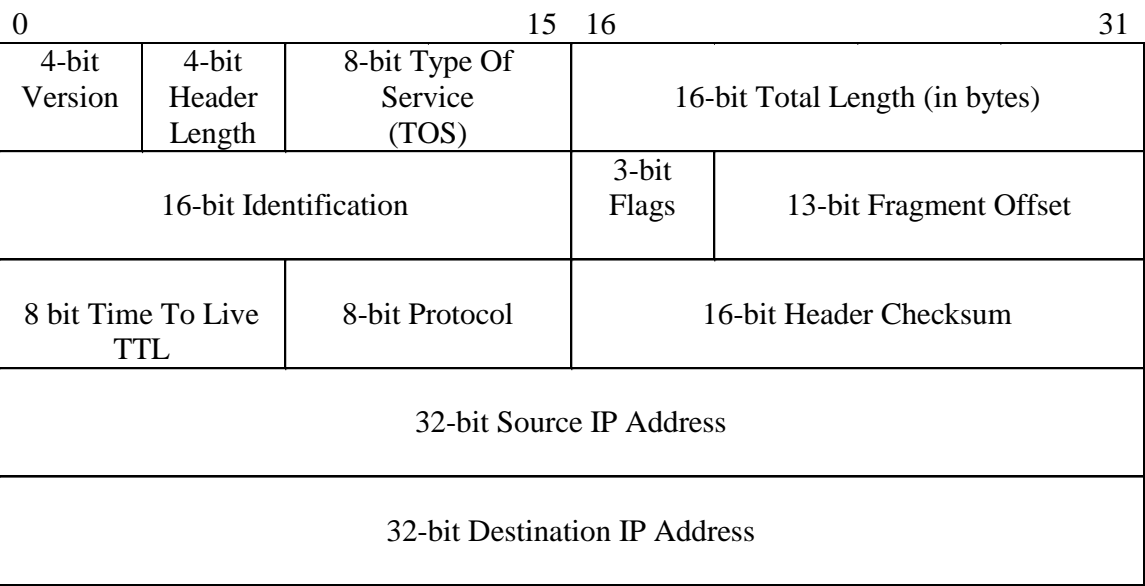
Network Layer Overview



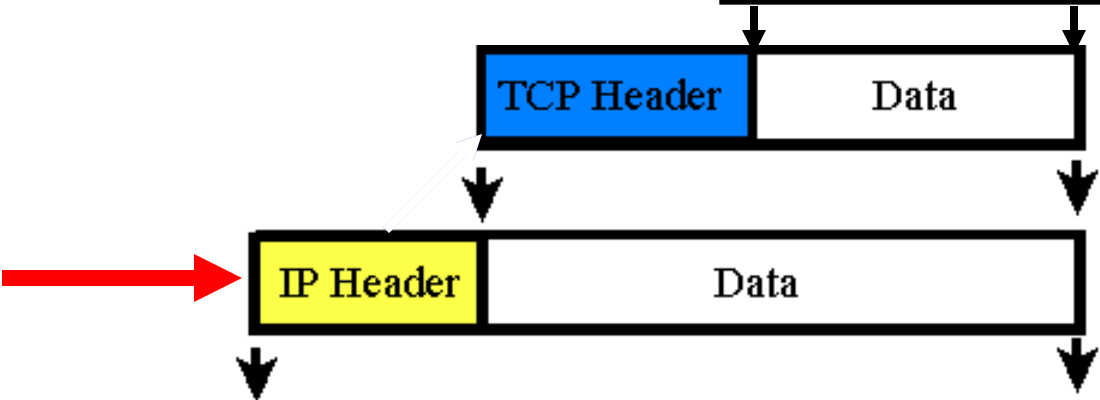
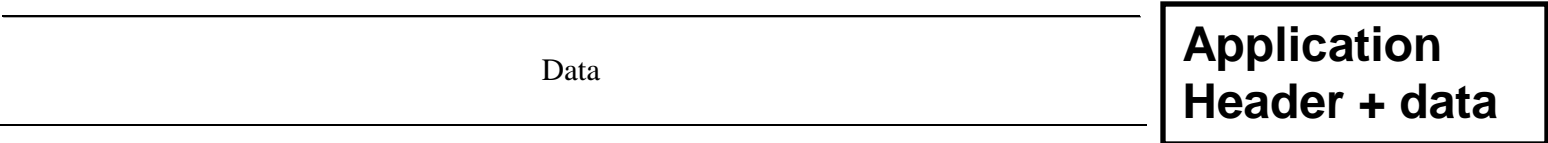
IP Packet (Data Gram) Header

- **TTL** -- Time-To-Live
- **protocol** -- the upper-layer (Layer 4) protocol sending the datagram
- **header checksum** -- an integrity check on the header
- **source IP address and destination IP address** -- 32-bit IP addresses
- **IP options** -- network testing, debugging, security, and other options
- **Data** – Upper layer headers and data

IP Header



Options (if any)



IP's TTL – Time To Live field

IP Header

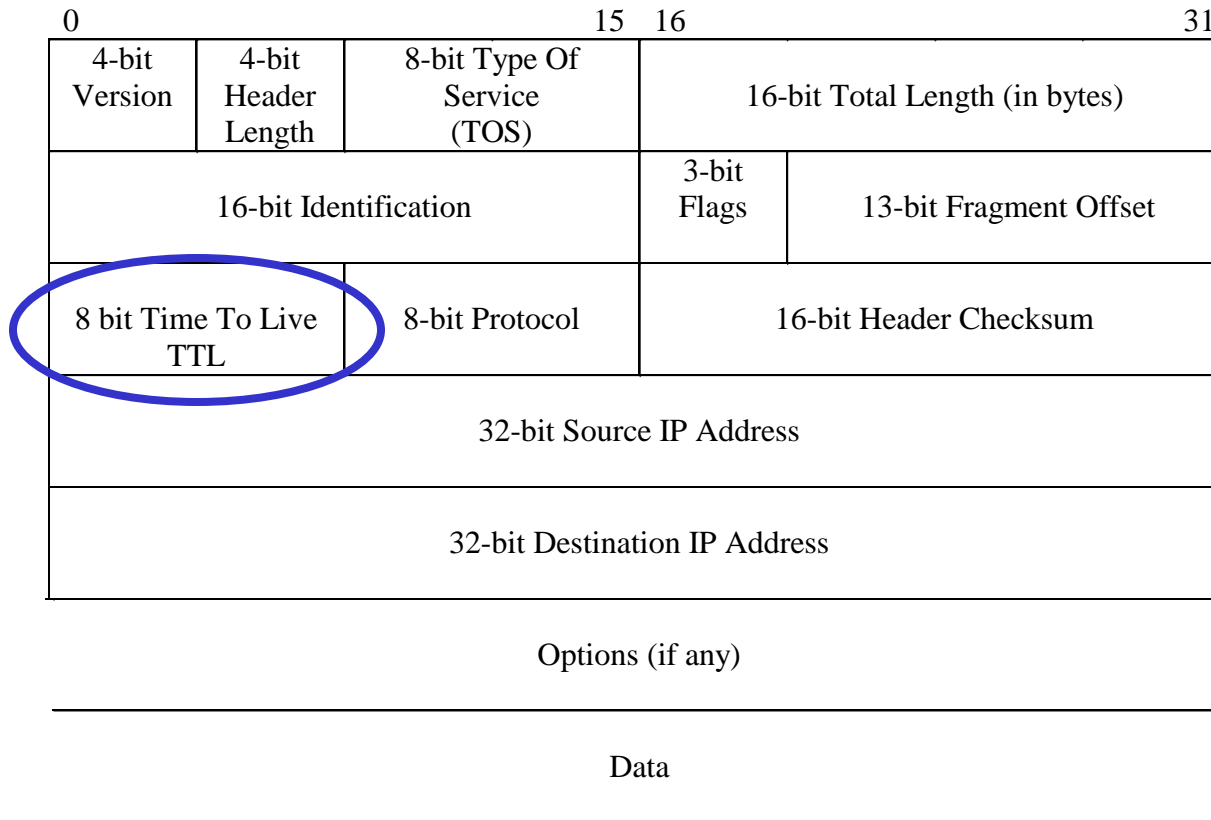
0		15		16		31	
4-bit Version	4-bit Header Length	8-bit Type Of Service (TOS)		16-bit Total Length (in bytes)			
16-bit Identification				3-bit Flags	13-bit Fragment Offset		
8 bit Time To Live TTL		8-bit Protocol		16-bit Header Checksum			
32-bit Source IP Address							
32-bit Destination IP Address							

Options (if any)

Data

IP's TTL – Time To Live field

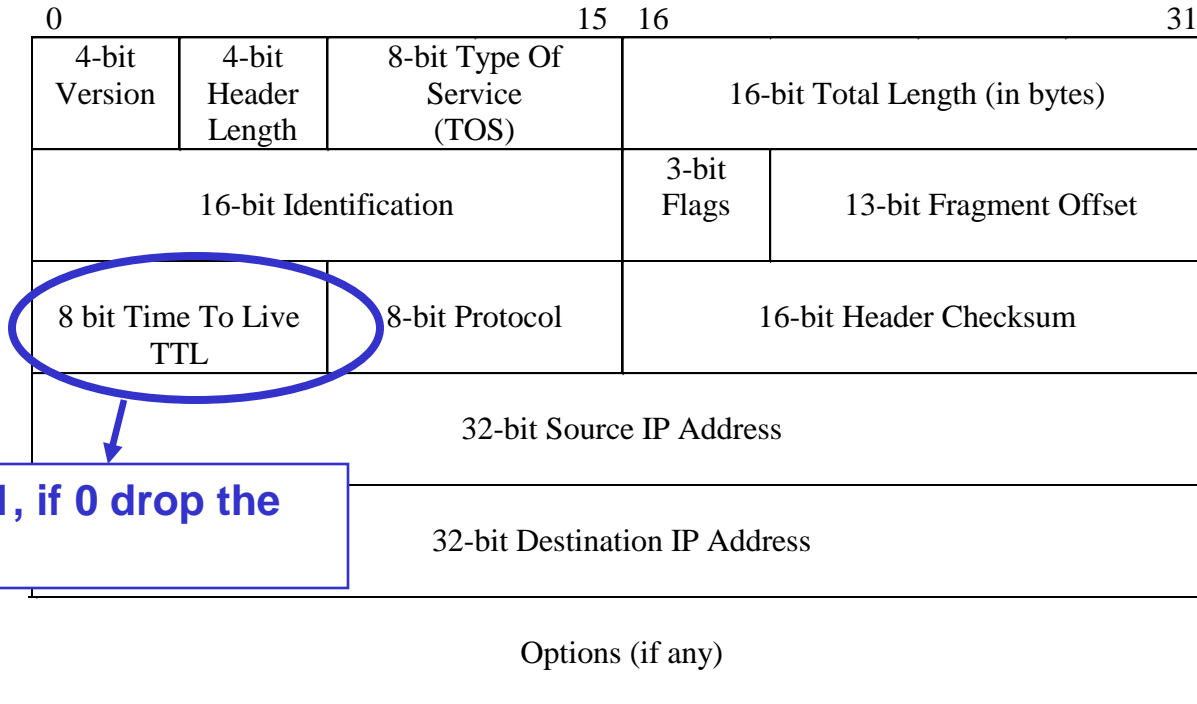
IP Header



- When a packet is first generated a value is entered into the TTL field.
- Originally, the TTL field was the number of seconds, but this was difficult to implement and rarely supported.
- Now, the TTL is now set to a specific value which is then decremented by each router.

IP's TTL – Time To Live field

IP Header



Data

- If the router decrements the TTL field to 0, it will then drop the packet (unless the packet is destined specifically for the router, i.e. ping, telnet, etc.).
- Common operating system TTL values are:
 - UNIX: **255**
 - Linux: **64 or 255** depending upon vendor and version
 - Microsoft Windows 95: **32**
 - Other Microsoft Windows operating systems: **128**

http://www.switch.ch/docs/ttl_default.html

TTL Overview - Disclaimer:

The following list is a best effort overview of some widely used TCP/IP stacks. The information was provided by vendors and many helpful system administrators. We would like to thank all these contributors for their precious help ! SWITCH cannot, however, take any responsibility that the provided information is correct. Furthermore, SWITCH cannot be made liable for any damage that may arise by the use of this information.

OS Version	"safe"	tcp_ttl	udp_ttl
AIX	n	60	30
DEC Pathworks V5	n	30	30
FreeBSD 2.1R	y	64	64
HP/UX 9.0x	n	30	30
HP/UX 10.01	y	64	64
Irix 5.3	y	60	60
Irix 6.x	y	60	60
Linux	y	64	64
MacOS/MacTCP 2.0.x	y	60	60
OS/2 TCP/IP 3.0	y	64	64
OSF/1 V3.2A	n	60	30
Solaris 2.x	y	255	255
SunOS 4.1.3/4.1.4	y	60	60
Ultrix V4.1/V4.2A	n	60	30
VMS/Multinet	y	64	64
VMS/TCPware	y	60	64
VMS/Wollongong 1.1.1.1	n	128	30
VMS/UCX (latest rel.)	y	128	128
MS WfW	n	32	32
MS Windows 95	n	32	32
MS Windows NT 3.51	n	32	32
MS Windows NT 4.0	y	128	128

Assigned Numbers (RFC
1700, J. Reynolds, J.
Postel, October 1994):

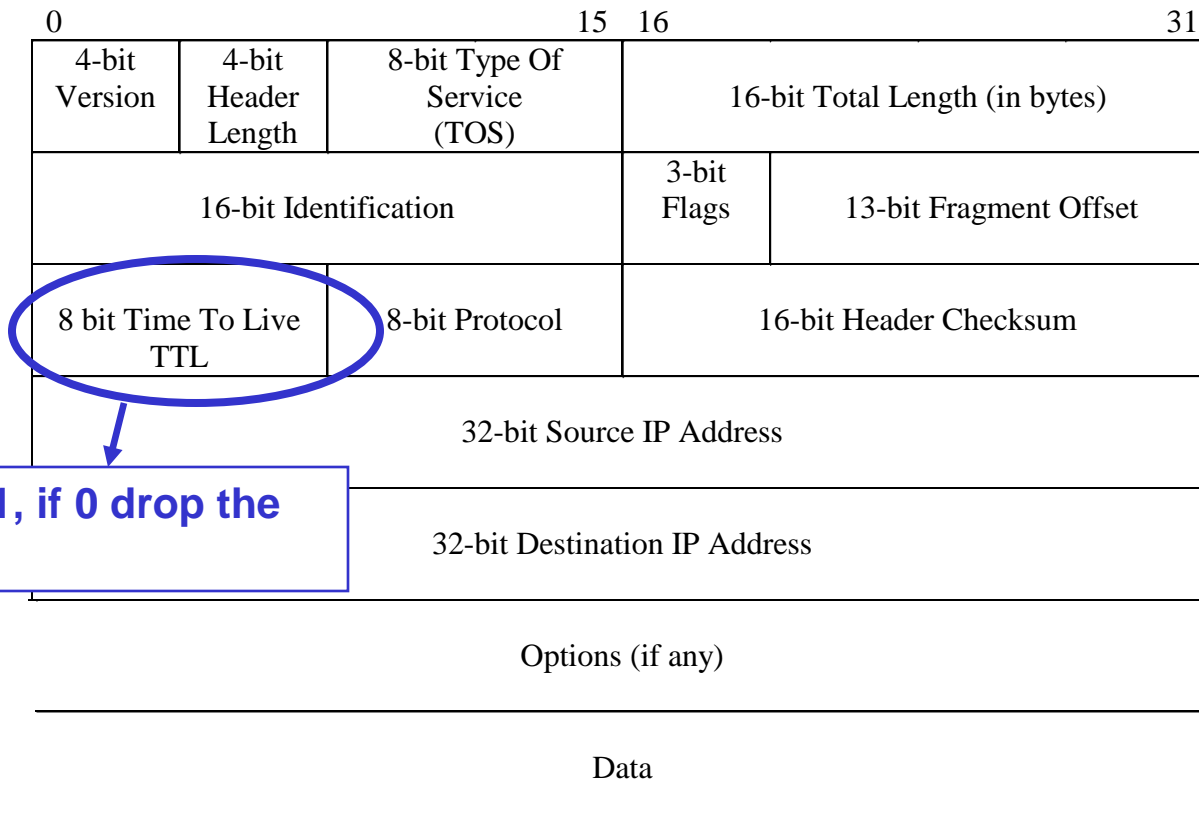
IP TIME TO LIVE PARAMETER

The current
recommended default
time to live (TTL)
for the Internet
Protocol (IP) is 64.

Safe: TCP and UDP
initial TTL values
should be set to a
"safe" value of at
least 60 today.

IP's TTL – Time To Live field

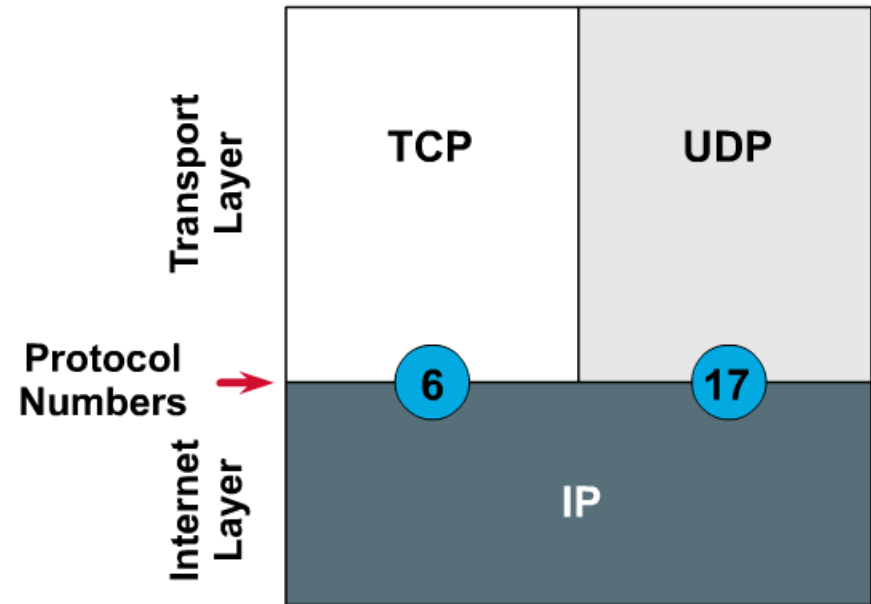
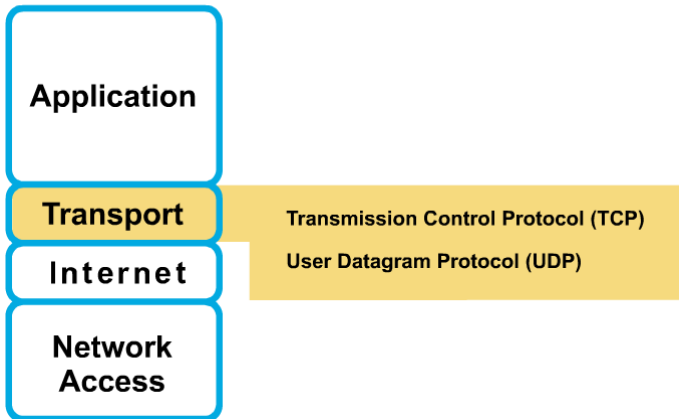
IP Header



- The idea behind the TTL field is that IP packets can not travel around the Internet forever, from router to router.
- Eventually, the packet's TTL which reach 0 and be dropped by the router, even if there is a routing loop somewhere in the network.

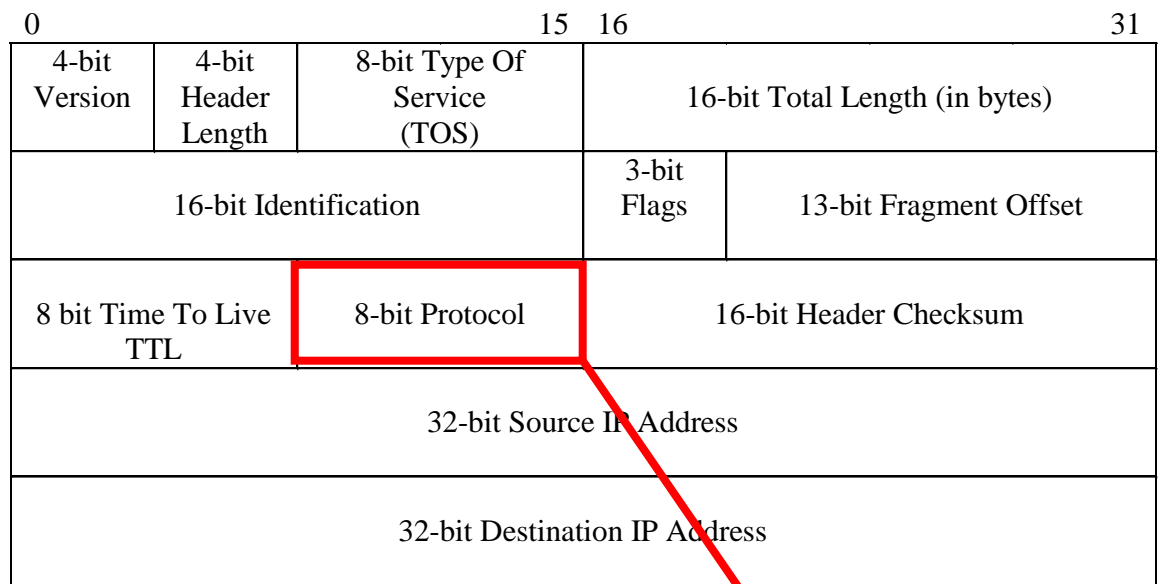
The Protocol Field

Transport Layer Overview



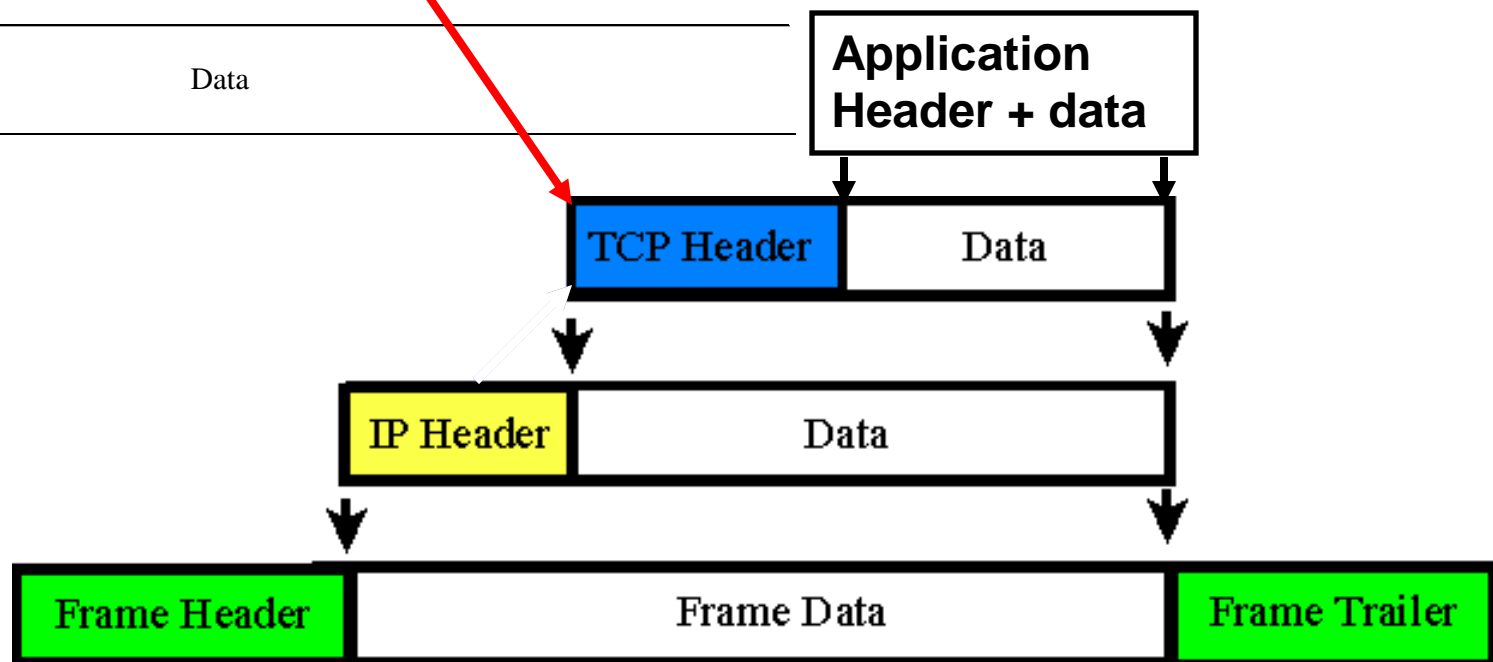
- The protocol field determines the Layer 4 protocol being carried within an IP datagram.
- Although much of the IP traffic uses TCP, other protocols can also use UDP, other transport layers, or UDP.
- Each IP header must identify the destination Layer 4 protocol for the datagram.
- Transport layer protocols are numbered, similarly to **port numbers**.
- IP includes the protocol number in the protocol field.

IP Header



Options (if any)

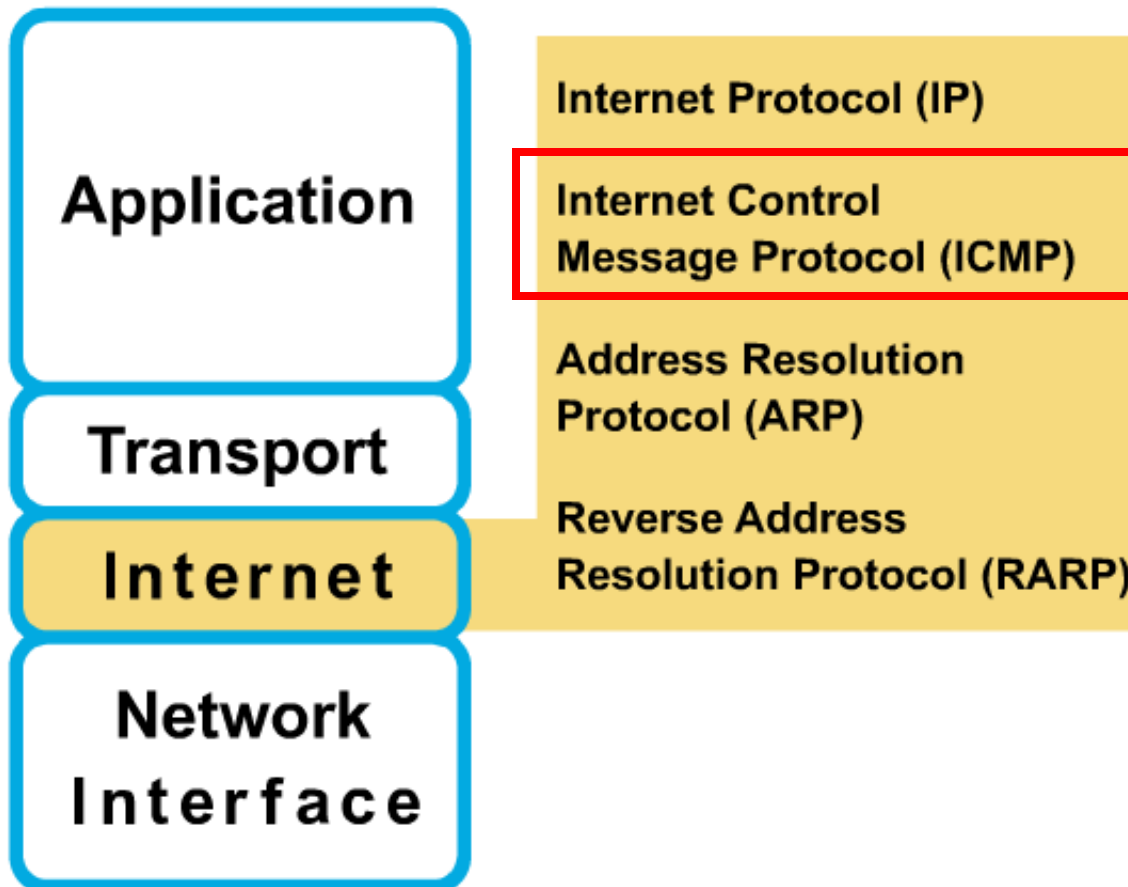
Data



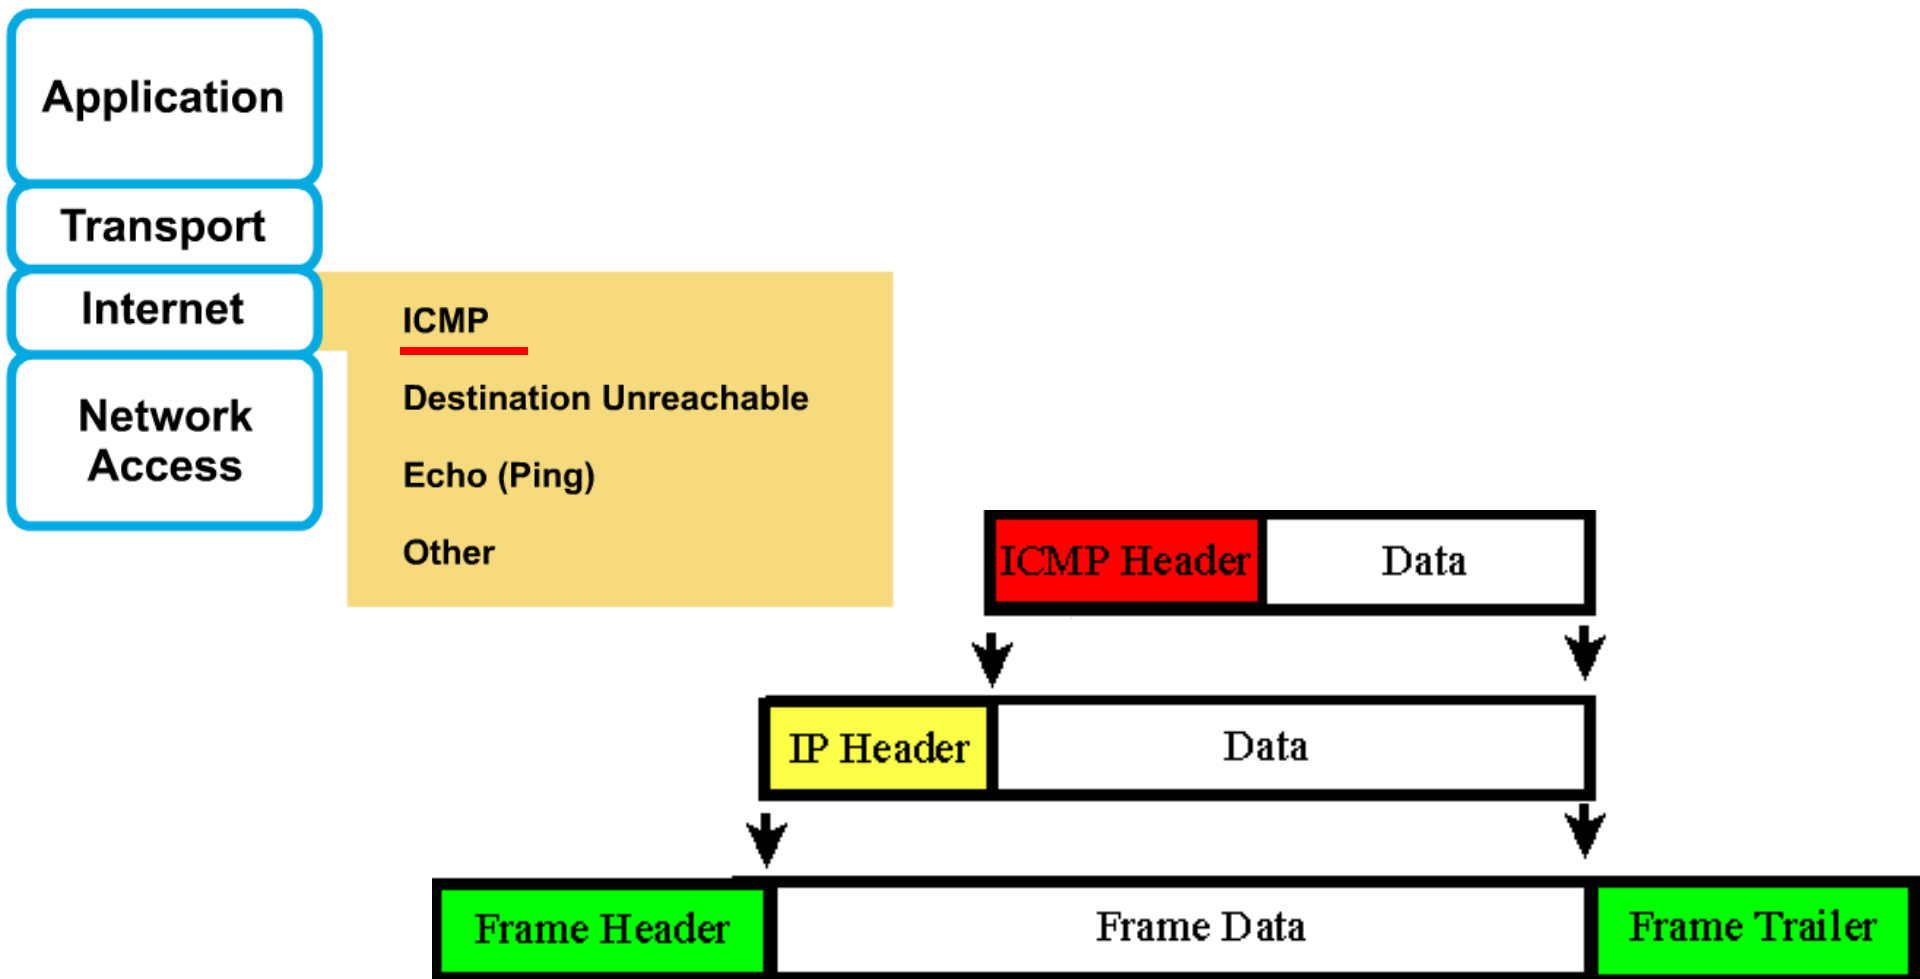
ICMP – Internet Control Message Protocol

Cabrillo College

Network Layer Overview

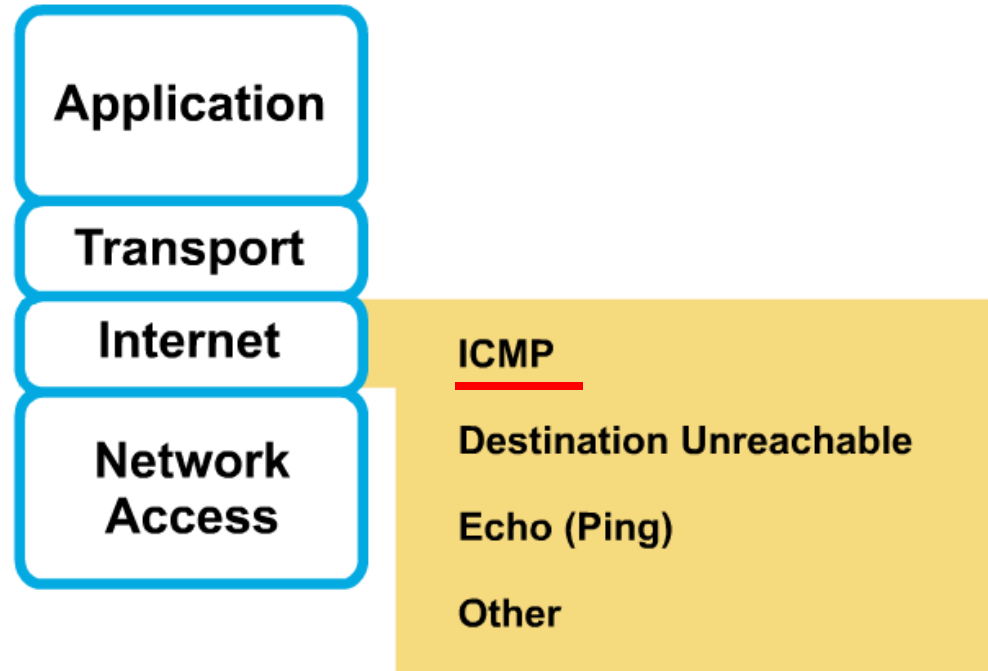


Internet Control Message Protocol



- All TCP/IP hosts implement ICMP. ICMP messages are carried in IP datagrams and are used to send error and control messages.

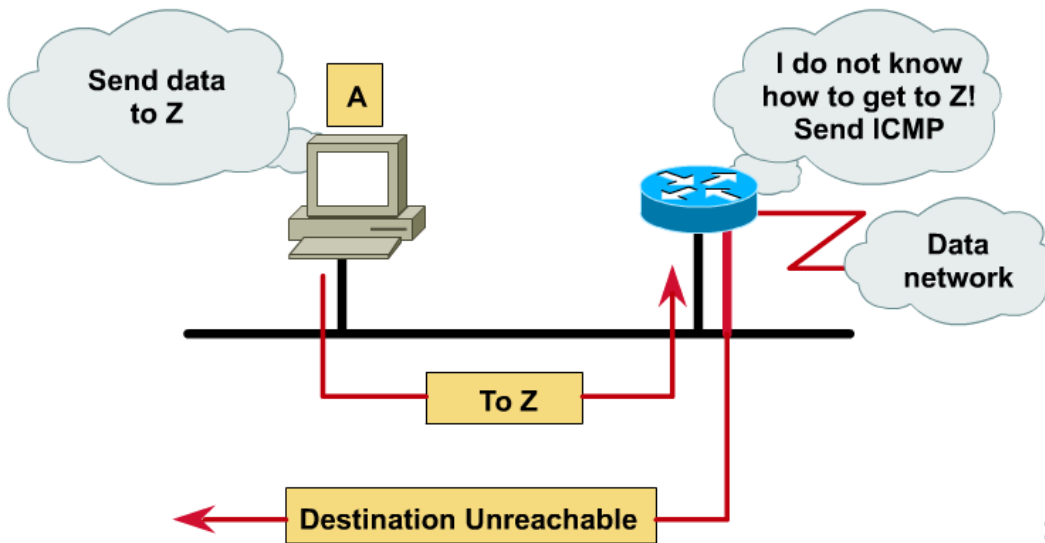
Internet Control Message Protocol



ICMP uses the following types of defined messages.

- Destination Unreachable
 - Time to Live Exceeded
 - Parameter Problem
 - Source Quench
 - Redirect
- Echo
 - Echo Reply
 - Timestamp
 - Timestamp Reply
 - Information Request
 - Information Reply
 - Address Request
 - Address Reply

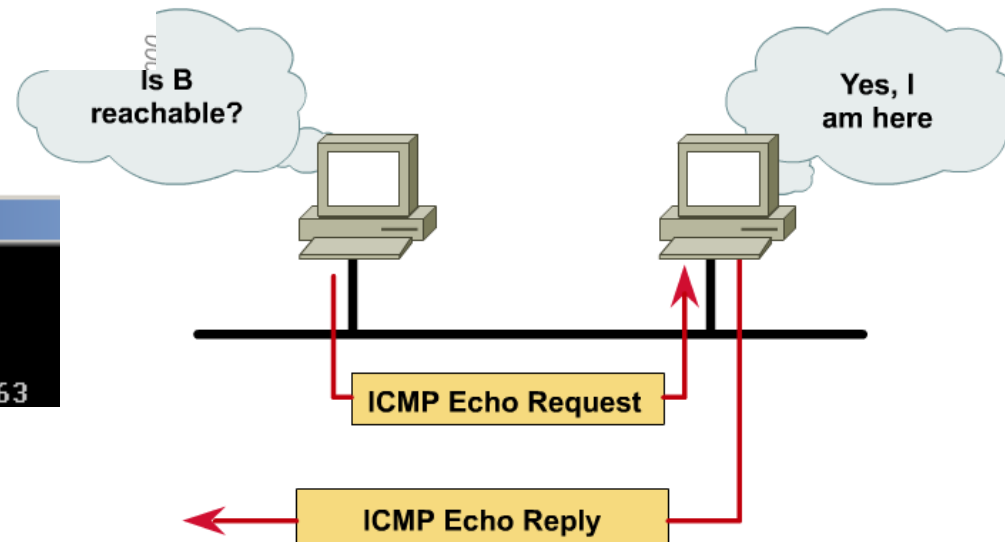
ICMP Testing



```
C:\WINDOWS\system32\cmd.exe
C:\>ping 192.168.100.1
Pinging 192.168.100.1 with 32 bytes of data:
Reply from 192.168.100.1: bytes=32 time=2ms TTL=63
```

Ping: ICMP Echo Request and Echo Reply

P Testing



- We will discuss ping, echo request and echo reply, in detail in the presentation ICMP – Understanding Ping and Traceroute.

ICMP Echo Request (ping)

No. ↓	Time	Source	Destination	Protocol	Info
5	7.166694	192.168.1.100	192.168.100.1	ICMP	Echo (ping) request
6	7.169161	192.168.100.1	192.168.1.100	ICMP	Echo (ping) reply

[-] Frame 5 (74 bytes on wire (58 bytes captured) on interface 0)
[-] Ethernet II, Src: 192.168.1.100 (00:0a:e4:d4:4c:f3), Dst: 192.168.1.1 (00:0f:66:09:4e:0f) Destination: 192.168.1.1 (00:0f:66:09:4e:0f) Source: 192.168.1.100 (00:0a:e4:d4:4c:f3) Type: IP (0x0800)
[-] Internet Protocol, Src: 192.168.1.100 (192.168.1.100), Dst: 192.168.100.1 (192.168.100.1) Version: 4 ← Header length: 20 bytes Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00) Total Length: 60 Identification: 0x3c22 (15394) Flags: 0x00 Fragment offset: 0 Time to live: 128 ← Protocol: ICMP (0x01) ← Header checksum: 0x17e9 [correct] Source: 192.168.1.100 (192.168.1.100) ← Destination: 192.168.100.1 (192.168.100.1) ←
[-] Internet Control Message Protocol Type: 8 (Echo (ping) request) Code: 0 Checksum: 0x425c [correct] Identifier: 0x0200 Sequence number: 0x0900 Data (32 bytes)

ICMP Echo Reply (ping)

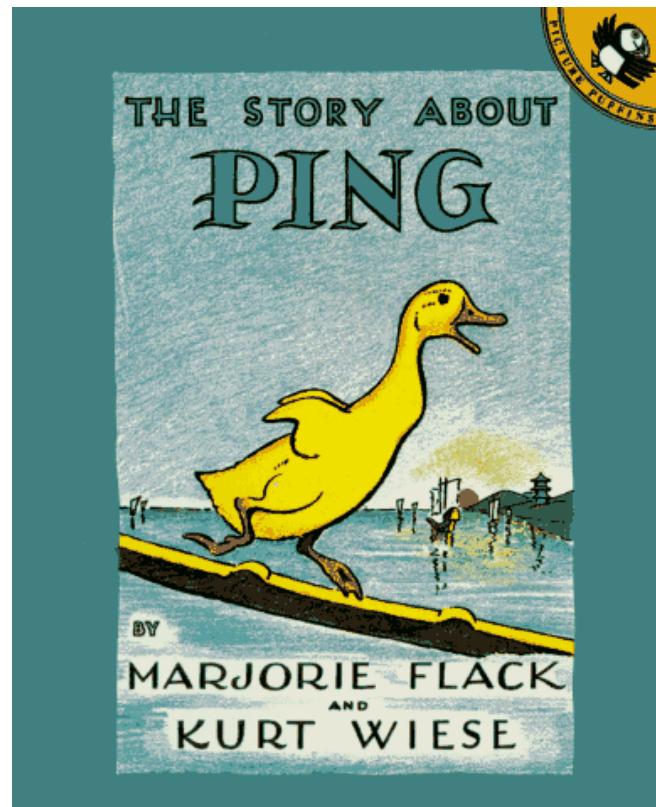
No. ↓	Time	Source	Destination	Protocol	Info
5	7.166694	192.168.1.100	192.168.100.1	ICMP	Echo (ping) request
6	7.169161	192.168.100.1	192.168.1.100	ICMP	Echo (ping) reply

[-] Frame 6 (74 bytes on wire, 74 bytes captured)
[-] Ethernet II, Src: 192.168.1.1 (00:0f:66:09:4e:0f), Dst: 192.168.1.100 (00:0a:e4:d4:4c:f3) Destination: 192.168.1.100 (00:0a:e4:d4:4c:f3) Source: 192.168.1.1 (00:0f:66:09:4e:0f) Type: IP (0x0800)
[-] Internet Protocol, Src: 192.168.100.1 (192.168.100.1), Dst: 192.168.1.100 (192.168.1.100) Version: 4 Header length: 20 bytes [-] Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00) Total Length: 60 Identification: 0x0036 (54) [-] Flags: 0x00 Fragment offset: 0 Time to live: 63 Protocol: ICMP (0x01) [-] Header checksum: 0x94d5 [correct] Source: 192.168.100.1 (192.168.100.1) Destination: 192.168.1.100 (192.168.1.100)
[-] Internet Control Message Protocol Type: 0 (Echo (ping) reply) Code: 0 Checksum: 0x4a5c [correct] Identifier: 0x0200 Sequence number: 0x0900 Data (32 bytes)

For more information on Ping

Here are two options for more information on Ping:

- See my PowerPoint presentation: **ICMP – Understanding Ping and Trace**
- Read the book: **The Story About Ping** ☺
by Marjorie Flack, Kurt Wiese (See a Amazon.com customer review on next slide – very funny!)



Review of Story of Ping on Amazon.com

8271 of 8518 people found the following review helpful:

Ping! I love that duck!, January 25, 2000

Reviewer: John E. Fracisco (El Segundo, CA USA)

Using deft allegory, the authors have provided an insightful and intuitive explanation of one of Unix's most venerable networking utilities. Even more stunning is that they were clearly working with a very early beta of the program, as their book first appeared in 1933, years (decades!) before the operating system and network infrastructure were finalized.

The book describes networking in terms even a child could understand, choosing to anthropomorphize the underlying packet structure. The ping packet is described as a duck, who, with other packets (more ducks), spends a certain period of time on the host machine (the wise-eyed boat). At the same time each day (I suspect this is scheduled under cron), the little packets (ducks) exit the host (boat) by way of a bridge (a bridge). From the bridge, the packets travel onto the internet (here embodied by the Yangtze River).

The title character -- er, packet, is called Ping. Ping meanders around the river before being received by another host (another boat). He spends a brief time on the other boat, but eventually returns to his original host machine (the wise-eyed boat) somewhat the worse for wear.

If you need a good, high-level overview of the ping utility, this is the book. I can't recommend it for most managers, as the technical aspects may be too overwhelming and the basic concepts too daunting.

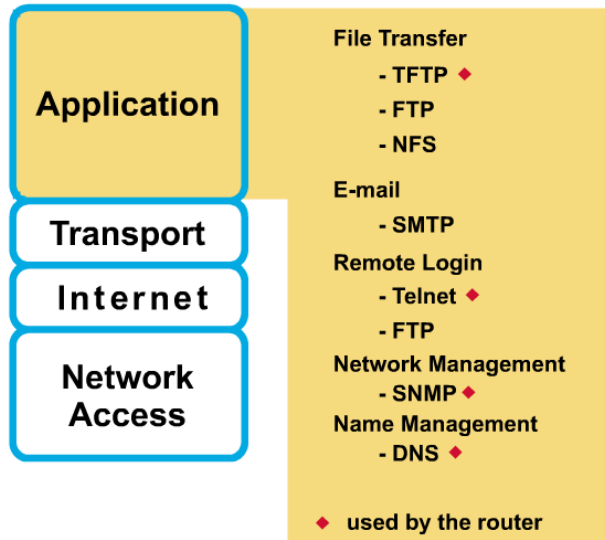
Problems With This Book

As good as it is, The Story About Ping is not without its faults. There is no index, and though the ping(8) man pages cover the command line options well enough, some review of them seems to be in order. Likewise, in a book solely about Ping, I would have expected a more detailed overview of the ICMP packet structure.

But even with these problems, The Story About Ping has earned a place on my bookshelf, right between Stevens' Advanced Programming in the Unix Environment, and my dog-eared copy of Dante's seminal work on MS Windows, Inferno. Who can read that passage on the Windows API ("Obscure, profound it was, and nebulous, So that by fixing on its depths my sight -- Nothing whatever I discerned therein."), without shaking their head with deep understanding. But I digress. --*This text refers to the School & Library Binding edition.*

Ping – A TCP/IP Application

Application Layer



ping

Command Output

```
C:\WINDOWS\Desktop>ping cisco.netacad.net

Pinging cisco.netacad.net [4.22.32.50] with 32 bytes of data:

Reply from 4.22.32.50: bytes=32 time=53ms TTL=4
Reply from 4.22.32.50: bytes=32 time=38ms TTL=4
Reply from 4.22.32.50: bytes=32 time=44ms TTL=4
Reply from 4.22.32.50: bytes=32 time=33ms TTL=4

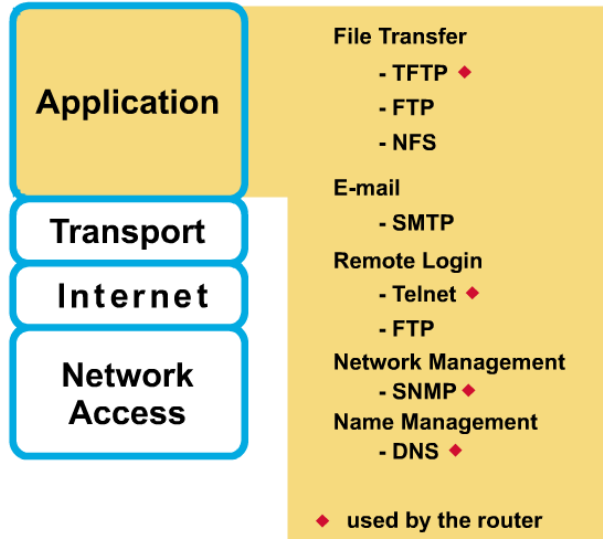
Ping statistics for 4.22.32.50:
    Packets: Sent = 4, Received = 4, Lost = 0 (0% loss),
    Approximate round trip times in milli-seconds:
        Minimum = 33ms, Maximum = 53ms, Average = 42ms

C:\WINDOWS\Desktop>_
```

- **PING** (Packet Internet Groper) is a diagnostic utility used to determine whether a computer is properly connected to devices/Internet.
- More in a later presentation!

Traceroute – A TCP/IP Application

Application Layer



tracert

Command Output

```
C:\WINDOWS\Desktop>tracert cisco.netacad.net
```

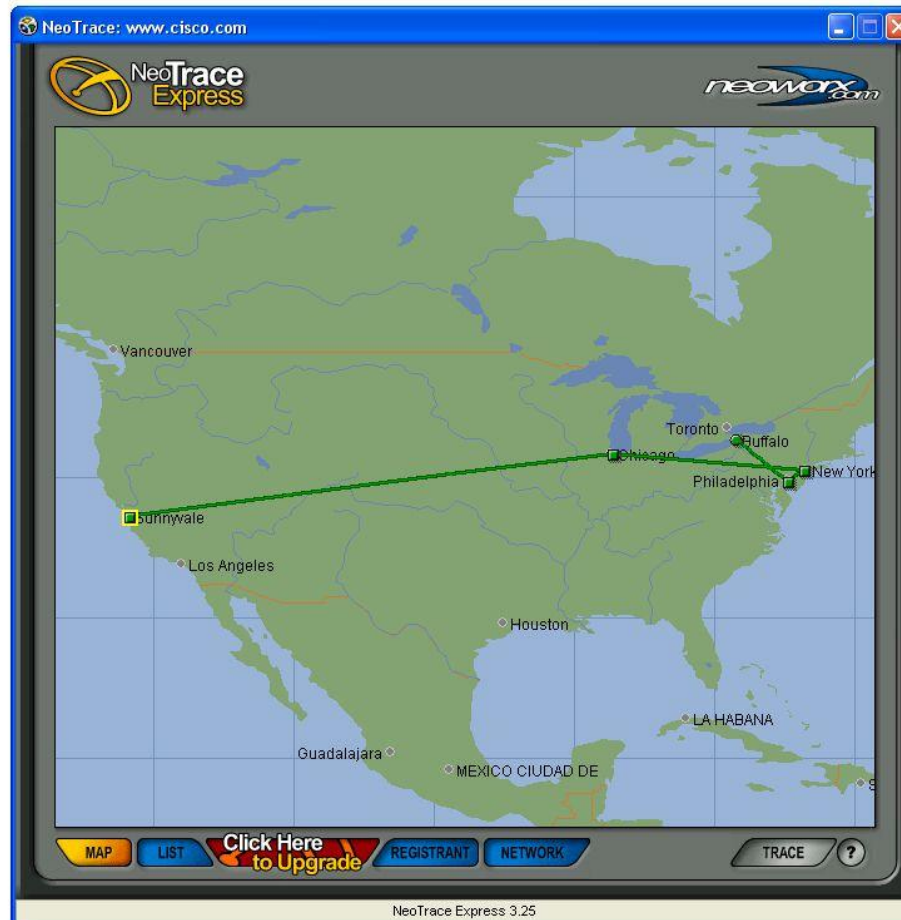
```
Tracing route to cisco.netacad.net [4.22.32.50]  
over a maximum of 30 hops:
```

1	1 ms	1 ms	<10 ms	two-little.cisco.com [171.70.134.3]
2	1 ms	<10 ms	<10 ms	b2-bomber.cisco.com [171.68.2.217]
3	1 ms	1 ms	1 ms	gsr-knoc.cisco.com [171.68.2.33]
4	1 ms	1 ms	1 ms	gaza-gw2.cisco.com [171.68.2.254]
5	1 ms	2 ms	1 ms	sj-wall-1.cisco.com [198.92.1.137]
6	2 ms	2 ms	3 ms	barrnet-gw.cisco.com [192.31.7.37]
7	5 ms	4 ms	4 ms	s2-1-1.paloalto-cr18.bbnplanet.net [4.1.142.237]
8	5 ms	4 ms	8 ms	p3-2.paloalto-nbr2.bbnplanet.net [4.0.3.85]
9	6 ms	6 ms	5 ms	p4-0.sanjosel-nbr1.bbnplanet.net [4.0.1.2]
10	14 ms	11 ms	12 ms	p4-2.lsajcal-nbr2.bbnplanet.net [4.0.1.18]
11	13 ms	12 ms	16 ms	p5-0.lsajcal-nbr1.bbnplanet.net [4.24.4.21]
12	15 ms	13 ms	15 ms	p10-0-0.lal-br1.bbnplanet.net [4.24.4.93]
13	29 ms	27 ms	24 ms	s0-0-0.phnyaz2-cr1.bbnplanet.net [4.0.3.186]
14	30 ms	26 ms	27 ms	s0.unicon.bbnplanet.net [4.1.110.10]
15	+	+	+	Request timed out

- **Traceroute** is a program that is available on many systems, and is similar to PING, except that traceroute provides more information than PING.
- Traceroute traces the path a packet takes to a destination, and is used to debug routing problems.
- More in a later presentation!

Traceroute – A TCP/IP Application

- Graphical Trace Programs like NeoTrace (now by McAfee)
- <http://www.networkingfiles.com/PingFinger/Neotraceexpress.htm>



Windows: tracer command

```
C:\WINDOWS\system32\cmd.exe
C:\>tracert www.ucsc.edu


Tracing route to ucsc.edu [128.114.124.7]
over a maximum of 30 hops:

 1      <1 ms      <1 ms      <1 ms      192.168.1.1
 2          *          *          *          Request timed out.
 3      11 ms      11 ms      10 ms      68.87.198.149
 4      10 ms      11 ms       9 ms      68.87.192.45
 5      21 ms      12 ms      17 ms      68.87.192.41
 6      43 ms      28 ms      16 ms      68.87.192.37
 7      37 ms      35 ms      12 ms      68.87.195.62
 8      15 ms      16 ms      36 ms      12.117.240.5
 9      17 ms      70 ms      20 ms      tbr1011001.sffca.ip.att.net [12.122.82.74]
10      23 ms      11 ms      11 ms      ggr1-p360.sffca.ip.att.net [12.123.13.65]
11      17 ms      17 ms      18 ms      64.200.151.17
12      14 ms      21 ms      29 ms      sntcca1wcx1-pos2-1.wcg.net [64.200.149.45]
13      20 ms      18 ms      20 ms      snfcca1wcx3-pos15-0.wcg.net [64.200.249.61]
14      23 ms      26 ms      39 ms      corp-for-ed-snfcca1wcx3-gige-10-1-33.wcg.net [64.
.200.198.98]
15      19 ms      40 ms      48 ms      dc-oak-isp--sfo-isp-t1.cenic.net [137.164.40.200]
16      25 ms      24 ms      25 ms      dc-ucsc--oak-isp-egm.cenic.net [137.164.23.62]
17      29 ms      24 ms      29 ms      isb-g-GE2-4.ucsc.edu [128.114.0.42]
18      40 ms      25 ms      26 ms      secure2-g-g1-0-26.ucsc.edu [128.114.101.202]
19      76 ms      24 ms      26 ms      fw-dc1.ucsc.edu [128.114.102.2]
20      86 ms      65 ms      23 ms      ucsc.edu [128.114.124.7]

Trace complete.
```

NeoTrace Map View

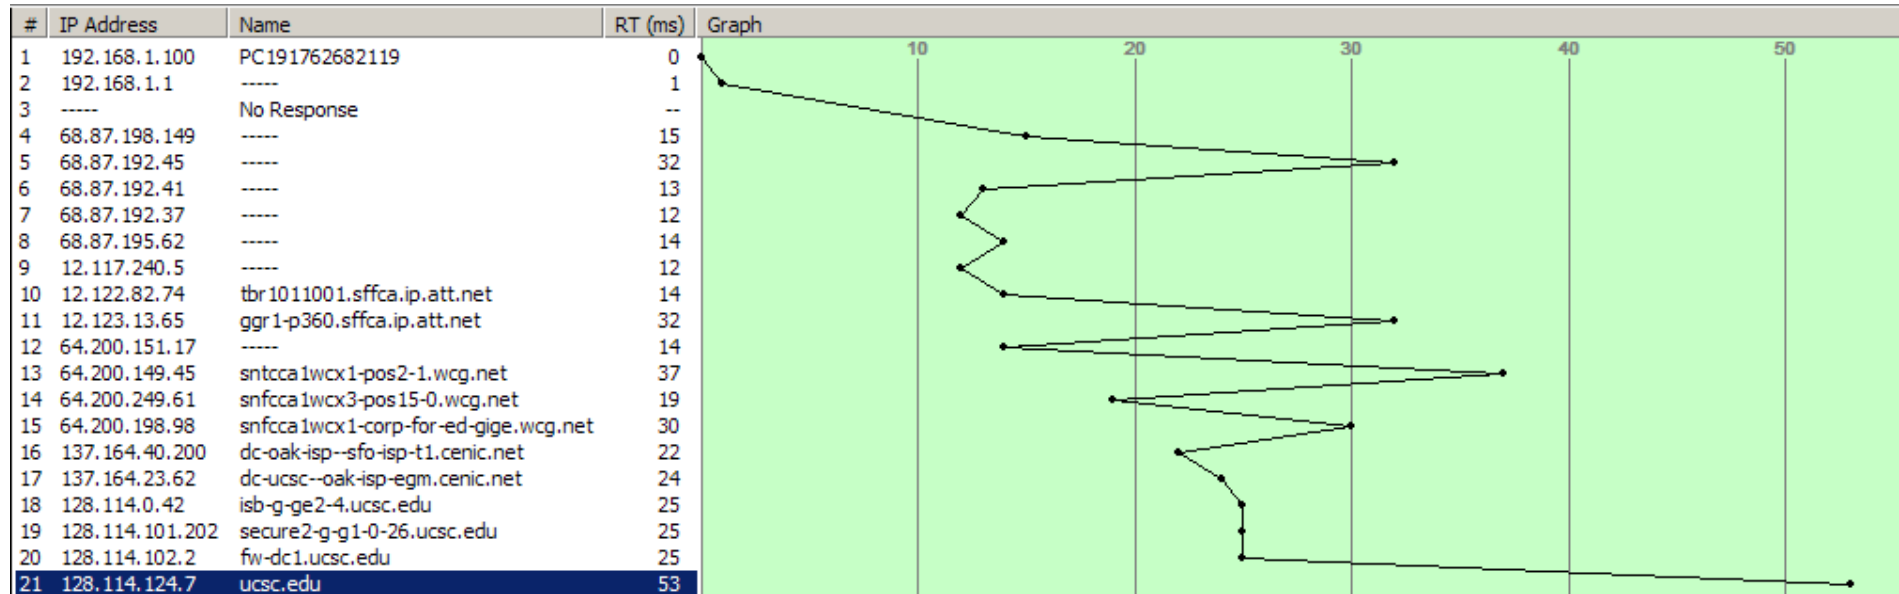
Cabrillo College



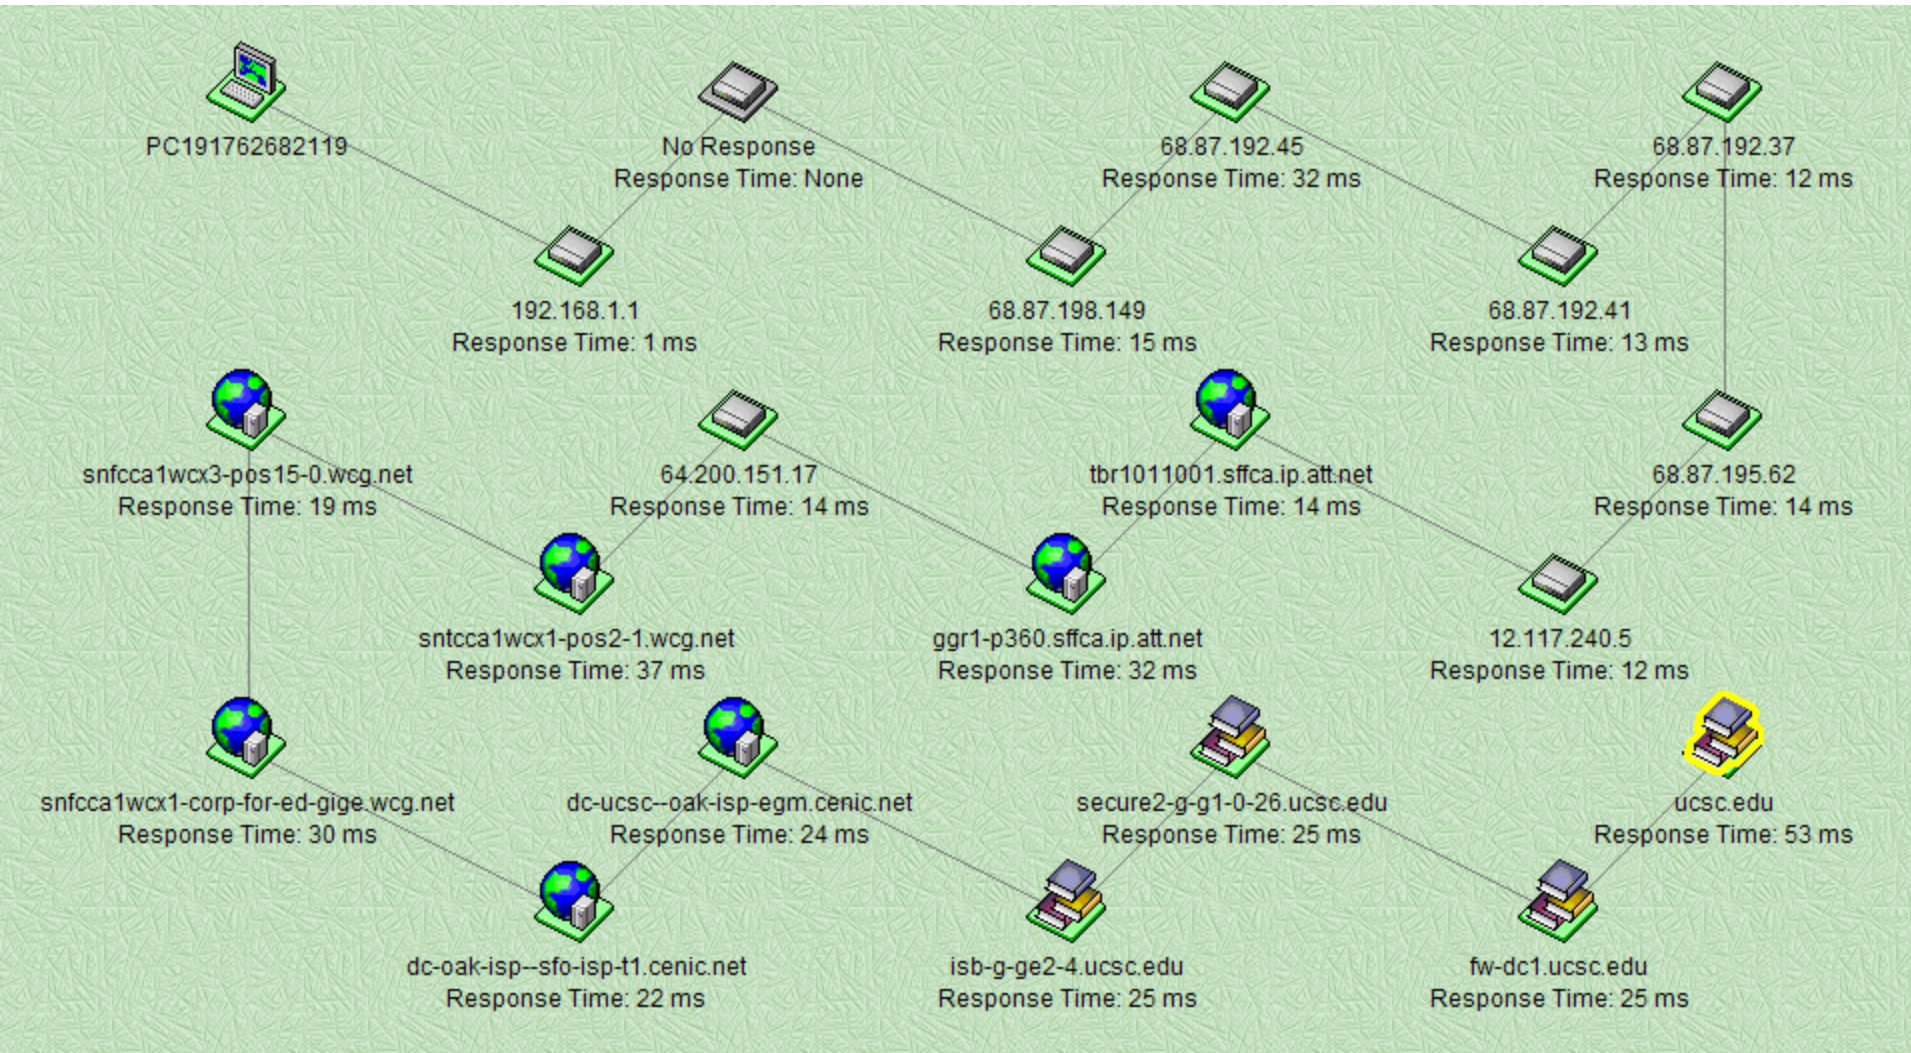
The map shows a network trace route across the United States. The route starts in San Jose, California, and ends in Philadelphia, Pennsylvania. The route is marked with a green line and passes through San Francisco, Chicago, Detroit, Toronto, and Montréal. The map is a satellite view with green land and blue water.

Node 10 of 21	Node 4 of 21
<p>tbr1011001.sffca.ip.att.net</p> <p>External Apps</p> <p>Previous Node 10 of 21 Next</p> <p>Name: tbr1011001.sffca.ip.att.net IP Address: 12.122.82.74 Location: San Francisco (37.775N, 122.417W) Network: Unknown</p> <p>Registrant: AT&T Corp 55 Corporate Drive Bridgewater, NJ 08807 US</p>	<p>68.87.198.149</p> <p>External Apps</p> <p>Previous Node 4 of 21 Next</p> <p>Name: Unknown IP Address: 68.87.198.149 Location: Philadelphia (39.950N, 75.167W) Network: Unknown</p> <p>Registrant contact information is not available</p>

NeoTrace List View

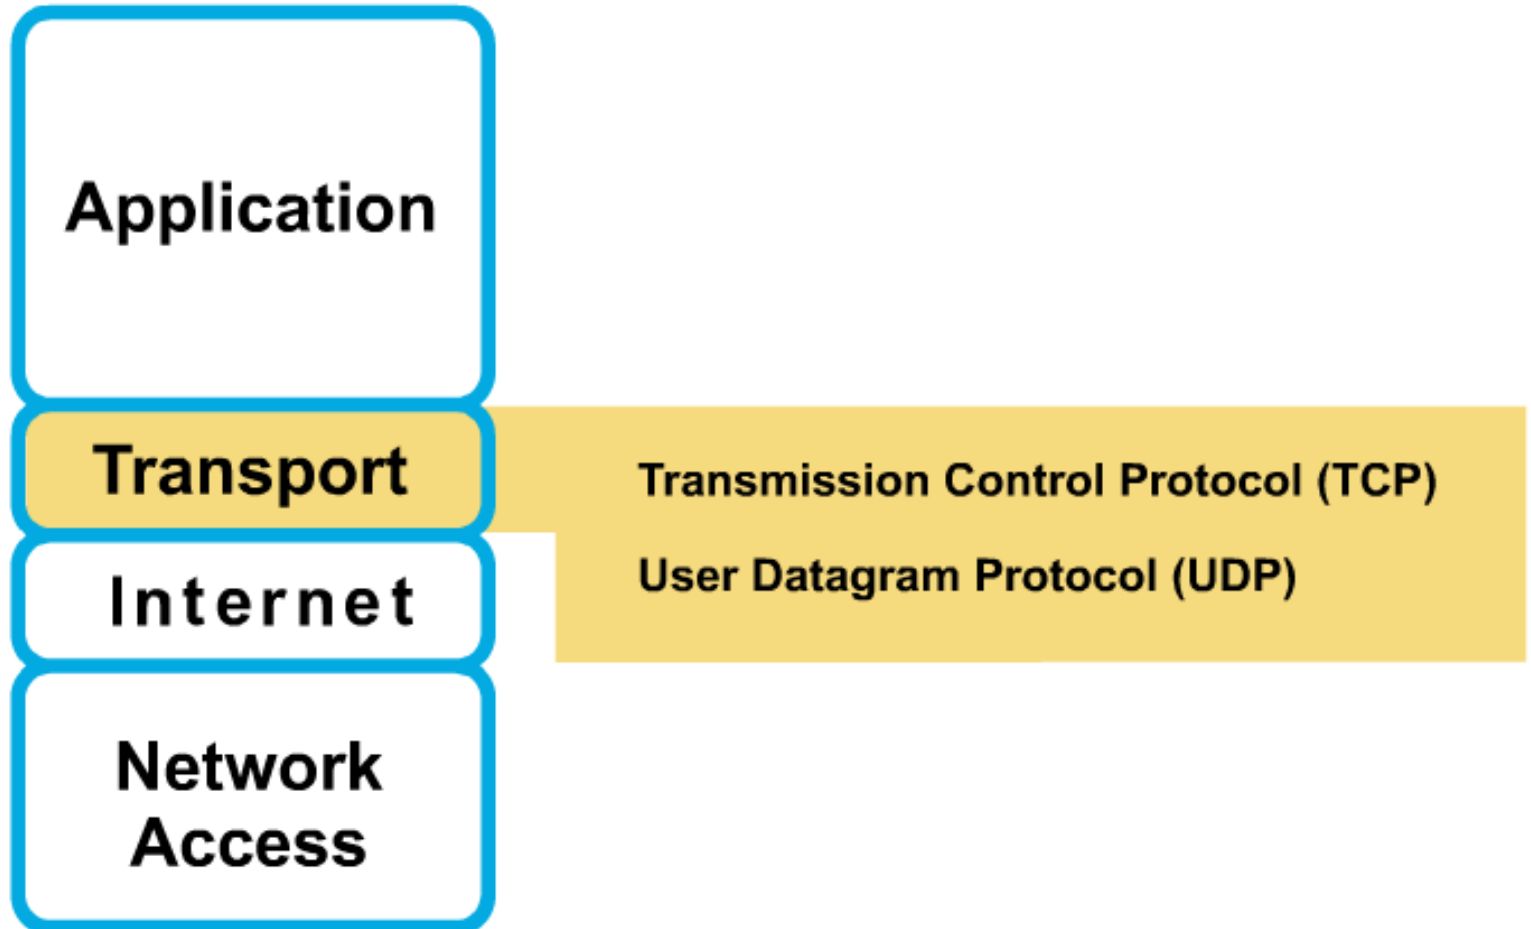


NeoTrace Node View



Layer 4: TCP/IP Transport Layer

Transport Layer Overview



Topics

Layer 3 Concepts

- TCP/IP and the Internet Layer
- Diagram the IP datagram
- Internet Control Message Protocol (ICMP)

TCP/IP protocol stack and the transport layer

- TCP and UDP segment format
- TCP and UDP port numbers
- TCP three-way handshake/open connection
- TCP simple acknowledgment and windowing

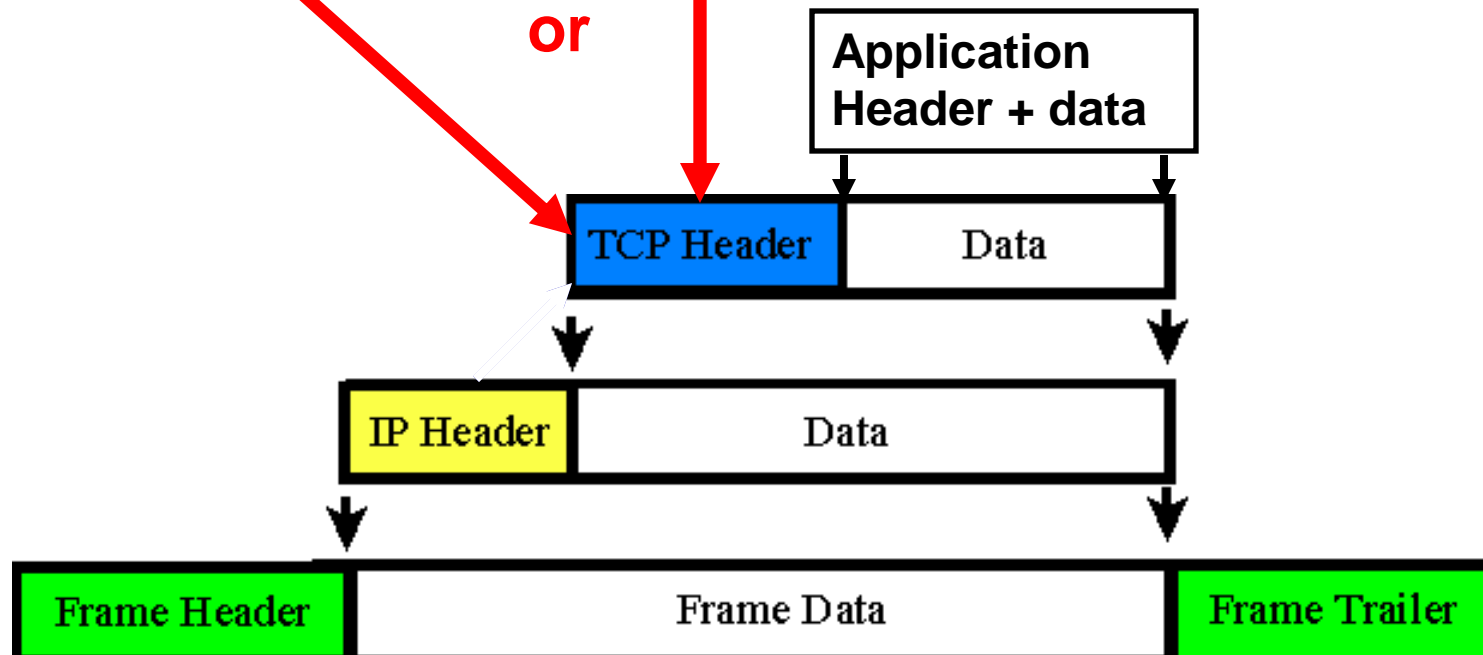
TCP Header

Source Port (16 bits)				Destination Port (16 bits)				
Sequence Number (32 bits)								
Acknowledgement Number (32 bits)								
Data Offset (4 bits)	Reserved (6 bits)	URG	ACK	PSH	RST	SYN	FIN	Window (16 bits)
Checksum (16 bits)							Urgent Pointer (16 bits)	
Options and Padding								

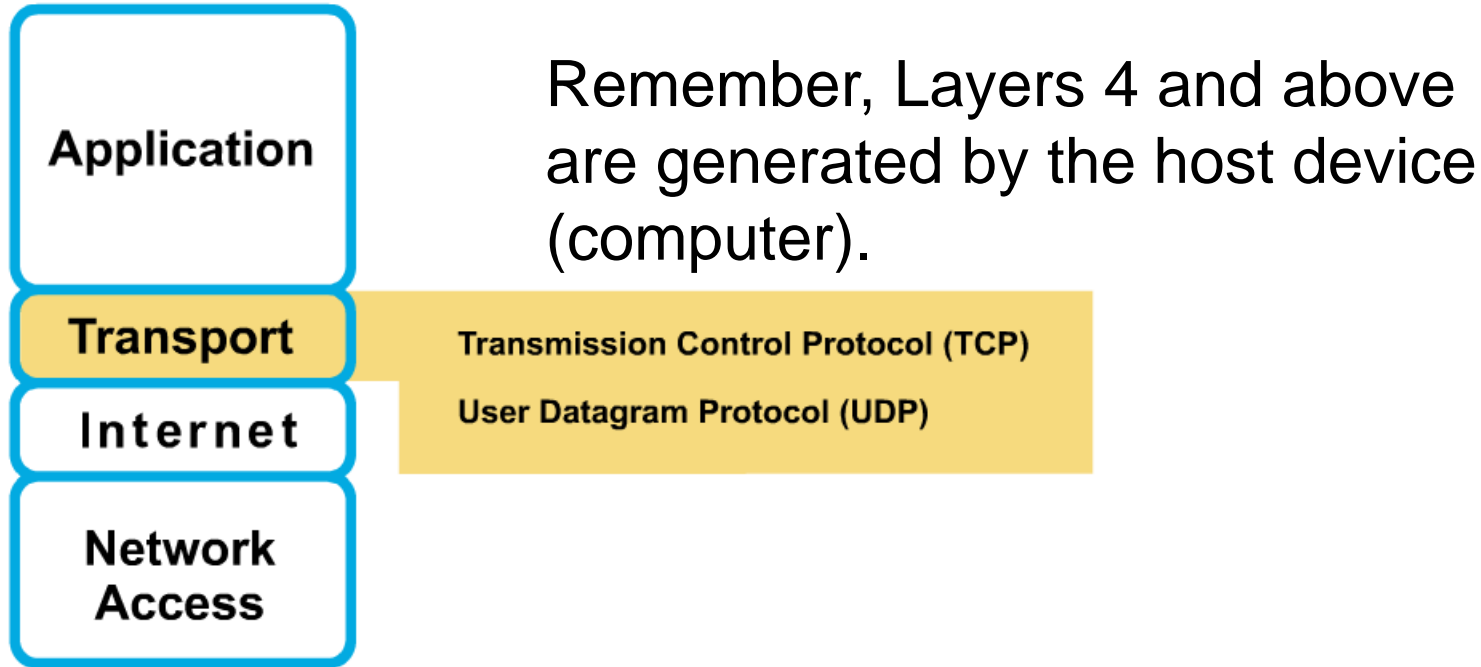
UDP Header

Source Port (16 bits)	Destination Port (16 bits)
Length (16 bits)	Checksum (16 bits)
Data....	

or

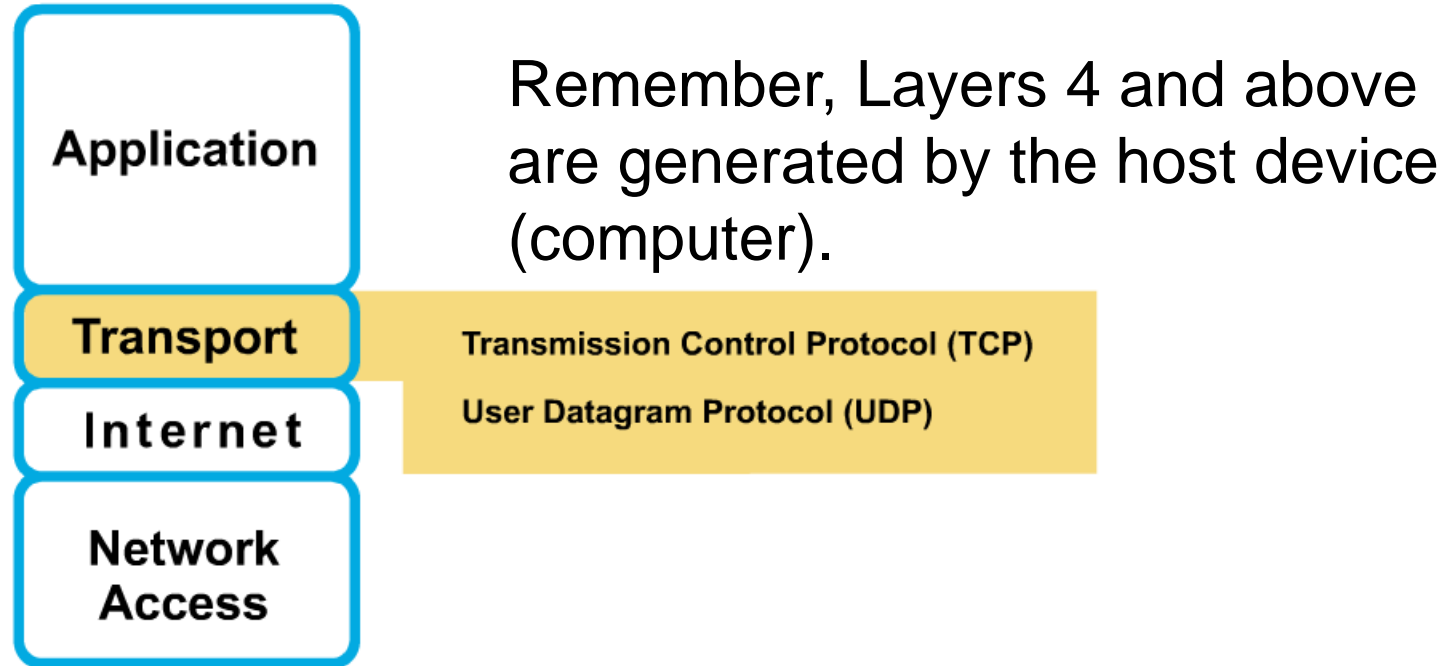


Transport Layer Overview



- The transport layer enables a user's device to **segment** several upper-layer applications for placement on the same Layer 4 data stream, and enables a receiving device to reassemble the upper-layer application segments.
- The Layer 4 data stream is a logical connection between the endpoints of a network, and provides transport services from a host to a destination.
- This service is sometimes referred to as **end-to-end service**.

Transport Layer Overview



The transport layer also provides two protocols

- **TCP** – Transmission Control Protocol
- **UDP** – User Datagram Protocol

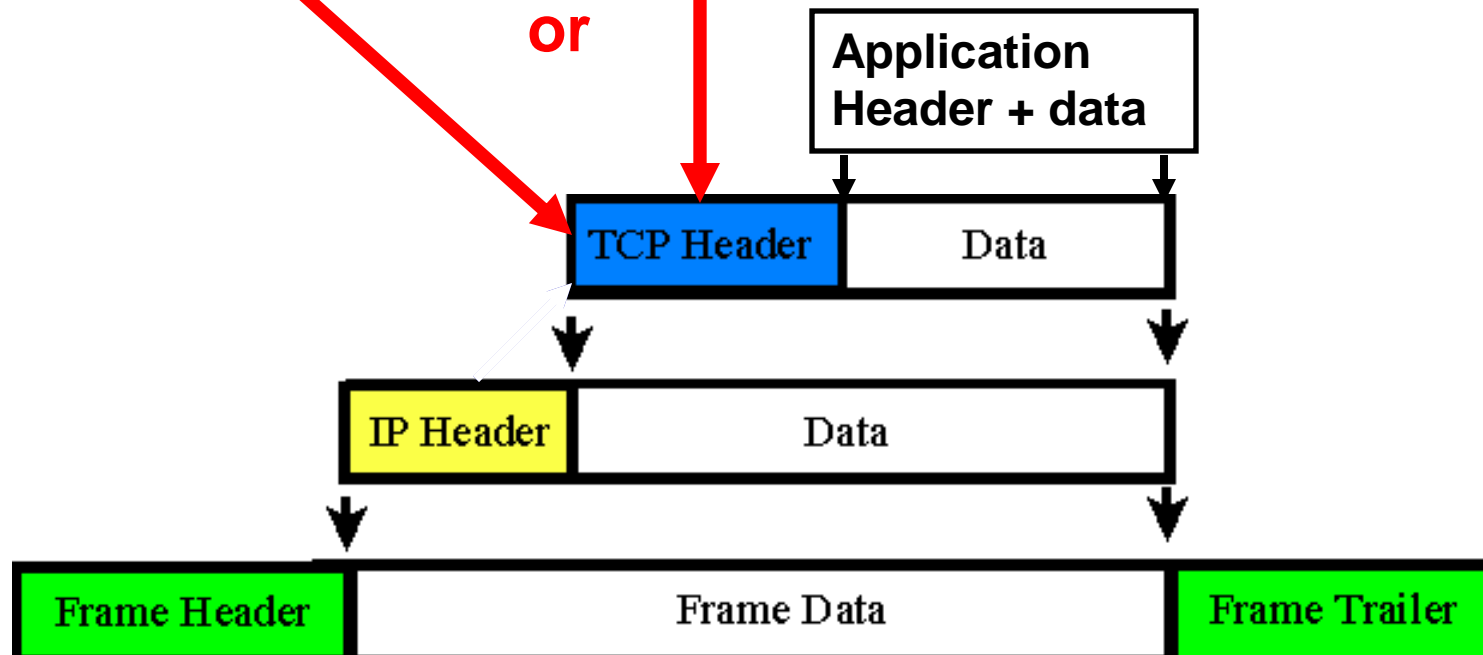
TCP Header

Source Port (16 bits)				Destination Port (16 bits)				
Sequence Number (32 bits)								
Acknowledgement Number (32 bits)								
Data Offset (4 bits)	Reserved (6 bits)	URG	ACK	PSH	RST	SYN	FIN	Window (16 bits)
Checksum (16 bits)							Urgent Pointer (16 bits)	
Options and Padding								

UDP Header

Source Port (16 bits)	Destination Port (16 bits)
Length (16 bits)	Checksum (16 bits)
Data....	

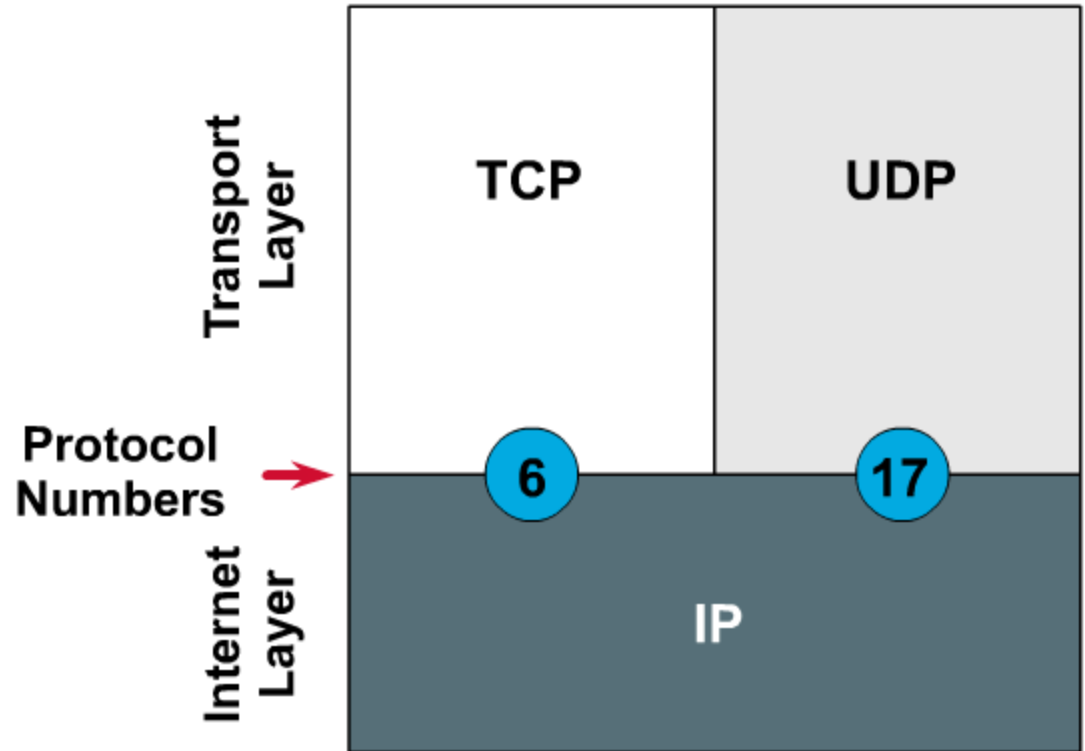
or



The Protocol Field

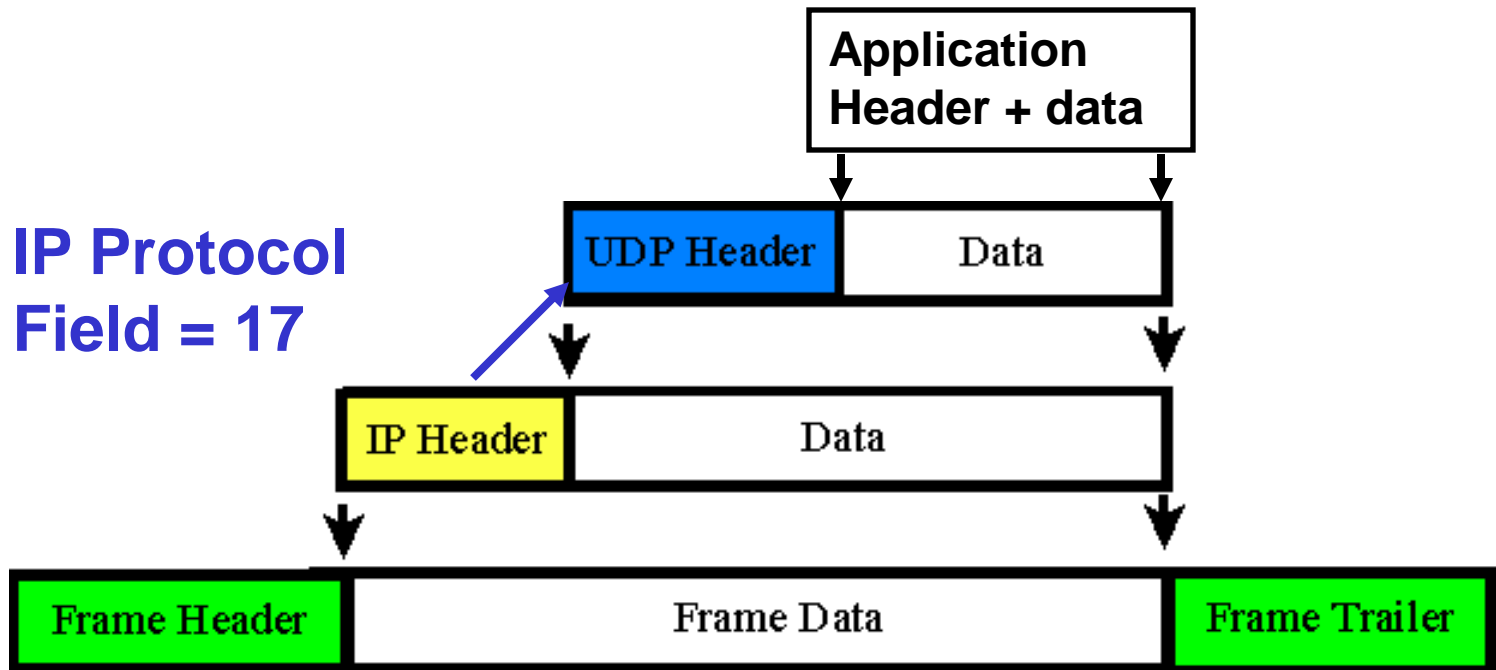
IP Header

0				15				16				31											
4-bit Version				4-bit Header Length				8-bit Type Of Service (TOS)				16-bit Total Length (in bytes)											
16-bit Identification								3-bit Flags				13-bit Fragment Offset											
8 bit Time To Live TTL				8-bit Protocol				16-bit Header Checksum															
												32-bit Source IP Address											
												32-bit Destination IP Address											
Options (if any)																							
Data																							

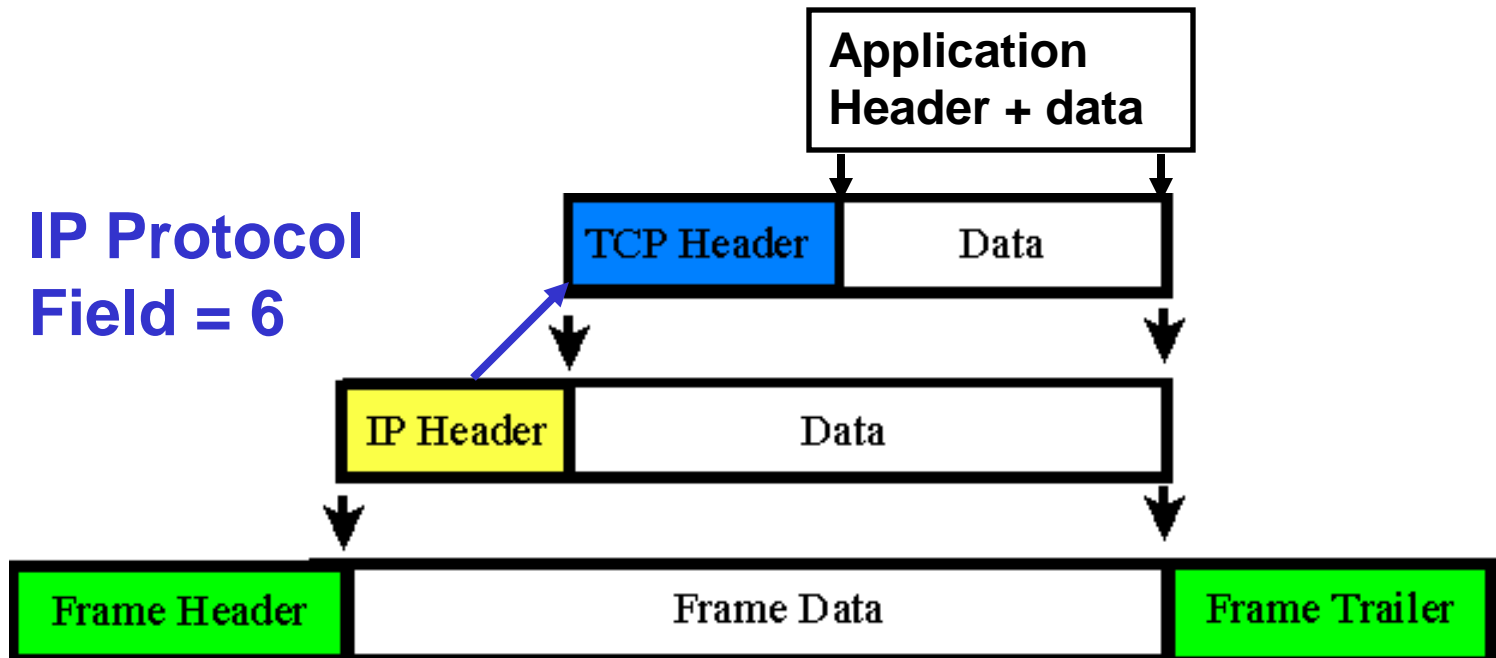


IP Packet has a **Protocol field** that specifies whether the segment is TCP or UDP.

**IP Protocol
Field = 17**



**IP Protocol
Field = 6**



Topics

Layer 3 Concepts

- TCP/IP and the Internet Layer
- Diagram the IP datagram
- Internet Control Message Protocol (ICMP)

TCP/IP protocol stack and the transport layer

- TCP and UDP segment format
- TCP and UDP port numbers
- TCP three-way handshake/open connection
- TCP simple acknowledgment and windowing

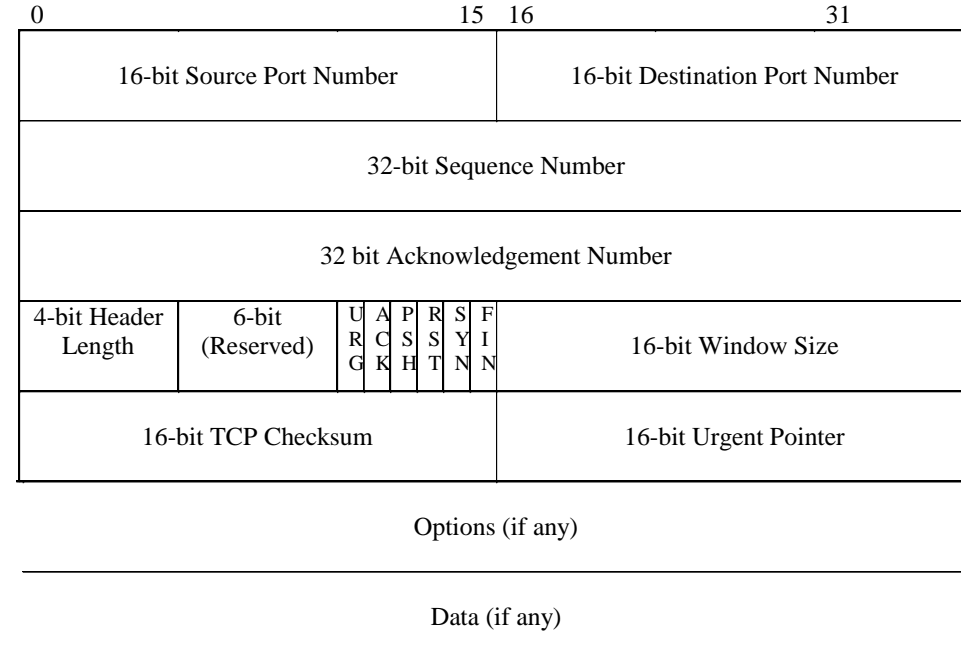
TCP Segment Header

0				15				16				31			
16-bit Source Port Number								16-bit Destination Port Number							
32-bit Sequence Number															
32 bit Acknowledgement Number															
4-bit Header Length		6-bit (Reserved)		U	A	P	R	S	F	16-bit Window Size					
				R	C	S	S	Y	I						
				G	K	H	T	N	N						
16-bit TCP Checksum								16-bit Urgent Pointer							

Options (if any)

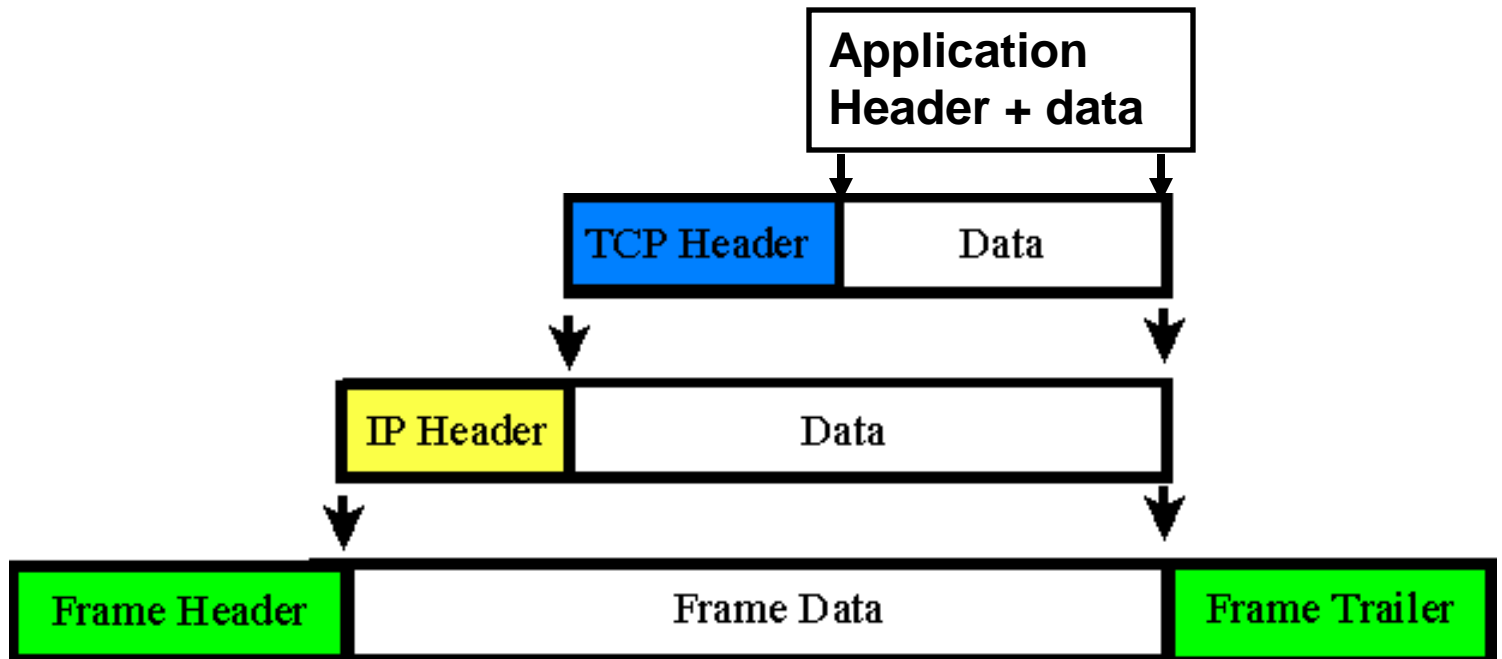
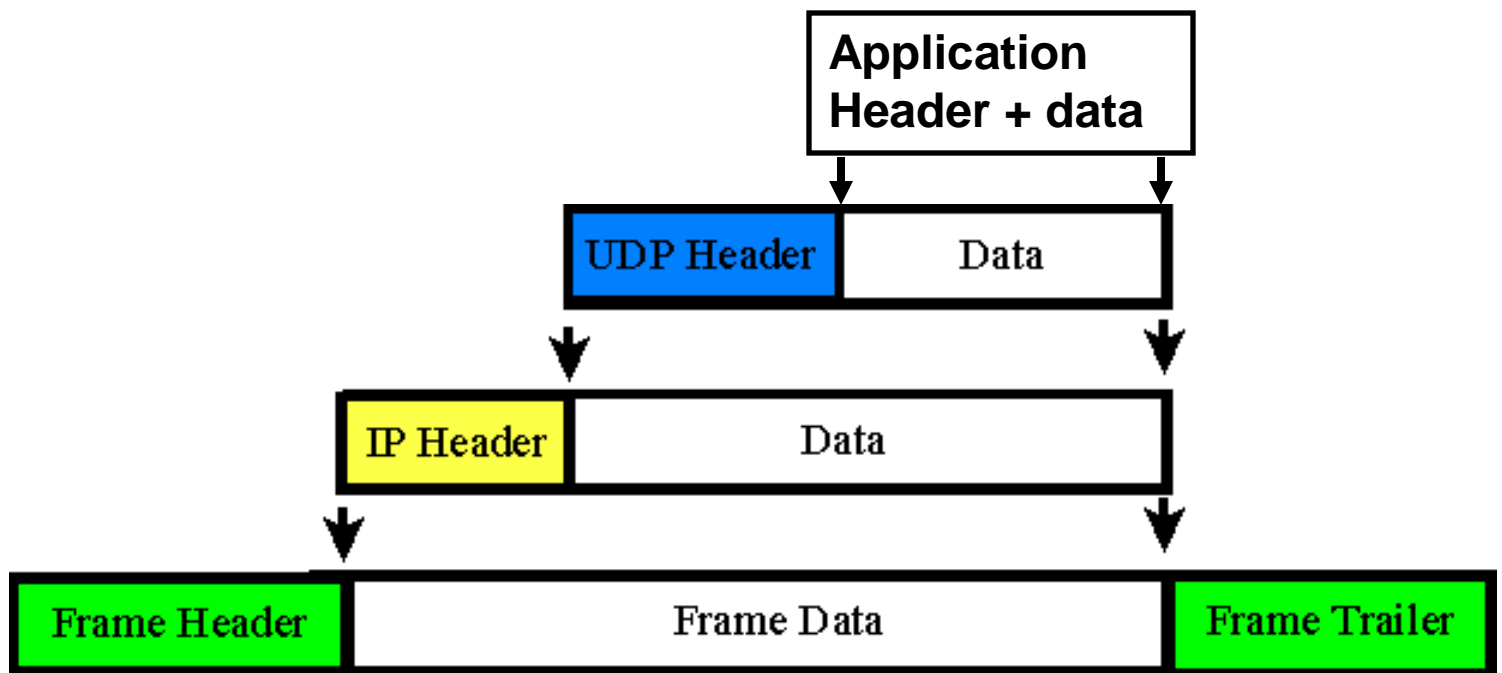
Data (if any)

TCP Segment Header

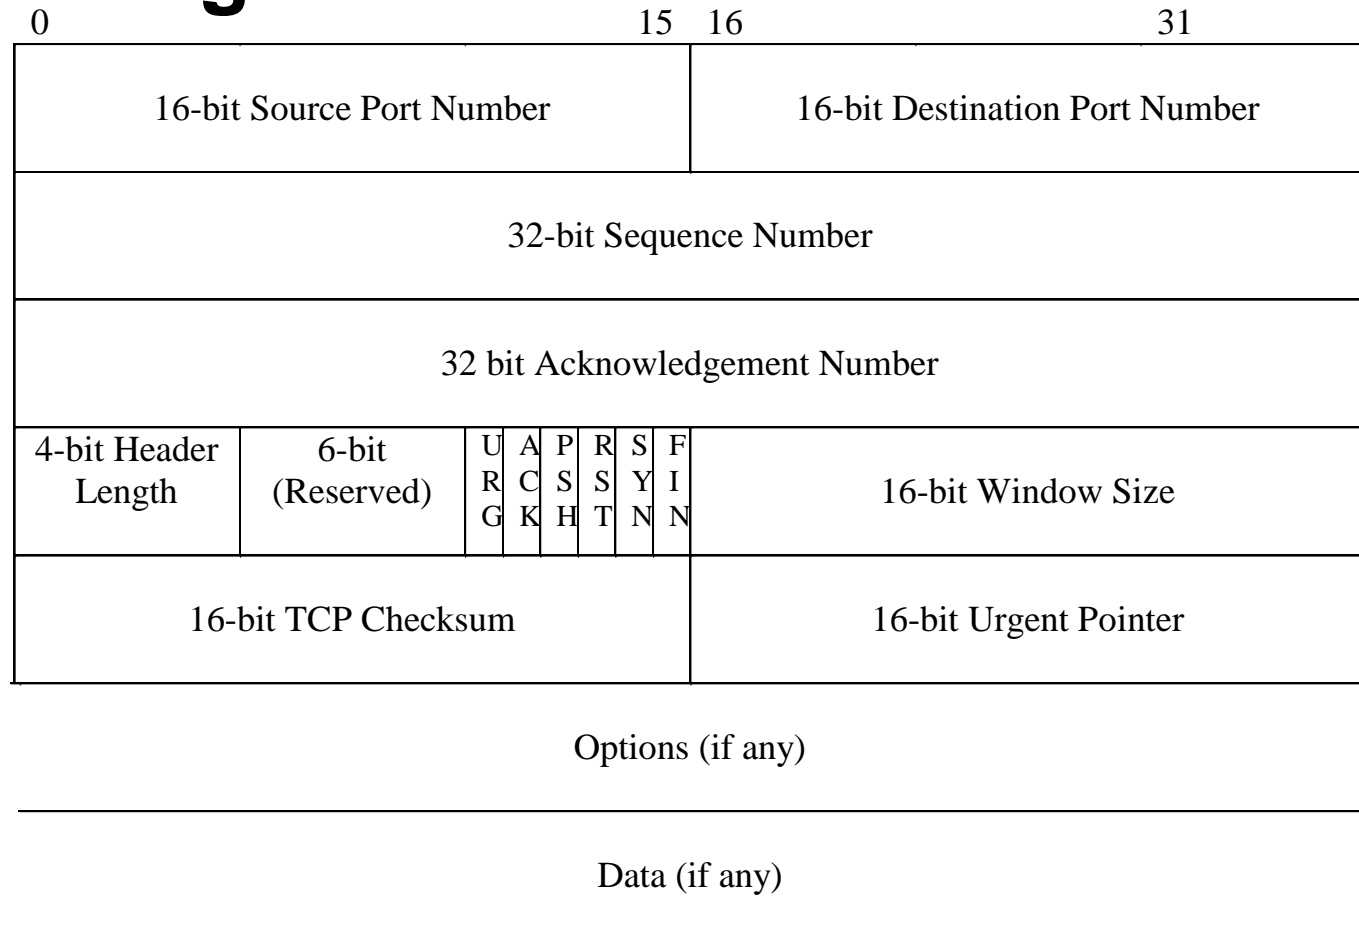


TCP (Transmission Control Protocol)

- **Connection-oriented, reliable protocol**
- Provides:
 - 1. flow control** by providing sliding windows,
 - 2. reliability** by providing sequence numbers and acknowledgments.
- TCP re-sends anything that is not received and supplies a **virtual circuit** between end-user applications.
- The advantage of TCP is that it provides **guaranteed delivery** of the segments.



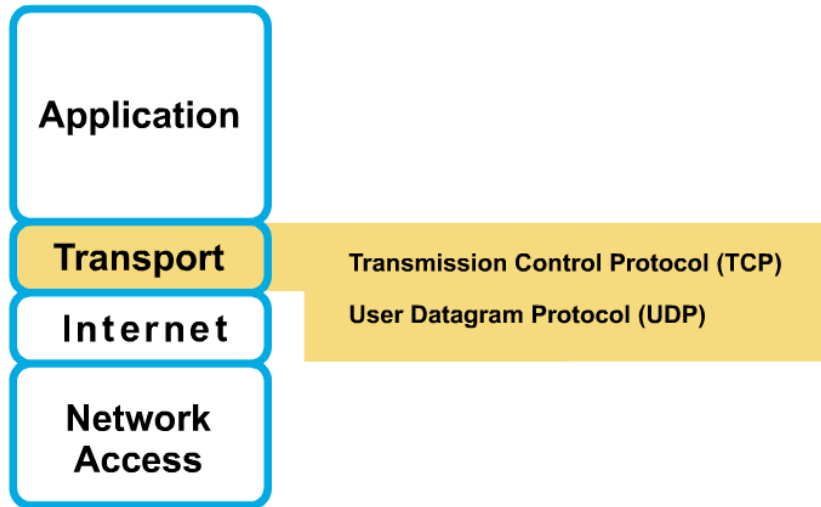
TCP Segment Header



Some of the protocols that use TCP are:

- HTTP
- Telnet
- FTP

Transport Layer Overview



0				15				16				31			
16-bit Source Port Number								16-bit Destination Port Number							
32-bit Sequence Number															
32 bit Acknowledgement Number															
4-bit Header Length		6-bit (Reserved)		U	A	P	R	S	F	16-bit Window Size					
				R	C	S	S	Y	I						
				G	K	H	T	N	N						
16-bit TCP Checksum								16-bit Urgent Pointer							
Options (if any)															
Data (if any)															

- **source port** -- the number of the calling port
- **destination port** -- the number of the called port
- **sequence number** -- the number used to ensure correct sequencing of the arriving data
- **acknowledgment number** -- the next expected TCP octet
- **HLEN** -- the number of 32-bit words in the header
- **reserved** -- set to 0
- **code bits** -- the control functions (e.g. setup and termination of a session)
- **window** -- the number of octets that the sender is willing to accept
- **checksum** -- the calculated checksum of the header and data fields
- **urgent pointer** -- indicates the end of the urgent data
- **option** -- one currently defined: maximum TCP segment size
- **data** -- upper-layer protocol data

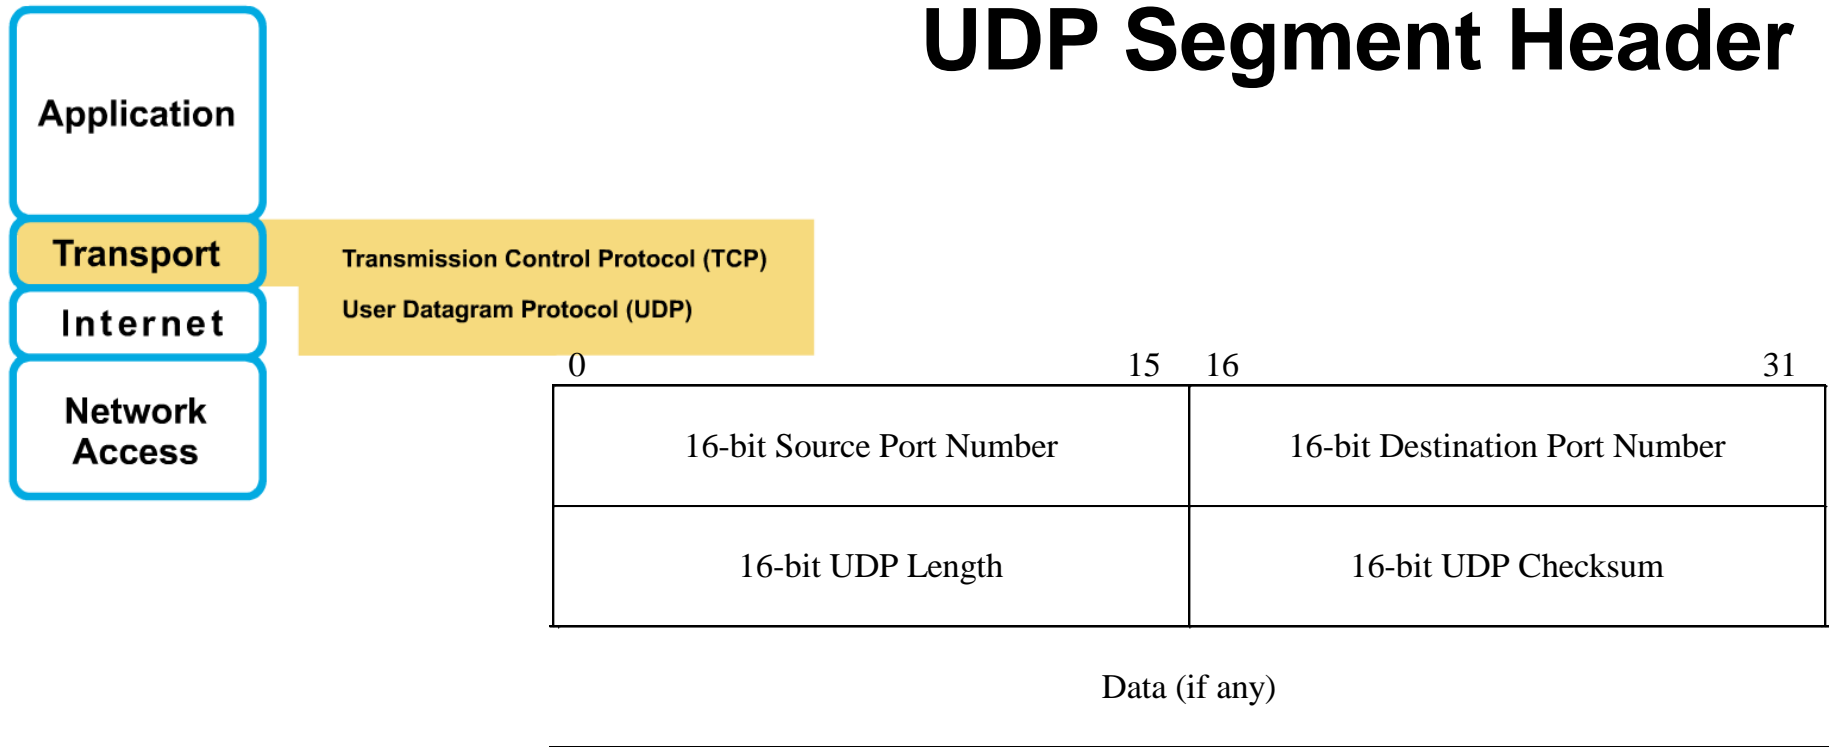
UDP Segment Header

0	15	16	31
16-bit Source Port Number		16-bit Destination Port Number	
16-bit UDP Length		16-bit UDP Checksum	

Data (if any)

- **UDP -- *connectionless and unreliable***; although responsible for transmitting messages, no software checking for segment delivery is provided at this layer.
- No flow control, no reliability.
- The advantage that UDP provides is **speed**.
- Since UDP provides no acknowledgments, less traffic is sent across the network, making the transfer faster.
- Protocols that use UDP include the following:
 - TFTP
 - SNMP
 - Network File System (NFS)
 - Domain Name System (DNS)

UDP Segment Header



- **source port** -- the number of the calling port
- **destination port** -- the number of the called port
- **UDP length** -- the length of the UDP header
- **checksum** -- the calculated checksum of the header and data fields
- **data** -- upper-layer protocol data

Topics

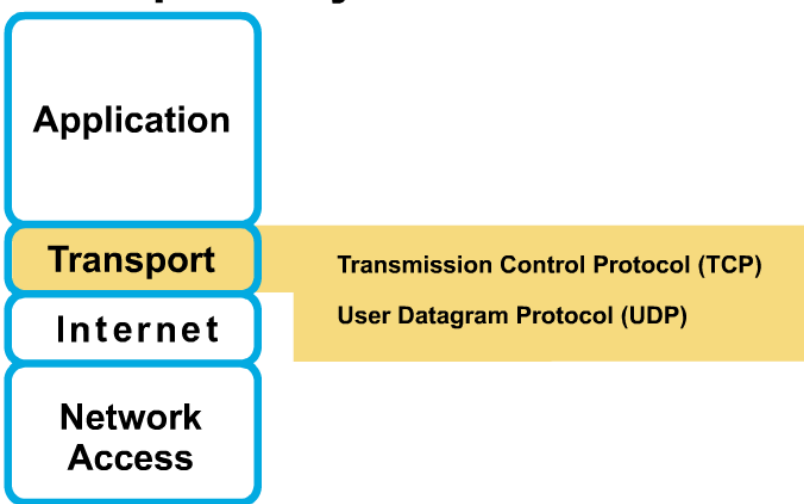
Layer 3 Concepts

- TCP/IP and the Internet Layer
- Diagram the IP datagram
- Internet Control Message Protocol (ICMP)

TCP/IP protocol stack and the transport layer

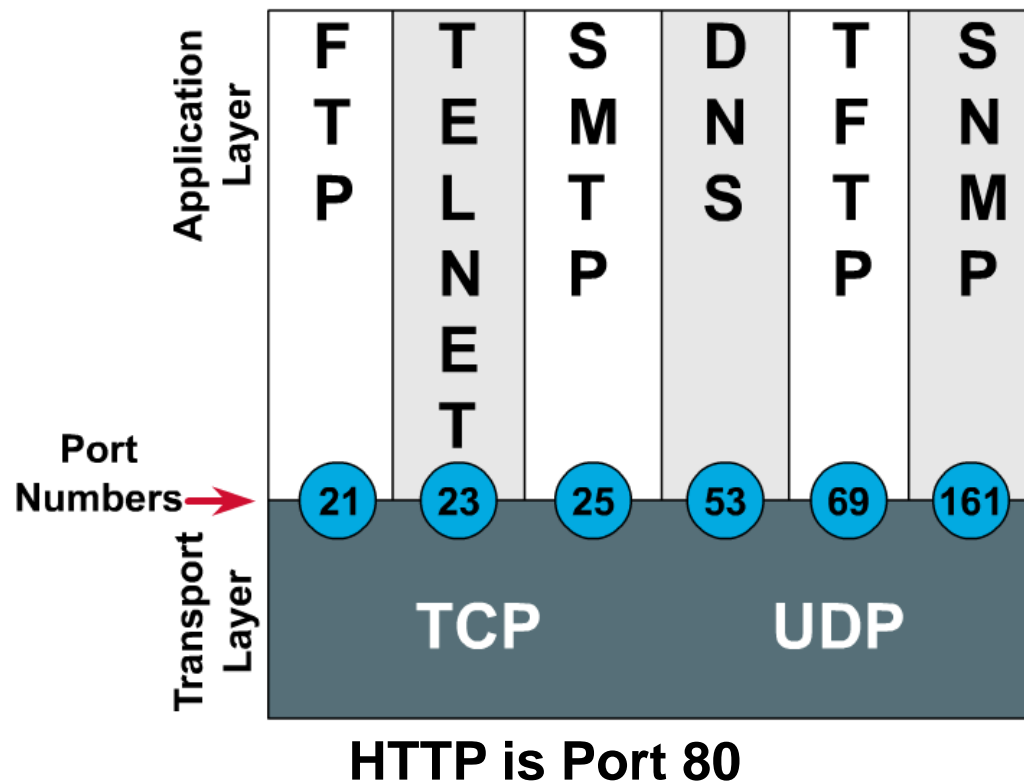
- TCP and UDP segment format
- TCP and UDP port numbers
- TCP three-way handshake/open connection
- TCP simple acknowledgment and windowing

Transport Layer Overview



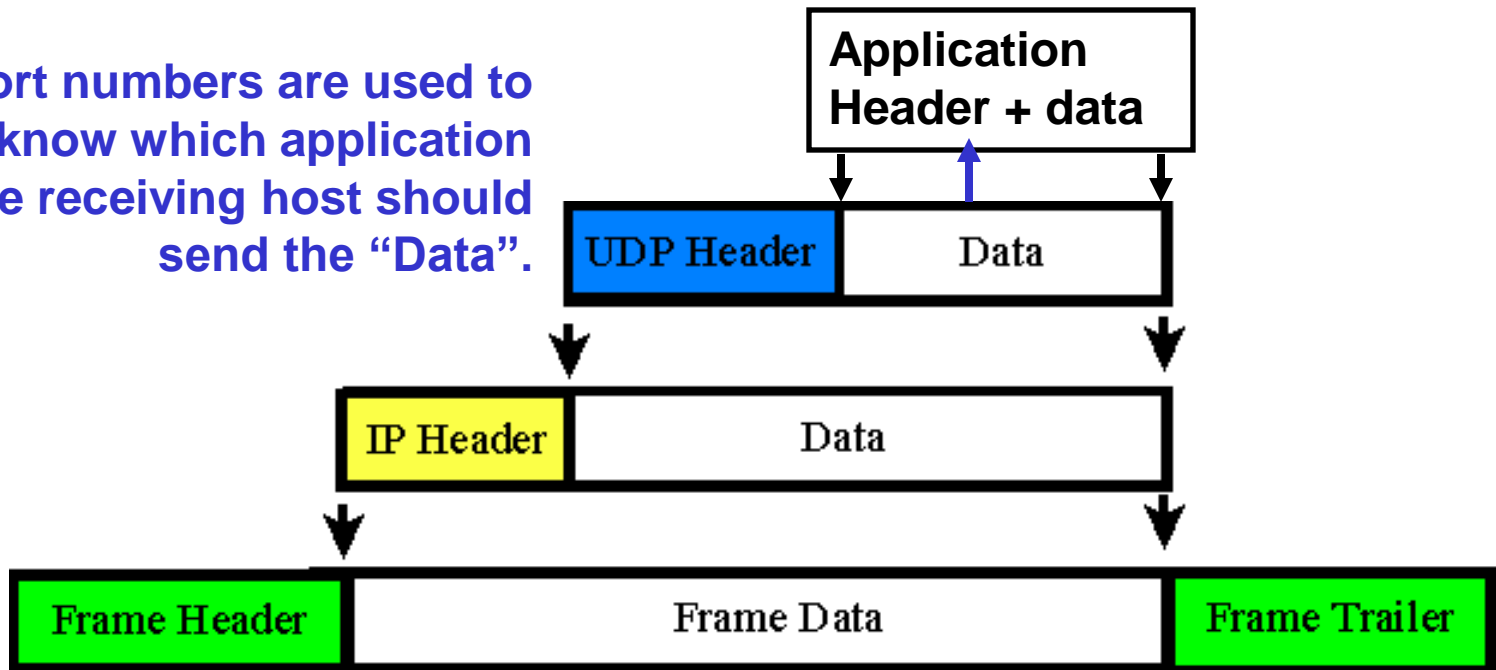
0				15				16				31			
16-bit Source Port Number								16-bit Destination Port Number							
32-bit Sequence Number															
32 bit Acknowledgement Number															
4-bit Header Length		6-bit (Reserved)		U	A	P	R	S	F	16-bit Window Size					
				R	C	S	H	T	I						
				G	K	H	T	N	N						
16-bit TCP Checksum								16-bit Urgent Pointer							
Options (if any)															
Data (if any)															

Port Numbers

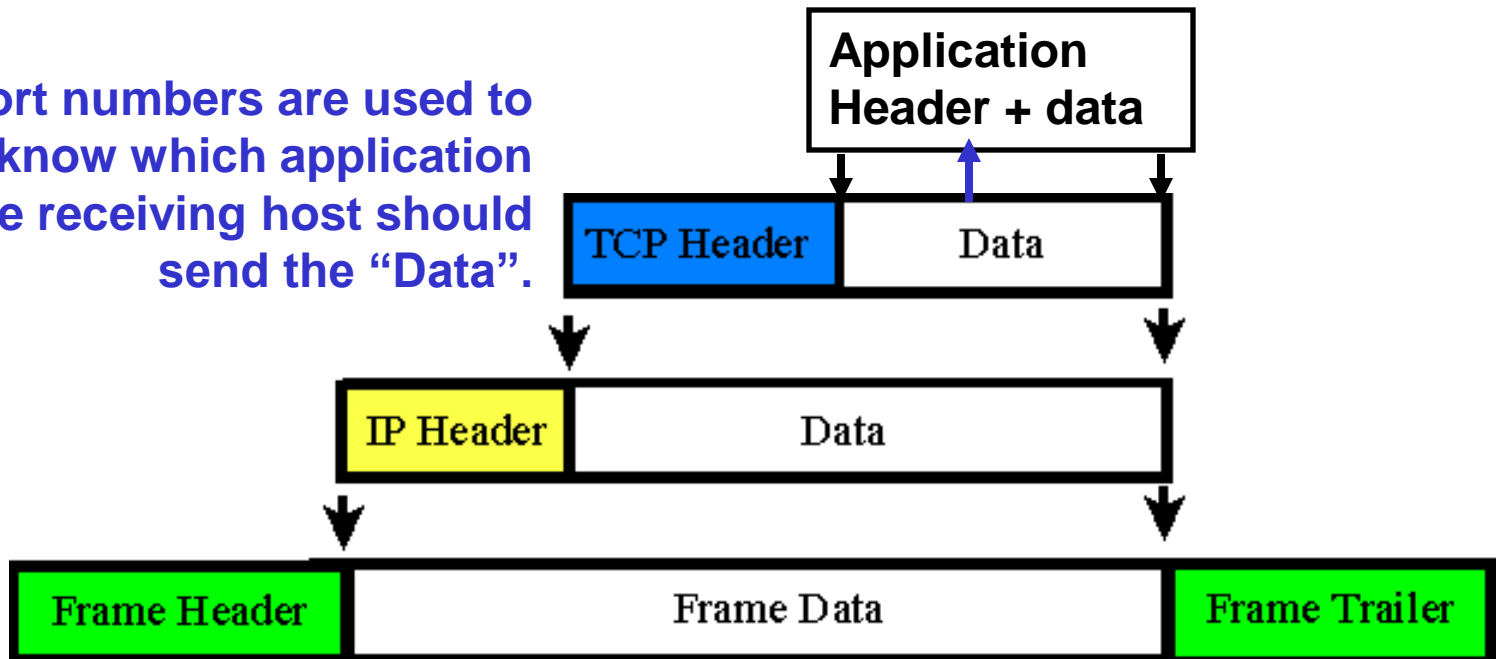


- Both **TCP** and **UDP** use ports (or sockets) numbers to pass information to the upper layers.

Port numbers are used to know which application the receiving host should send the “Data”.

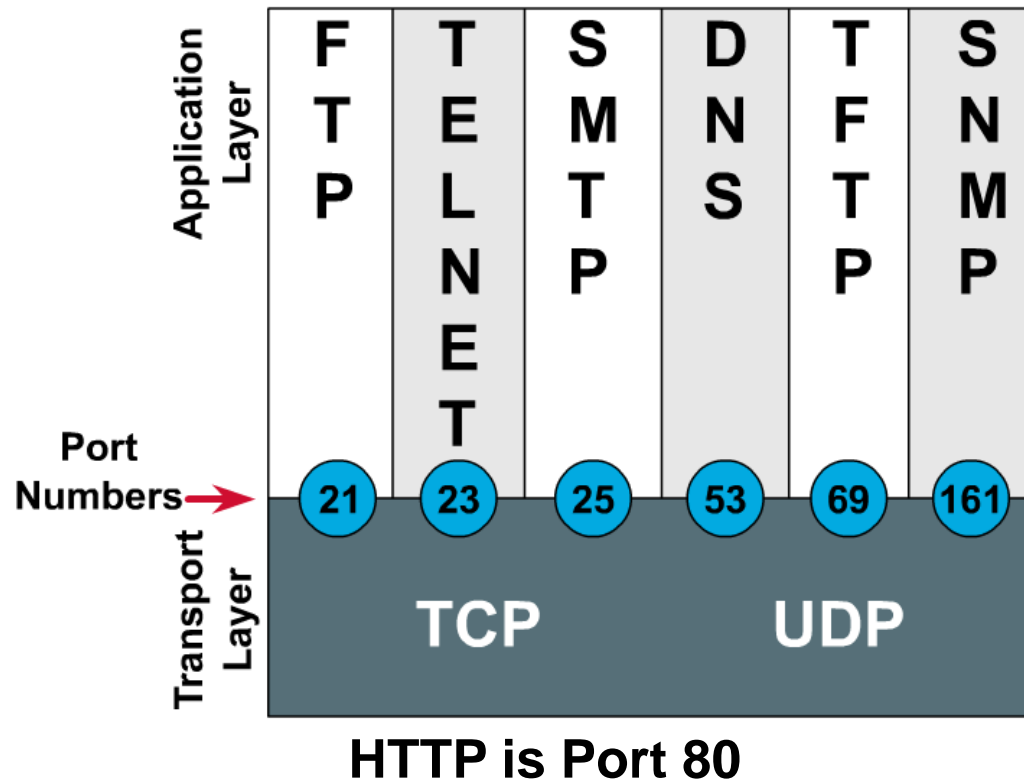
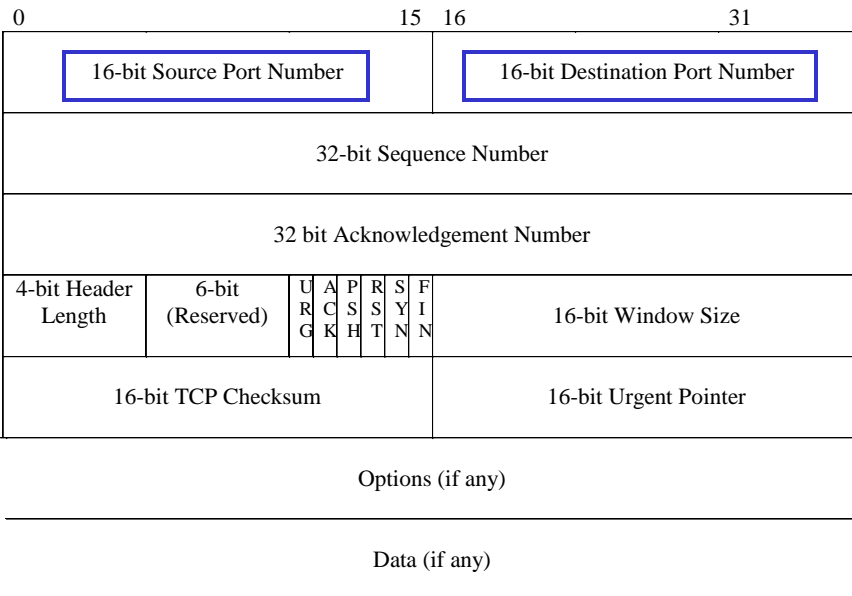


Port numbers are used to know which application the receiving host should send the “Data”.



Port Numbers

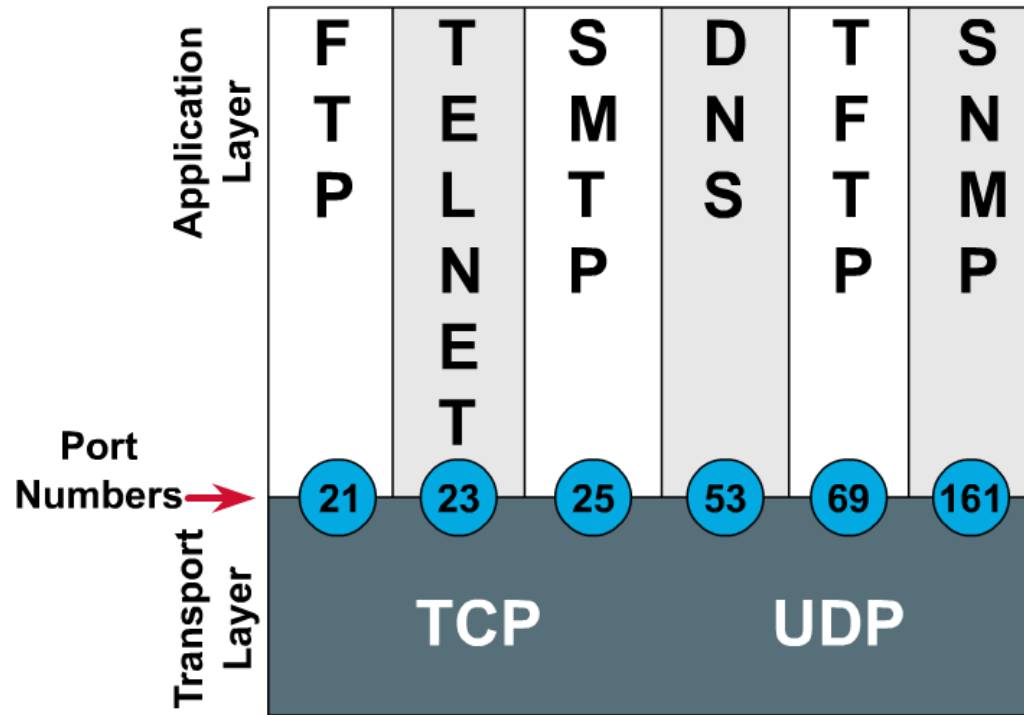
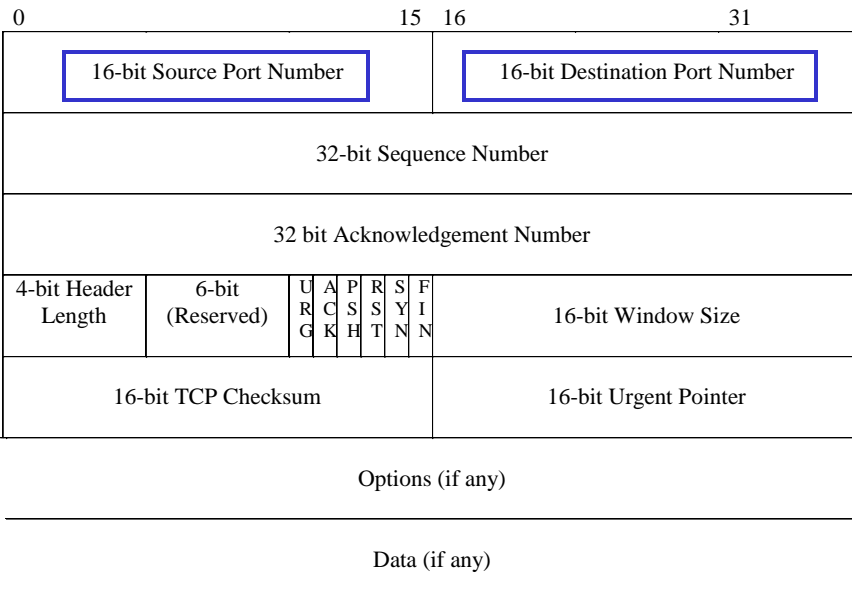
TCP Header



- Application software developers have agreed to use the **well-known port numbers** that are defined in RFC 1700.
- For example, any conversation bound for an **FTP** application uses the standard port number **21**.

Port Numbers

TCP Header

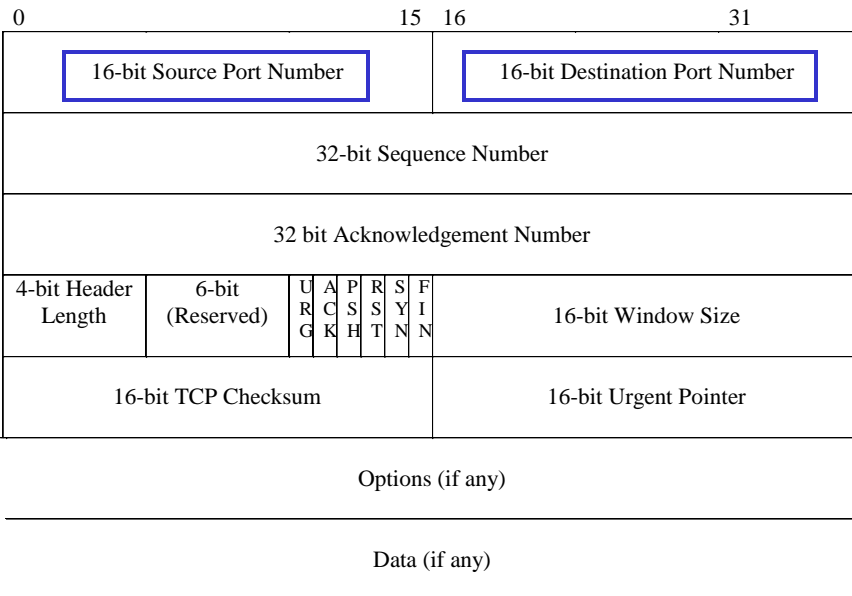


HTTP is Port 80

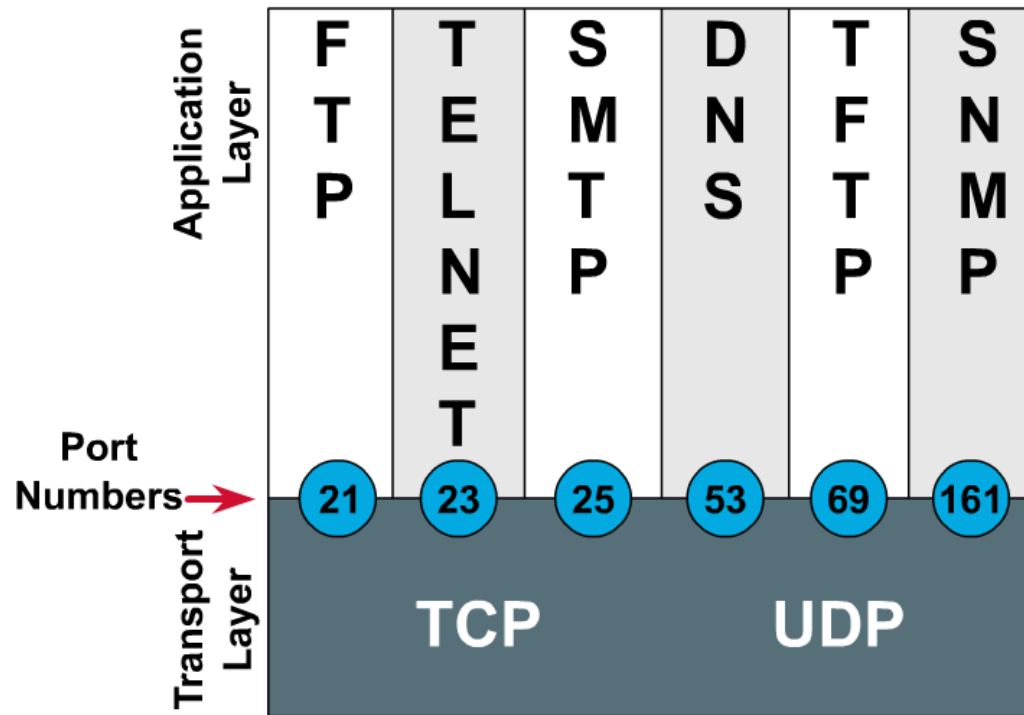
- Conversations that do not involve an application with a well-known port number are, instead, assigned port numbers that are randomly selected from within a specific range.
- These port numbers are used as source and destination addresses in the TCP segment.

Port Numbers

TCP Header



HTTP is Port 80



- Some ports are reserved in both TCP and UDP, although applications might not be written to support them.
- The range for assigned ports managed by the IANA is 0-1023.:
<http://www.iana.org/assignments/port-numbers>
 - The **Well Known Ports** are those from **0 through 1023**. (This is updated information as of 11-13-2002. Before then, 0 – 255 were considered well known ports.)
 - The **Registered Ports** are those from **1024 through 49151**
 - The **Dynamic and/or Private Ports** are those from **49152 through 65535**

Telnet

```
amit@orpheus:~  
[amit@orpheus amit]$ telnet 192.168.100.222  
Trying 192.168.100.222...  
Connected to 192.168.100.222.  
Escape character is '^]'.  
  
*****  
*                               Welcome to ARC Linux                               *  
*****  
  
BusyBox v1.00-rc3 (2004.09.22-12:13+0000) Built-in shell (ash)  
Enter 'help' for a list of built-in commands.  
  
/ # ls  
bin    dev    etc    nfs    proc   root   sbin   tmp    var  
/ # ls bin/  
[      chvt      egrep    install  ping     tee      vi  
addgroup clear     env      kill     ps       test     wc  
adduser  cmp      expr     killall  pwd      time     wget  
ash      cp       false    ln        reset    top      who  
awk      cut      fgrep    logger    rm       touch    whoami  
basename date     find     login     rmdir    tr       xargs  
boa      dd       free     ls        sed      true     yes  
bunzip2  delgroup grep     mkdir     sh       tty      zcat  
busybox  deluser  gunzip   mknod     sleep    umount  
bzipcat  df       gzip     more      sort     uname  
cat      dirname head     mount     su       uniq  
chgrp    dmesg   hexdump  mv        sync     unzip  
chmod    du      hostname netstat   tail     uptime  
chown    echo    id       passwd    tar      usleep  
/ # █
```

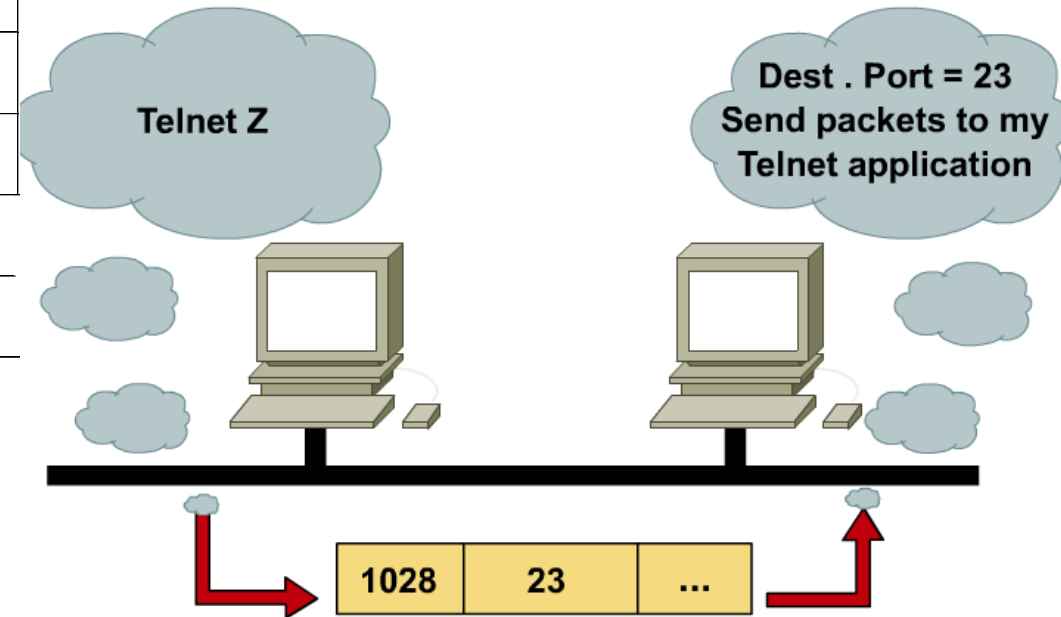

TCP Header

31

16-bit Source Port Number				16-bit Destination Port Number			
32-bit Sequence Number							
32 bit Acknowledgement Number							
4-bit Header Length	6-bit (Reserved)	U R G	A C K	P R S T	S S Y N	F I N	16-bit Window Size
16-bit TCP Checksum				16-bit Urgent Pointer			
Options (if any)							
Data (if any)							

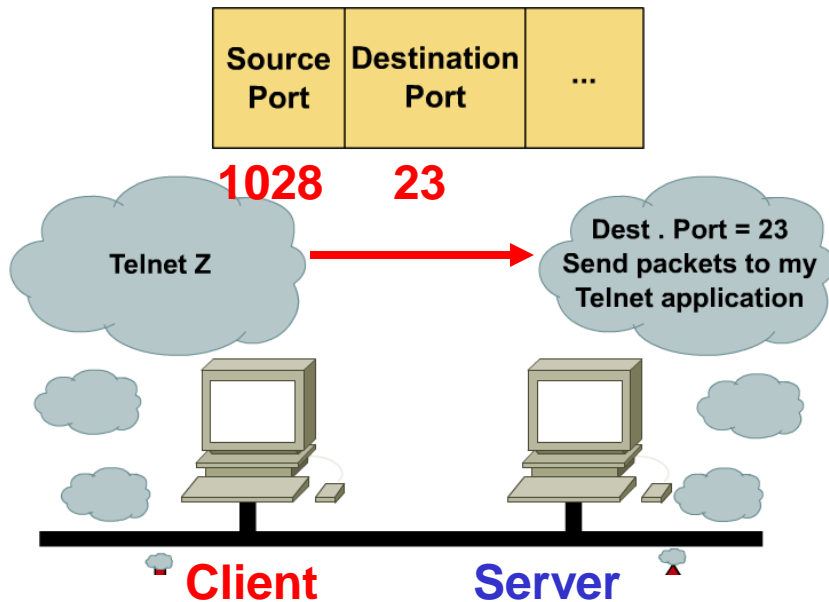
TCP/UDP Port Numbers

Source Port	Destination Port	...
-------------	------------------	-----

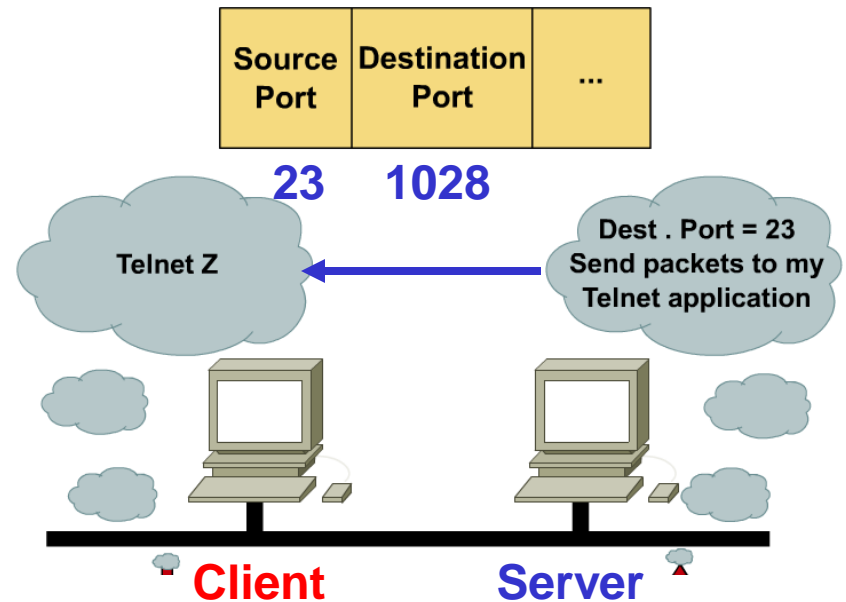


- End systems use port numbers to select the proper application.
- Originating source port numbers, usually some numbers larger than 1023, are dynamically assigned by the source host.

TCP/UDP Port Numbers



TCP/UDP Port Numbers



Notice the difference in how source and destination port numbers are used with clients and servers:

Client (initiating Telnet service):

- Destination Port = 23 (telnet)
- Source Port = 1028 (dynamically assigned)

Server (responding to Telnet service):

- Destination Port = 1028 (source port of client)
- Source Port = 23 (telnet)

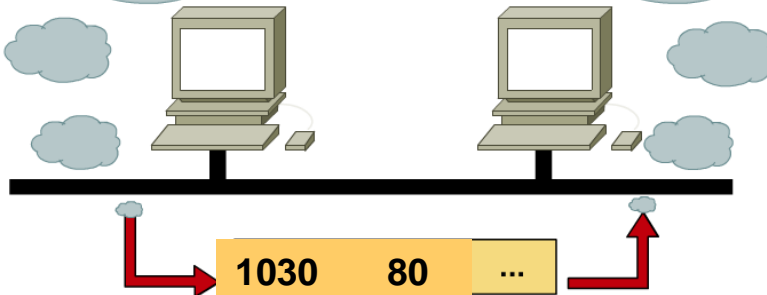
TCP/UDP Port Numbers

Source Port	Destination Port	...
-------------	------------------	-----

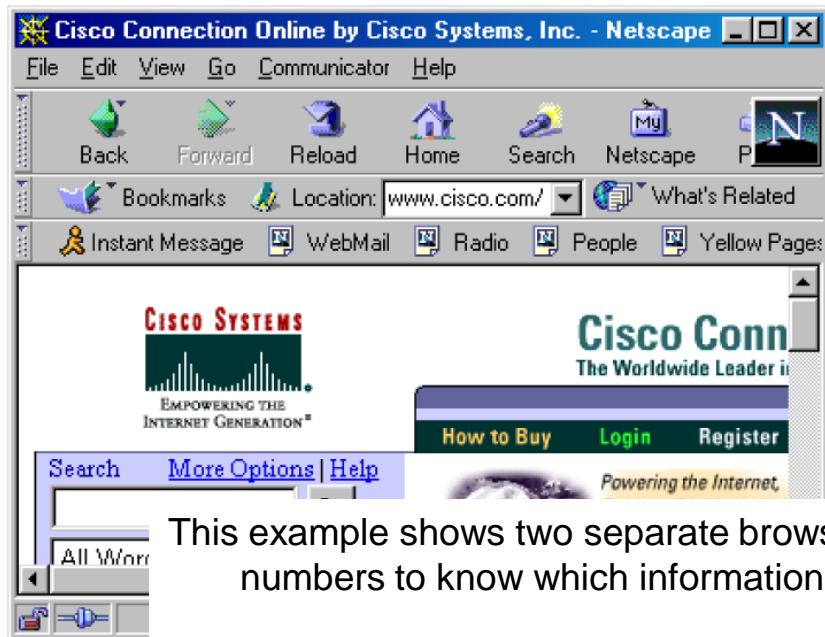
Second http session from the between the same client and server. Same destination port, but different source port to uniquely identify this web session.

http to
www.cisco.com

Dest. Port = 80 Send
packets to web
server application

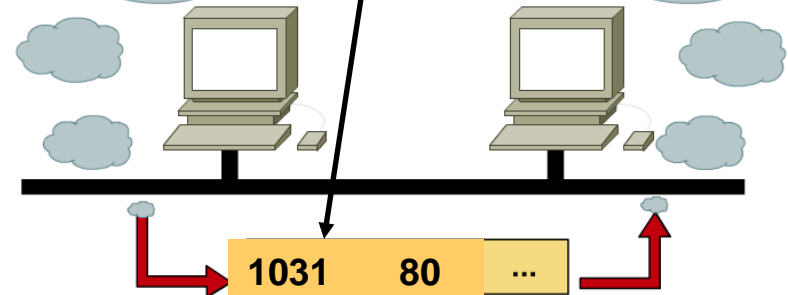


Netscape Navigator



http to
www.cisco.com

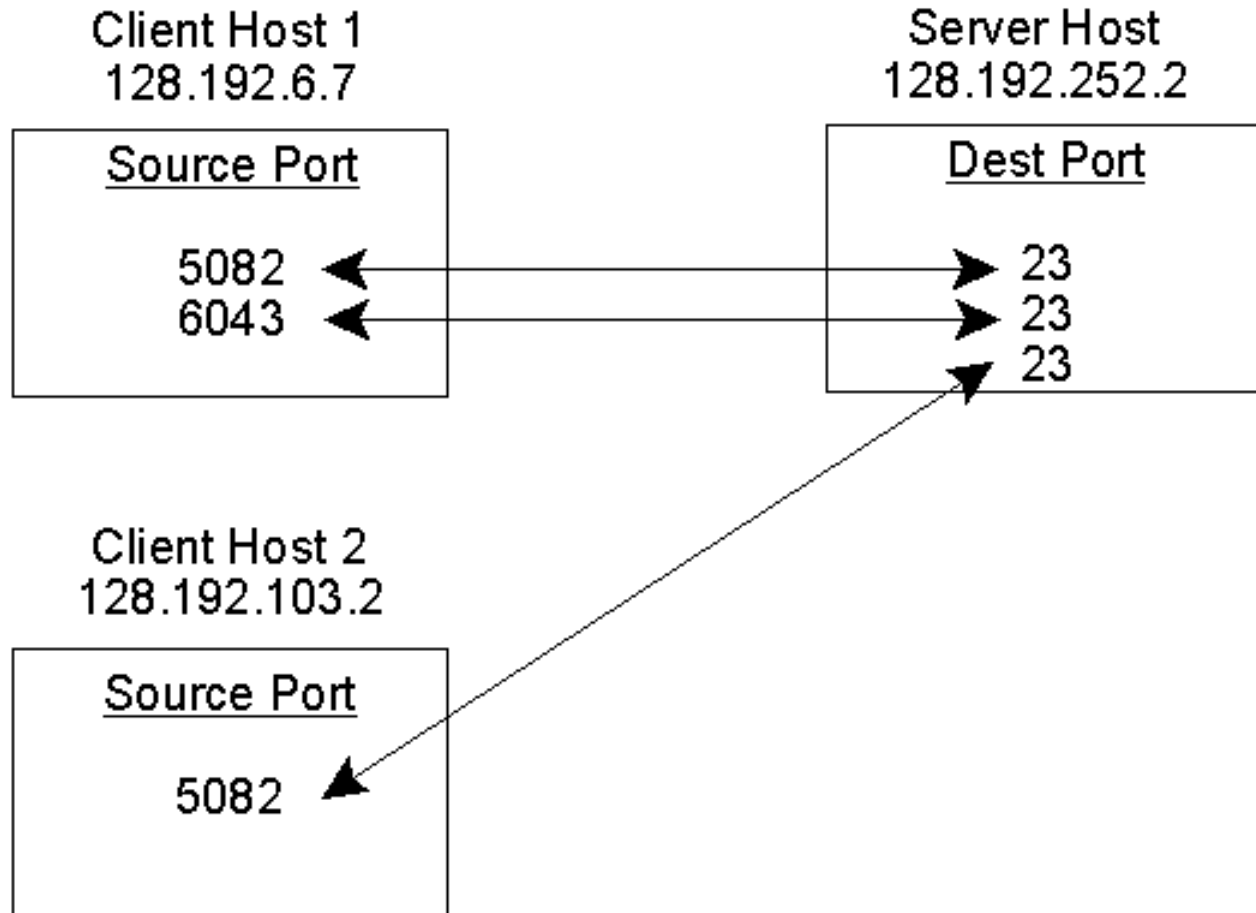
Dest. Port = 80 Send
packets to web
server application



Netscape Navigator

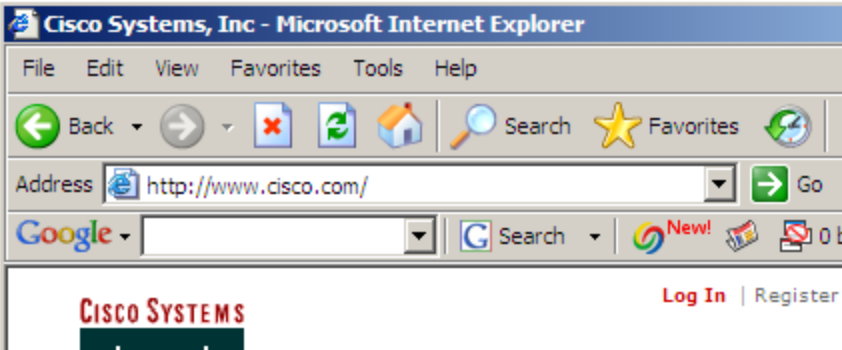
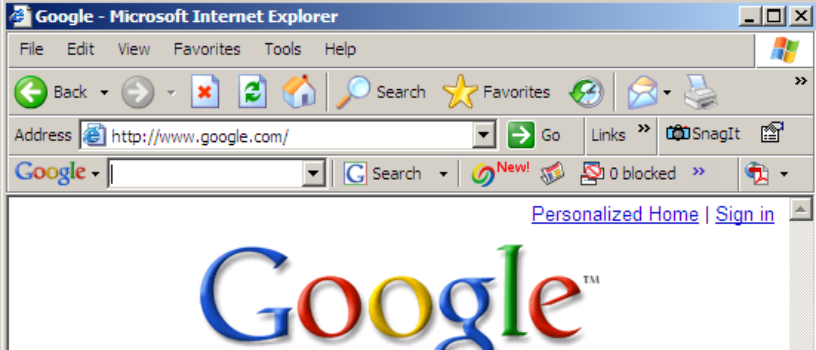


This example shows two separate browser windows to the same URL. TCP/IP uses source port numbers to know which information goes to which window.



What makes each connection unique?

- Connection defined by the pair of numbers:
 - **Source IP address, Source port**
 - **Destination IP address, Destination port**
- Different connections can use the same destination port on server host as long as the source ports or source IPs are different.



TCP
or
UDP

C:\WINDOWS\system32\cmd.exe

C:\>netstat -n

Active Connections

Source IP

Destination IP

Connection State

Source Port

Destination Port

Proto	Local Address	Foreign Address	State
TCP	172.17.150.112:1033	172.16.1.44:524	ESTABLISHED
TCP	172.17.150.112:1034	172.16.1.44:524	ESTABLISHED
TCP	172.17.150.112:1042	205.188.9.73:5190	ESTABLISHED
TCP	172.17.150.112:1050	64.12.165.95:5190	ESTABLISHED
TCP	172.17.150.112:1069	207.62.185.140:143	ESTABLISHED
TCP	172.17.150.112:1332	198.133.219.25:80	TIME_WAIT
TCP	172.17.150.112:1333	198.133.219.25:80	ESTABLISHED
TCP	172.17.150.112:1334	198.133.219.25:80	ESTABLISHED
TCP	172.17.150.112:1335	64.154.80.254:80	ESTABLISHED
TCP	172.17.150.112:1336	66.102.7.99:80	ESTABLISHED

C:\>

www.google.com

www.cisco.com

netstat -n

- **Note:** In actuality, when you open up a single html page, there are usually several TCP sessions created, not just one.
- Example of multiple TCP connections for a single http session.

Topics

Layer 3 Concepts

- TCP/IP and the Internet Layer
- Diagram the IP datagram
- Internet Control Message Protocol (ICMP)

TCP/IP protocol stack and the transport layer

- TCP and UDP segment format
- TCP and UDP port numbers
- TCP three-way handshake/open connection
- TCP simple acknowledgment and windowing

TCP Header

31

16-bit Source Port Number				16-bit Destination Port Number							
<div>32-bit Sequence Number</div>											
<div>32 bit Acknowledgement Number</div>											
4-bit Header Length		6-bit (Reserved)		U R G	A C K	P S H	R S T	S Y N	F I N	16-bit Window Size	
16-bit TCP Checksum						16-bit Urgent Pointer					
Options (if any)											
Data (if any)											

TCP Three-Way Handshake/ Open Connection



Send SYN
(seq = x)

Receive SYN
(seq = y,
ACK = x + 1)

Send ACK
(ack = y+1)



Receive SYN
(seq = x)

Send SYN
(seq = y,
ACK = x + 1)

Receive ACK
(ack = y+1)

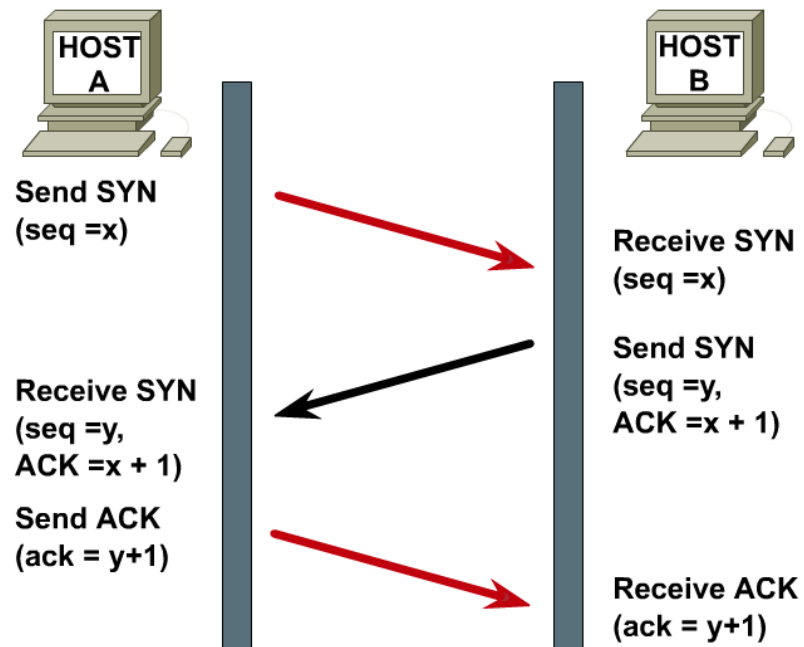
- For a connection to be established, the two end stations must synchronize on each other's TCP initial sequence numbers (ISNs).
- Sequence numbers are used to track the order of packets and to ensure that no packets are lost in transmission.
- The initial sequence number is the starting number used when a TCP connection is established.
- Exchanging beginning sequence numbers during the connection sequence ensures that lost data can be recovered.

TCP Header

31

16-bit Source Port Number				16-bit Destination Port Number				
<div>32-bit Sequence Number</div>								
<div>32 bit Acknowledgement Number</div>								
4-bit Header Length	6-bit (Reserved)	U	A	P	R	S	F	16-bit Window Size
		R	C	S	S	Y	I	
		G	K	H	T	N	N	
16-bit TCP Checksum				16-bit Urgent Pointer				
Options (if any)								
Data (if any)								

TCP Three-Way Handshake/ Open Connection



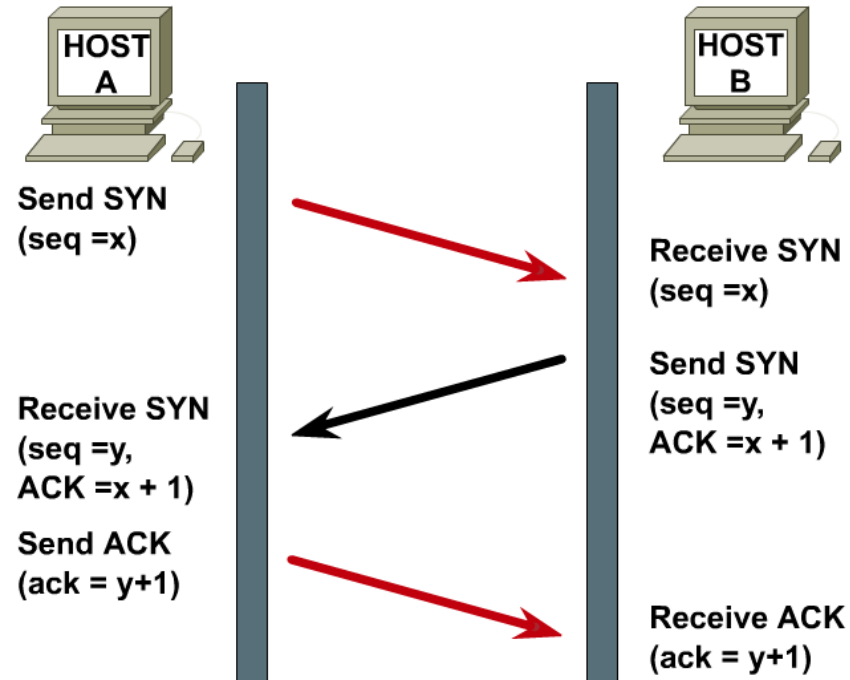
- **Synchronization** is accomplished by exchanging segments carrying the ISNs and a control bit called *SYN*, which stands for *synchronize*. (Segments carrying the SYN bit are also called SYNs.)
- Successful connection requires a suitable mechanism for choosing an initial sequence and a slightly involved handshake to exchange the ISNs.
- Synchronization requires that each side send its own ISN and receive a confirmation and ISN from the other side of the connection.
- Each side must receive the other side's ISN and send a confirming acknowledgment (ACK) in a specific order.

TCP Header

31

16-bit Source Port Number				16-bit Destination Port Number						
<div>32-bit Sequence Number</div>										
<div>32 bit Acknowledgement Number</div>										
4-bit Header Length		6-bit (Reserved)		U	A	P	R	S	F	16-bit Window Size
				R	C	S	S	Y	I	
				G	K	H	T	N	N	
16-bit TCP Checksum					16-bit Urgent Pointer					
Options (if any)										
Data (if any)										

TCP Three-Way Handshake/ Open Connection



- Because the second and third steps can be combined in a single message, the exchange is called a three-way handshake/open connection.
- A three-way handshake is necessary because TCPs may use different mechanisms for picking the ISN.
- The receiver of the first SYN has no way of knowing if the segment was an old delayed one unless it remembers the last sequence number used on the connection, which is not always possible, and so it must ask the sender to verify this SYN

TCP Header

31

16-bit Source Port Number				16-bit Destination Port Number						
<div>32-bit Sequence Number</div>										
<div>32 bit Acknowledgement Number</div>										
4-bit Header Length		6-bit (Reserved)		U	A	P	R	S	F	16-bit Window Size
				R	C	S	S	Y	I	
				G	K	H	T	N	N	
16-bit TCP Checksum					16-bit Urgent Pointer					
Options (if any)										
Data (if any)										

TCP Three-Way Handshake/ Open Connection



Send SYN
(seq =x)



Receive SYN
(seq =x)

Send SYN
(seq =y,
ACK =x + 1)

Receive SYN
(seq =y,
ACK =x + 1)

Send ACK
(ack = y+1)

Receive ACK
(ack = y+1)

- At this point, either side can begin communicating, and either side can break the communication because TCP is a peer-to-peer (balanced) communication method.

Client: Seq = 4264974716

No. ↓	↑	Source	Destination	Protocol	Info
66	4	172.17.150.112	207.62.187.7	TCP	1061 > 80 [SYN] Seq=4264974716 Ack=0 win=65535 Len=0 MSS=1260
67	4	207.62.187.7	172.17.150.112	TCP	80 > 1061 [SYN, ACK] Seq=1158257438 Ack=4264974717 win=5840 Len=0 MSS=1460
68	4	172.17.150.112	207.62.187.7	TCP	1061 > 80 [ACK] Seq=4264974717 Ack=1158257439 win=65535 Len=0
69	4	172.17.150.112	207.62.187.7	HTTP	GET /~rgraziani/ HTTP/1.1
70	4	207.62.187.7	172.17.150.112	TCP	80 > 1061 [ACK] Seq=1158257439 Ack=4264975052 win=6432 Len=0
Frame 66 (62 bytes on wire, 62 bytes captured)					
Ethernet II, Src: Wistron_d4:4c:f3 (00:0a:e4:d4:4c:f3), Dst: 172.17.128.1 (00:03:e3:7e:1d:c0)					
Internet Protocol, Src: 172.17.150.112 (172.17.150.112), Dst: 207.62.187.7 (207.62.187.7)					
Version: 4					
Header length: 20 bytes					
Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00)					
Total Length: 48					
Identification: 0x0372 (882)					
Flags: 0x04 (Don't Fragment)					
Fragment offset: 0					
Time to live: 128					
Protocol: TCP (0x06)					
Header checksum: 0x2a8e [correct]					
Source: 172.17.150.112 (172.17.150.112)					
Destination: 207.62.187.7 (207.62.187.7)					
Transmission Control Protocol, Src Port: 1061 (1061), Dst Port: http (80), Seq: 4264974716, Ack: 0, Len: 0					
Source port: 1061 (1061)					
Destination port: http (80)					
Sequence number: 4264974716					
Header length: 28 bytes					
Flags: 0x0002 (SYN)					
window size: 65535					
checksum: 0x5af7 [correct]					
Options: (8 bytes)					

Server: ACK = 4264974717

Seq = 1158257438

Cabrillo College

No. ↓	↑	Source	Destination	Protocol	Info
66	4	172.17.150.112	207.62.187.7	TCP	1061 > 80 [SYN] Seq=4264974716 Ack=0 win=65535 Len=0 MSS=1260
67	4	207.62.187.7	172.17.150.112	TCP	80 > 1061 [SYN, ACK] Seq=1158257438 Ack=4264974717 win=5840 Len=0 MSS=1460
68	4	172.17.150.112	207.62.187.7	TCP	1061 > 80 [ACK] Seq=4264974717 Ack=1158257439 win=65535 Len=0
69	4	172.17.150.112	207.62.187.7	HTTP	GET /~rgraziani/ HTTP/1.1
70	4	207.62.187.7	172.17.150.112	TCP	80 > 1061 [ACK] Seq=1158257439 Ack=4264975052 win=6432 Len=0

⊞	Frame 67 (62 bytes on wire, 62 bytes captured)
⊞	Ethernet II, Src: 172.17.128.1 (00:03:e3:7e:1d:c0), Dst: Wistron_d4:4c:f3 (00:0a:e4:d4:4c:f3)
⊞	Internet Protocol, Src: 207.62.187.7 (207.62.187.7), Dst: 172.17.150.112 (172.17.150.112) Version: 4 Header length: 20 bytes
⊞	Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00) Total Length: 48 Identification: 0x0000 (0)
⊞	Flags: 0x04 (Don't Fragment) Fragment offset: 0 Time to live: 62 Protocol: TCP (0x06)
⊞	Header checksum: 0x7000 [correct] Source: 207.62.187.7 (207.62.187.7) Destination: 172.17.150.112 (172.17.150.112)
⊞	Transmission Control Protocol, Src Port: http (80), Dst Port: 1061 (1061), Seq: 1158257438, Ack: 4264974717, Len: 0 Source port: http (80) Destination port: 1061 (1061) Sequence number: 1158257438 Acknowledgement number: 4264974717 Header length: 28 bytes
⊞	Flags: 0x0012 (SYN, ACK) window size: 5840 checksum: 0x6326 [correct]
⊞	Options: (8 bytes)
⊞	[SEQ/ACK analysis]

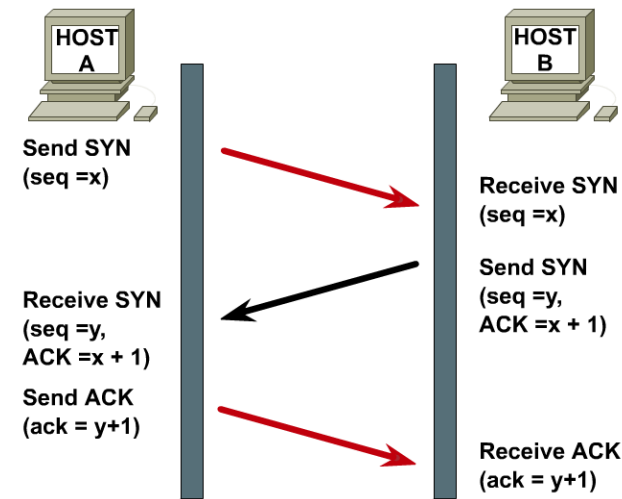
Client: ACK = 1158257439

No.	Source	Destination	Protocol	Info
66	4 172.17.150.112	207.62.187.7	TCP	1061 > 80 [SYN] Seq=4264974716 Ack=0 win=65535 Len=0 MSS=1260
67	4 207.62.187.7	172.17.150.112	TCP	80 > 1061 [SYN, ACK] Seq=1158257438 Ack=4264974717 win=5840 Len=0 MSS=1460
68	4 172.17.150.112	207.62.187.7	TCP	1061 > 80 [ACK] Seq=4264974717 Ack=1158257439 win=65535 Len=0
69	4 172.17.150.112	207.62.187.7	HTTP	GET /~rgraziani/ HTTP/1.1
70	4 207.62.187.7	172.17.150.112	TCP	80 > 1061 [ACK] Seq=1158257439 Ack=4264975052 win=6432 Len=0

[-] Frame 68 (54 bytes on wire, 54 bytes captured)
[-] Ethernet II, Src: wistron_d4:4c:f3 (00:0a:e4:d4:4c:f3), Dst: 172.17.128.1 (00:03:e3:7e:1d:c0)
[-] Internet Protocol, Src: 172.17.150.112 (172.17.150.112), Dst: 207.62.187.7 (207.62.187.7)
[-] Version: 4
[-] Header length: 20 bytes
[-] Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00)
[-] Total Length: 40
[-] Identification: 0x0374 (884)
[-] Flags: 0x04 (Don't Fragment)
[-] Fragment offset: 0
[-] Time to live: 128
[-] Protocol: TCP (0x06)
[-] Header checksum: 0x2a94 [correct]
[-] Source: 172.17.150.112 (172.17.150.112)
[-] Destination: 207.62.187.7 (207.62.187.7)
[-] Transmission Control Protocol, Src Port: 1061 (1061), Dst Port: http (80), Seq: 4264974717, Ack: 1158257439, Len: 0
[-] Source port: 1061 (1061)
[-] Destination port: http (80)
[-] Sequence number: 4264974717
[-] Acknowledgement number: 1158257439
[-] Header length: 20 bytes
[-] Flags: 0x0010 (ACK)
[-] window size: 65535
[-] Checksum: 0xa6ba [correct]
[-] [SEQ/ACK analysis]

Another example

TCP Three-Way Handshake/ Open Connection



- Only part of the TCP headers are displayed.
- Notice that the Maximum segment size and the negotiated Window size are also sent.

Packet 1: source: 130.57.20.10 dest.:130.57.20.1

TCP: ----- TCP header -----

TCP: Source port = 1026

TCP: Destination port = 524

TCP: Initial sequence number = 12952

TCP: Next expected Seq number= 12953

TCP:1. = SYN

TCP: Window = 8192

TCP: Checksum = 1303 (correct)

TCP: Maximum segment size = 1460 (TCP Option)

Packet 2: source: 130.57.20.1 dest: 130.57.20.10

TCP: ----- TCP header -----

TCP: Source port = 524

TCP: Destination port = 1026

TCP: Initial sequence number = 2744080

TCP: Next expected Seq number= 2744081

TCP: Acknowledgment number = 12953

TCP:1. = SYN

TCP: Window = 32768

TCP: Checksum = D3B7 (correct)

TCP: Maximum segment size = 1460 (TCP Option)

Packet 3: source: 130.57.20.10 dest: 130.57.20.1

TCP: ----- TCP header -----

TCP: Source port = 1026

TCP: Destination port = 524

TCP: Sequence number = 12953

TCP: Next expected Seq number= 12953

TCP: Acknowledgment number = 2744081

TCP: ...1 = Acknowledgment

TCP: Window = 8760

TCP: Checksum = 493D (correct)

TCP: No TCP options

*

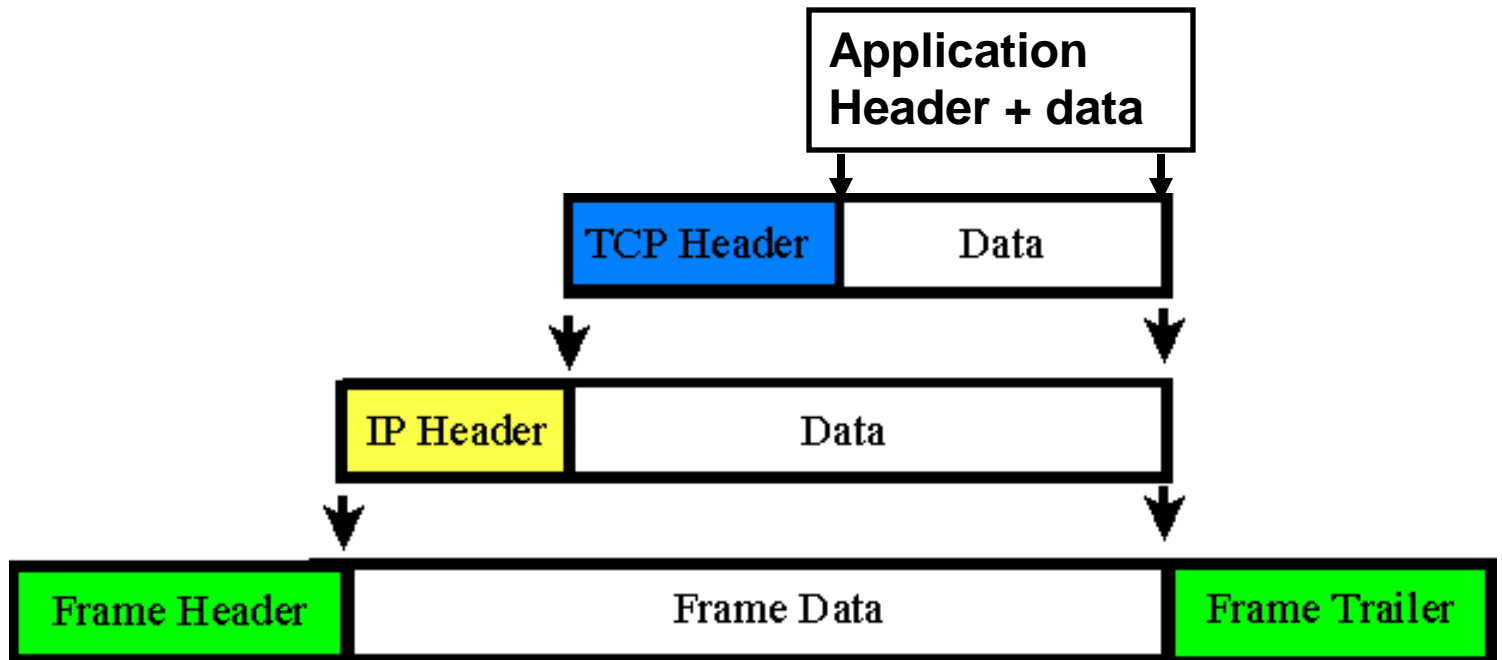
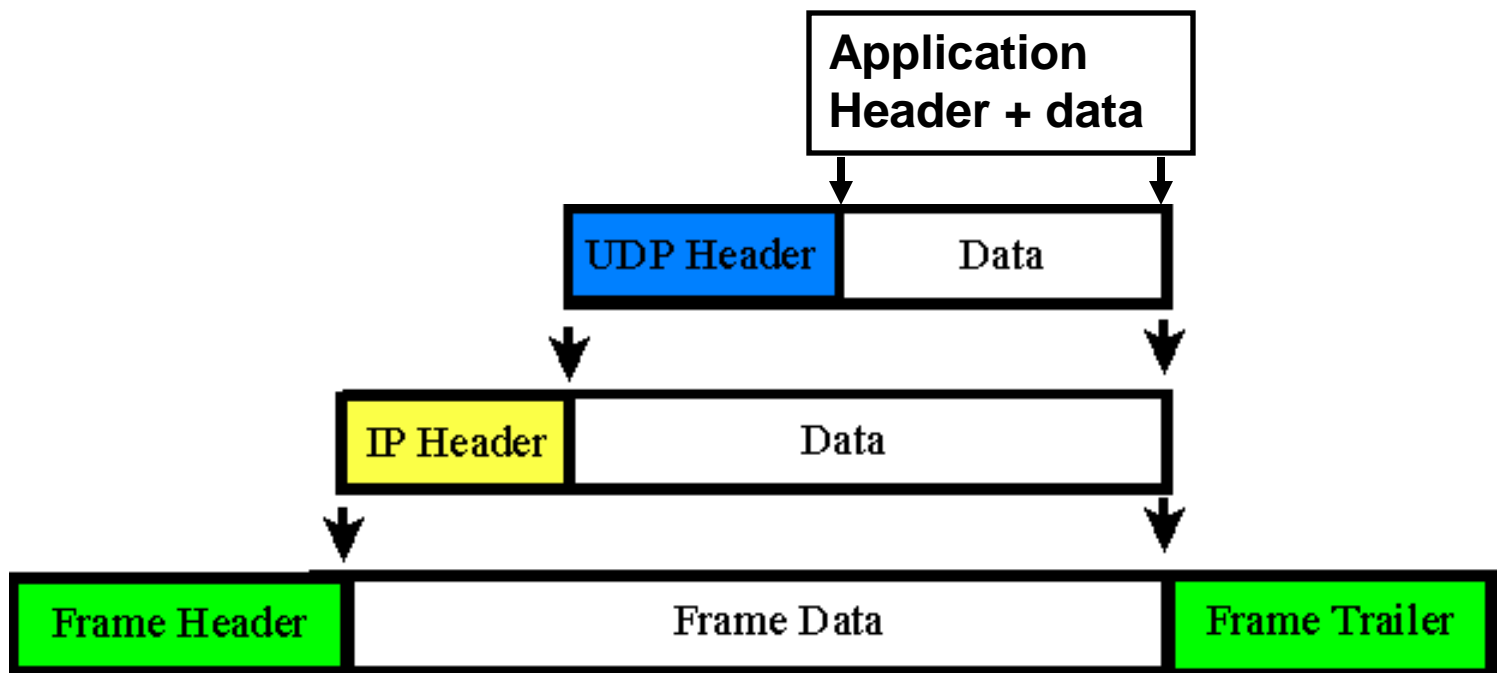
Topics

Layer 3 Concepts

- TCP/IP and the Internet Layer
- Diagram the IP datagram
- Internet Control Message Protocol (ICMP)

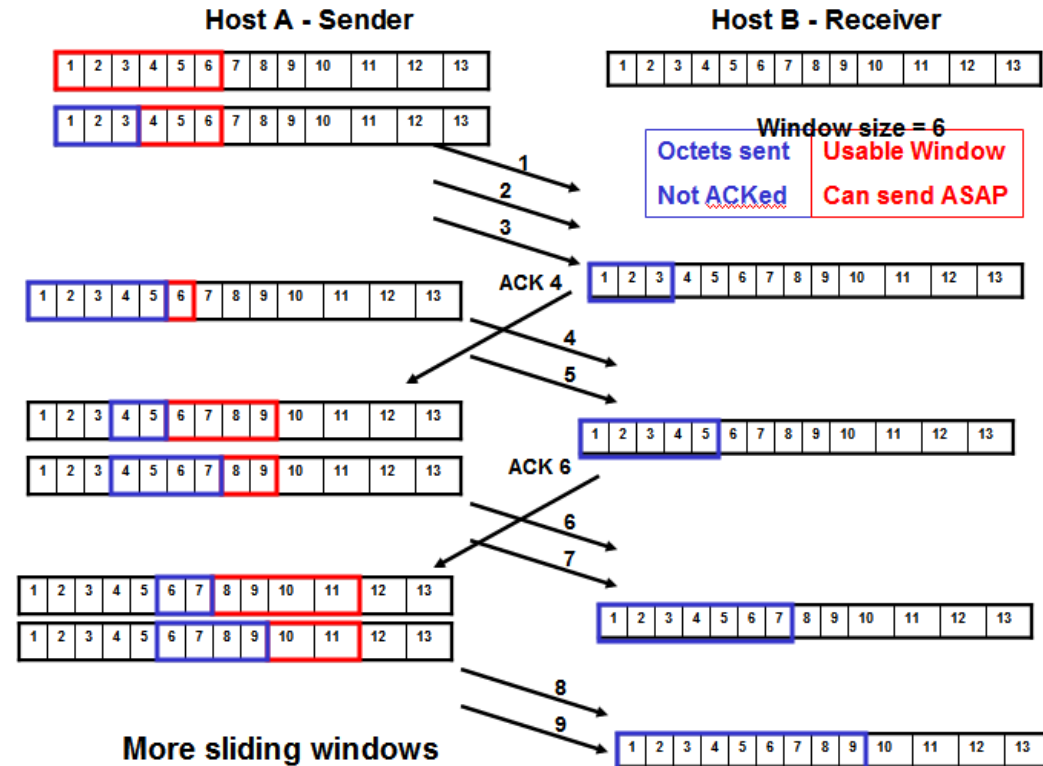
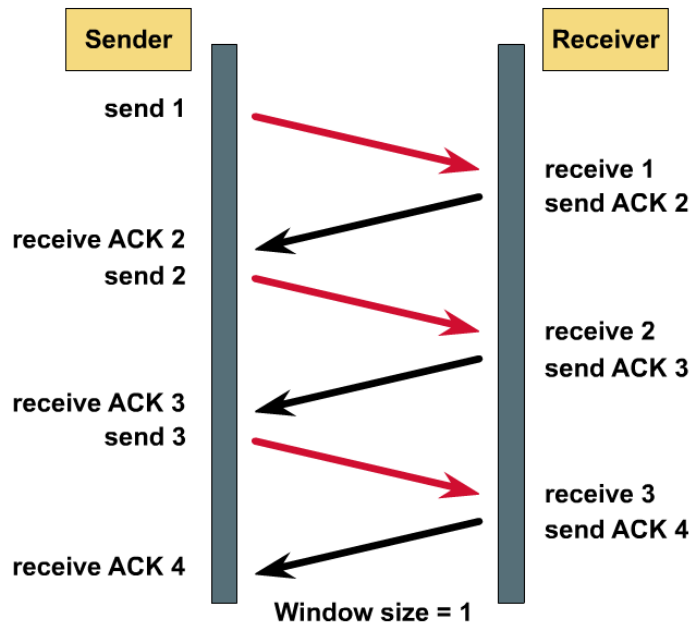
TCP/IP protocol stack and the transport layer

- TCP and UDP segment format
- TCP and UDP port numbers
- TCP three-way handshake/open connection
- TCP simple acknowledgment and windowing



TCP Windows and Sliding Windows

TCP Simple Acknowledgment



Over Simplification

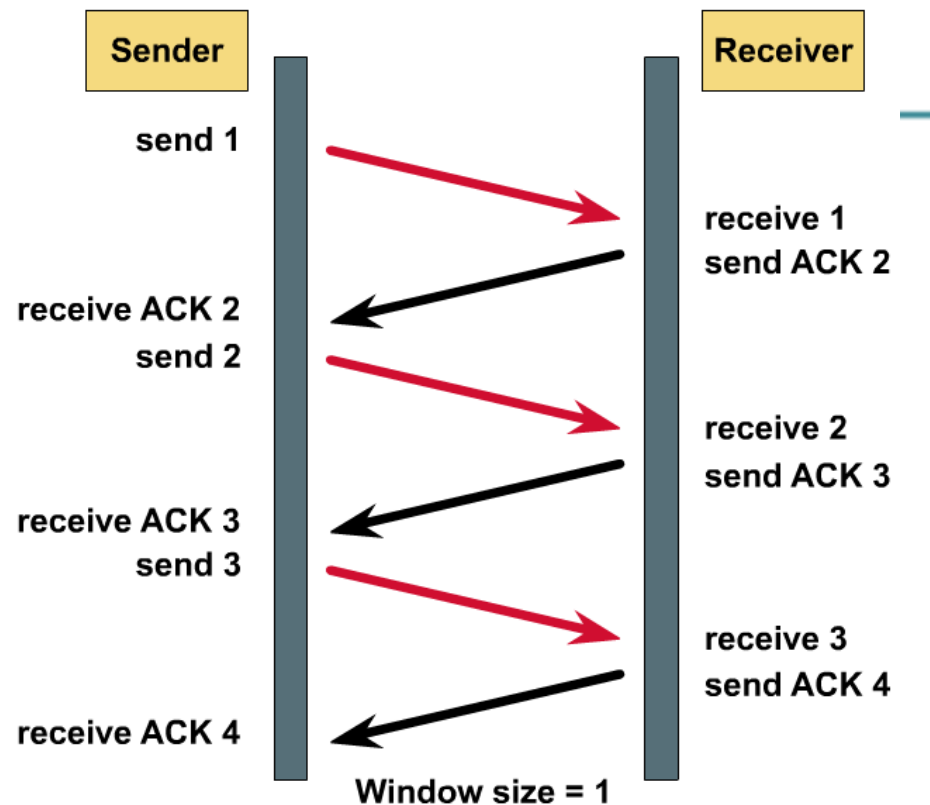
- Note: The following examples of Window Size, Sliding Windows, and Retransmission are very simplistic examples using 1 byte segments. This is meant only to introduce the reader to TCP and is not intended to give realistic examples.

TCP Header

31

16-bit Source Port Number				16-bit Destination Port Number			
<div>32-bit Sequence Number</div>							
<div>32 bit Acknowledgement Number</div>							
4-bit Header Length		6-bit (Reserved)		<div>U A P R S F R C S S Y I G K H T N N</div>		16-bit Window Size	
16-bit TCP Checksum				16-bit Urgent Pointer			
Options (if any)							
Data (if any)							

TCP Simple Acknowledgment



Flow Control and Reliability

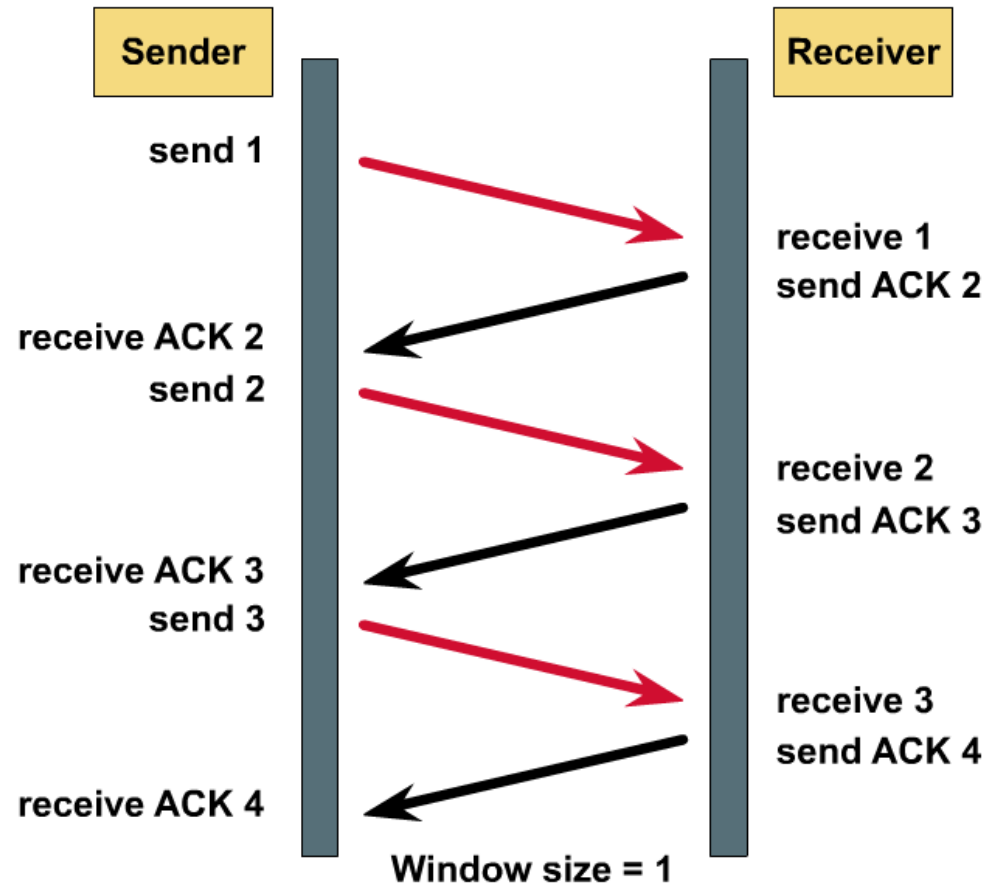
- To govern the flow of data between devices, TCP uses a peer-to-peer flow control mechanism.
- The receiving host's TCP layer reports a window size to the sending host's TCP layer.
- This **window size** specifies the number of bytes, starting with the acknowledgment number, that the receiving host's TCP layer is currently prepared to receive.

TCP Header

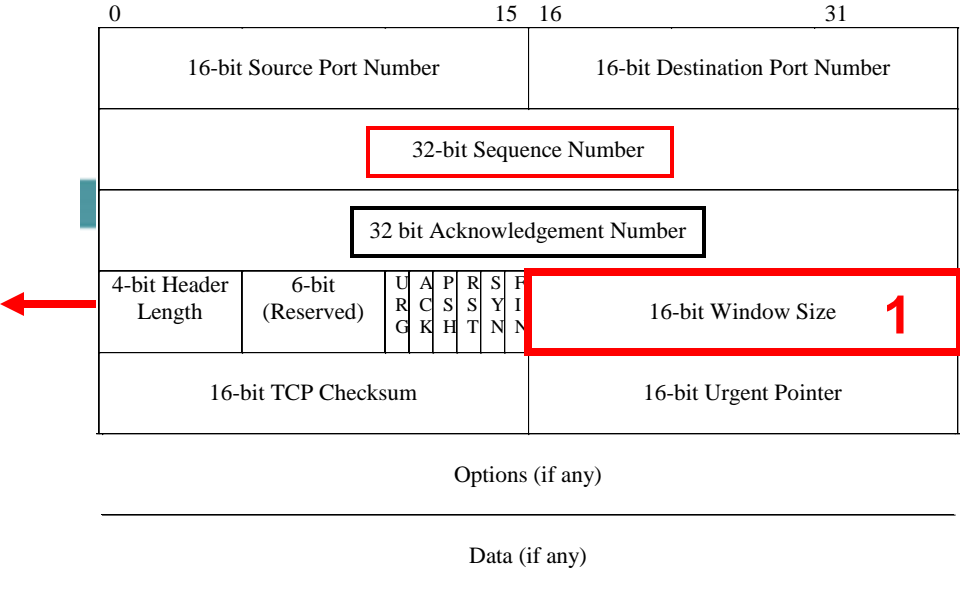
31

16-bit Source Port Number				16-bit Destination Port Number			
<div>32-bit Sequence Number</div>							
<div>32 bit Acknowledgement Number</div>							
4-bit Header Length		6-bit (Reserved)		<div>UAPRSF RCSSH GKHTNN</div>		<div>16-bit Window Size</div>	
16-bit TCP Checksum				16-bit Urgent Pointer			
Options (if any)							
Data (if any)							

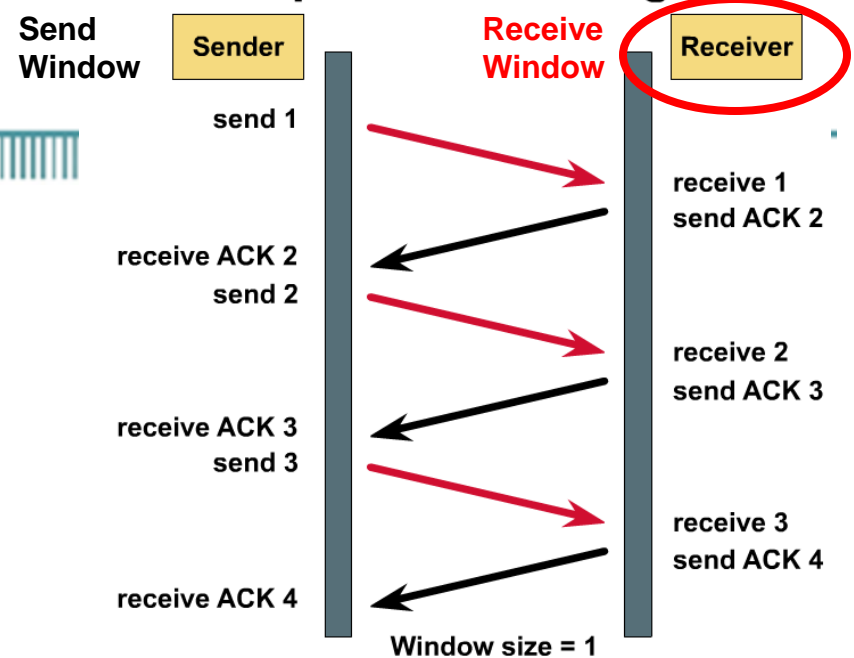
TCP Simple Acknowledgment



- **TCP** -- a connection-oriented, reliable protocol; provides **flow control** by providing sliding windows, and **reliability** by providing sequence numbers and acknowledgments.
- TCP re-sends anything that is not received and supplies a “TCP” **virtual circuit** between end-user applications.
- The advantage of TCP is that it provides **guaranteed delivery** of the segments.



TCP Simple Acknowledgment



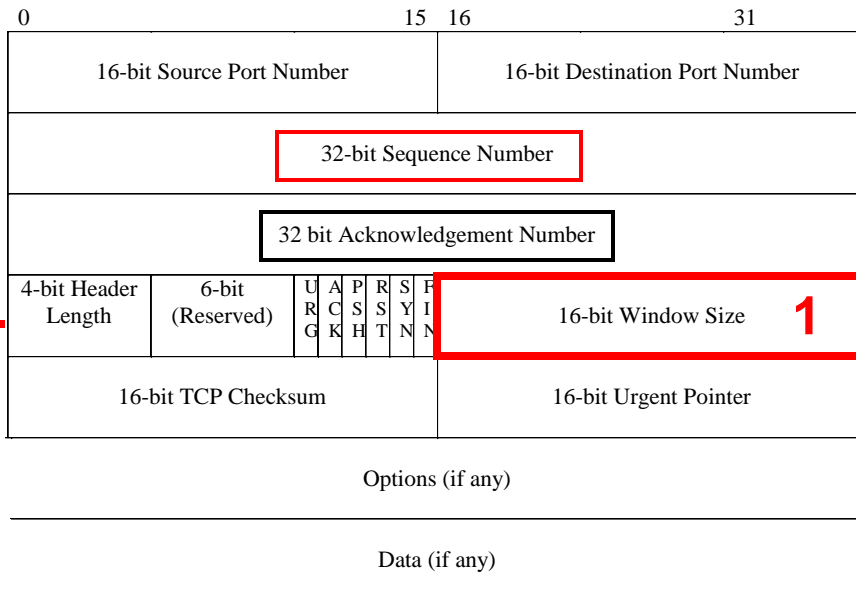
Receive Window

- The TCP Receive Window size is the amount of receive data (in bytes) that can be buffered by this host, at one time on a connection.
- The other (sending) host can send only that amount of data before getting an acknowledgment and window update from this (the receiving) host.

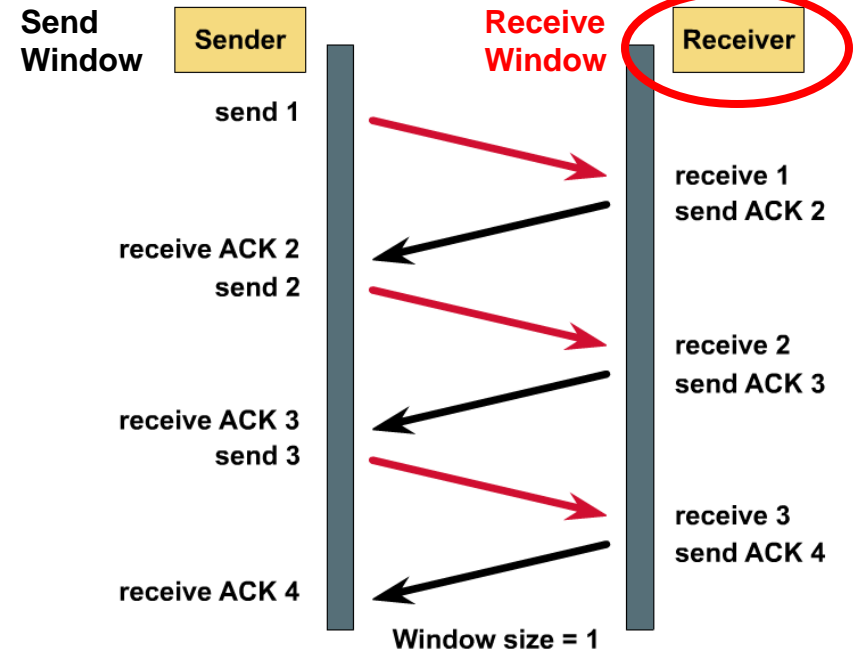
Send Window (not a TCP field)

- The TCP Receive Window size of the other host.
- How much data (in bytes) that can be sent by this host before receiving an acknowledgement from the other host.

TCP Window Size



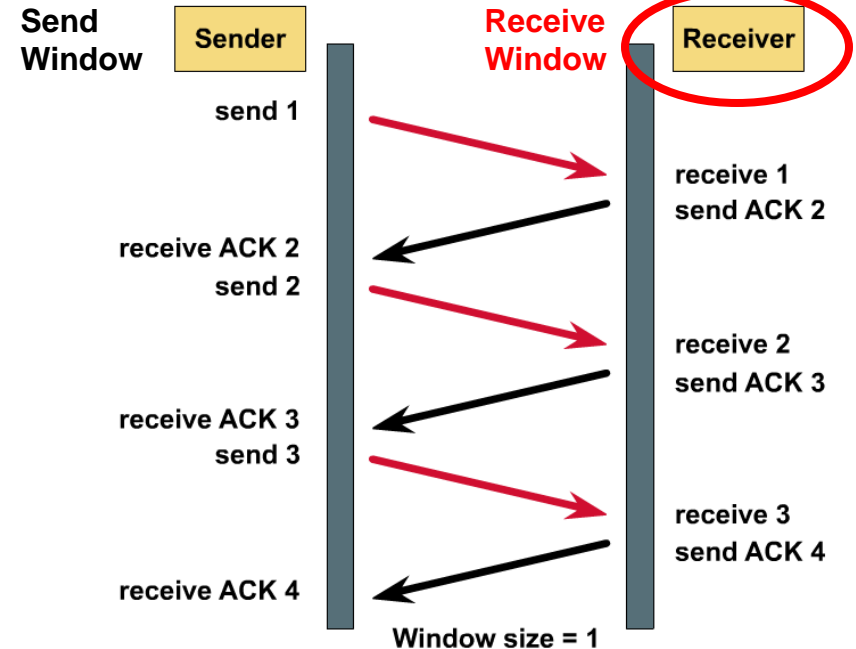
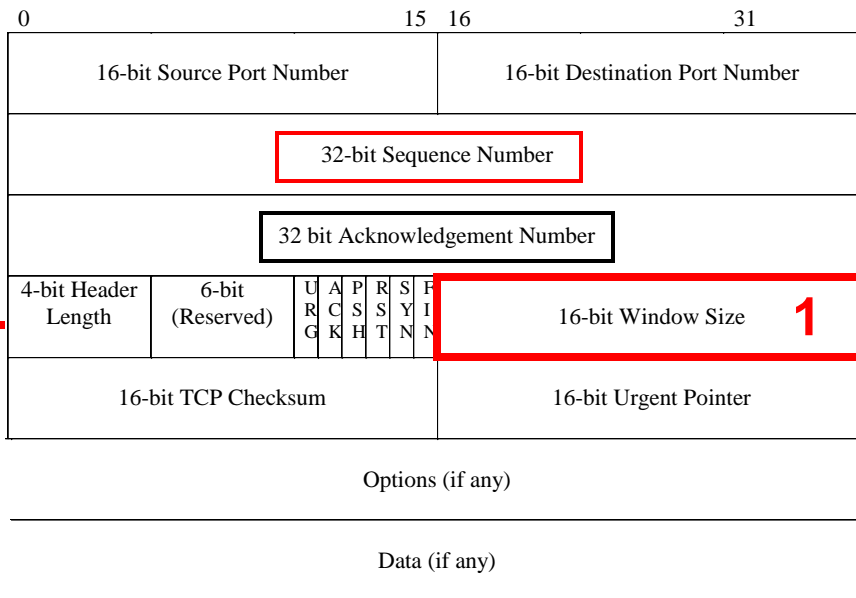
TCP Simple Acknowledgment



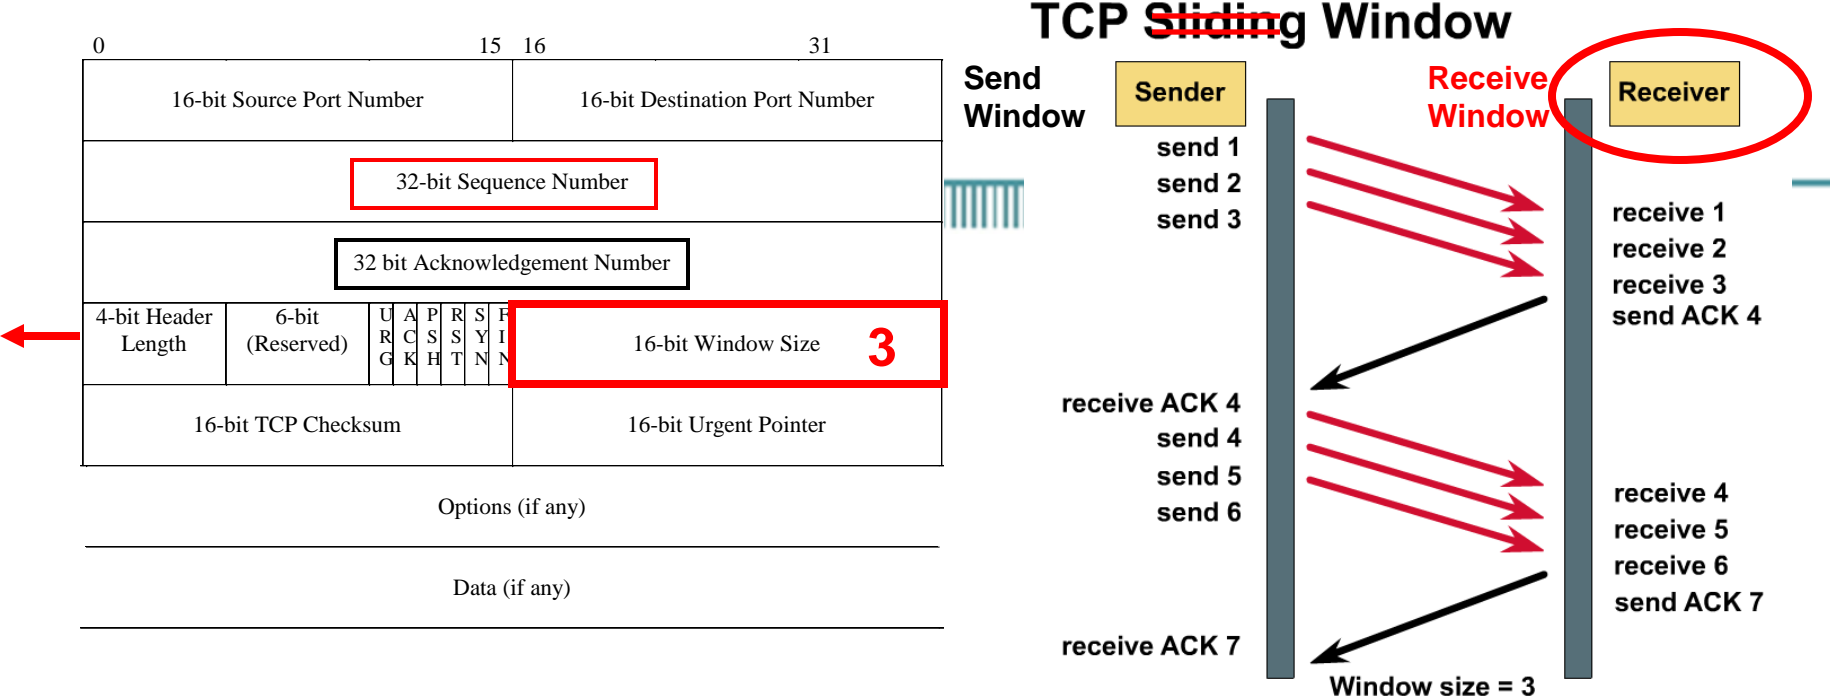
- After a host transmits the window-size number of bytes, it must receive an acknowledgment before any more data can be sent.
- The window size determines how much data the receiving station can accept at one time.

TCP Window Size

TCP Simple Acknowledgment

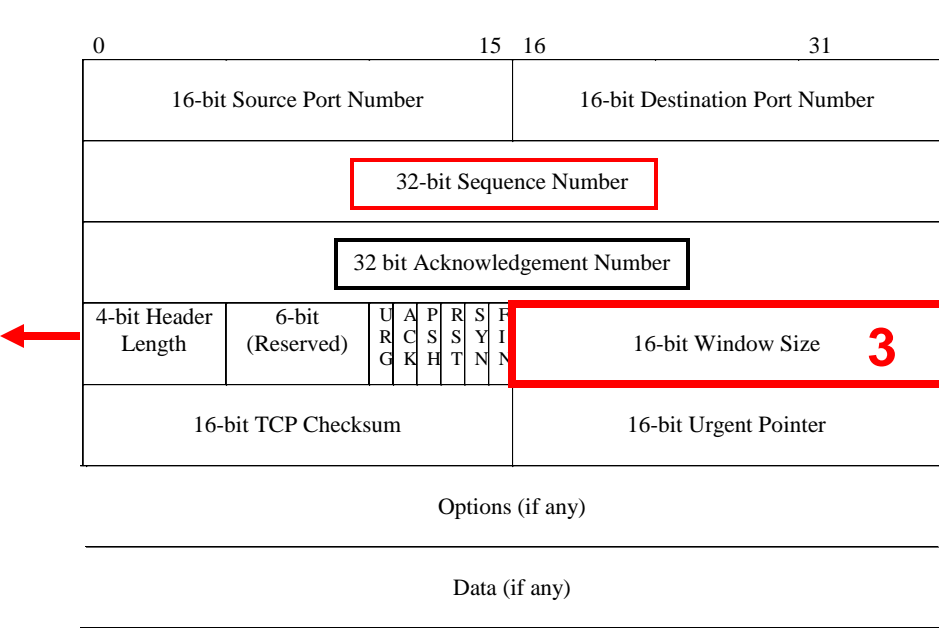


- With a **window size of 1**, each segment carries only one byte of data and must be acknowledged before another segment is transmitted.
- This results in **inefficient** host use of bandwidth.
- The purpose of windowing is to improve flow control and reliability.
- Unfortunately, with a window size of 1, you see a very inefficient use of bandwidth.

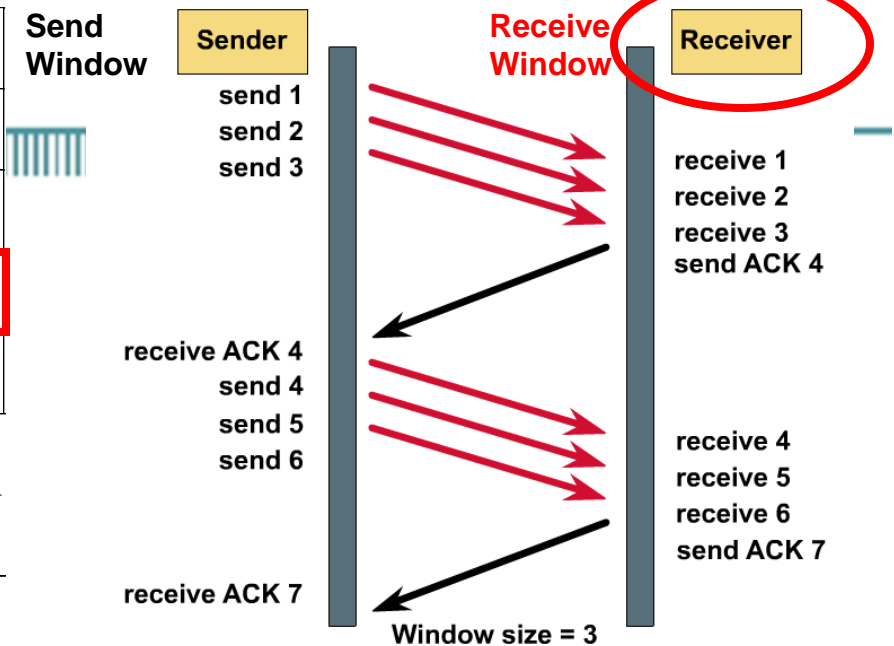


Receiver's TCP Window Size

- TCP uses a window size, number of bytes, that the receiver is willing to accept, and is usually controlled by the receiving process.
- TCP uses **expectational acknowledgments**, meaning that the acknowledgment number refers to the next byte that the sender of the acknowledgement expects to receive.
- A larger window size allows more data to be transmitted pending acknowledgment.
- Note: The sequence number being sent identifies the first byte of data in that segment.



TCP Sliding Window



TCP Window Size

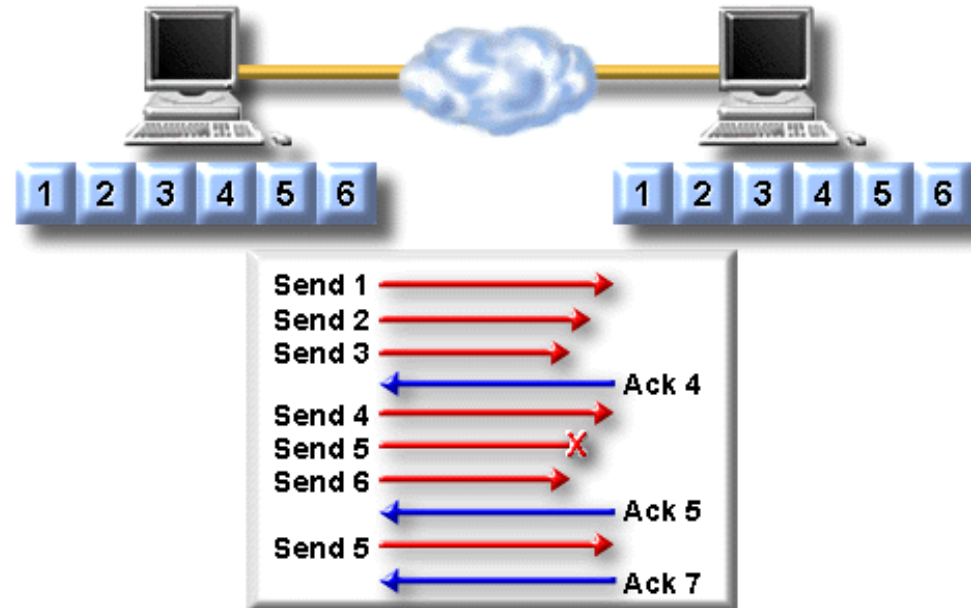
- TCP provides **full-duplex service**, which means data can be flowing in each direction, independent of the other direction.
- Window sizes, sequence numbers and acknowledgment numbers are independent of each other's data flow.
- Receiver sends acceptable window size to sender during each segment transmission (flow control)
 - if too much data being sent, acceptable window size is reduced
 - if more data can be handled, acceptable window size is increased
- This is known as a **Stop-and-Wait** windowing protocol.

0		15		16		31	
16-bit Source Port Number				16-bit Destination Port Number			
32-bit Sequence Number				TCP Header			
32 bit Acknowledgement Number							
4-bit Header Length	6-bit (Reserved)	URG	ACK	PUSH	RESET	SYN	FIN
		16-bit Window Size					
16-bit TCP Checksum				16-bit Urgent Pointer			
Options (if any)							
Data (if any)							

TCP Header

3

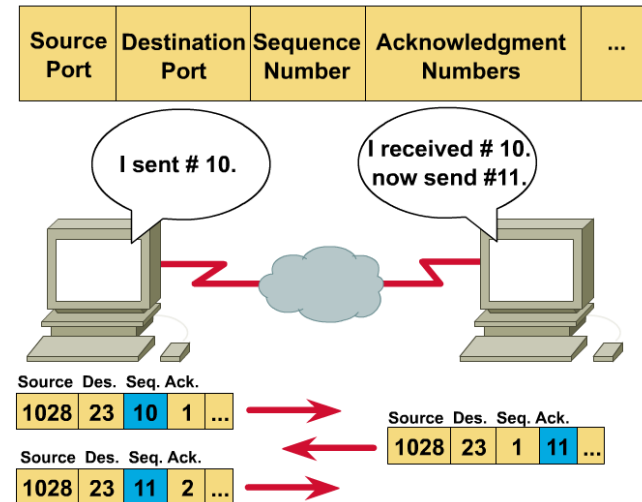
An Acknowledgment Technique



- Packets may be dropped along the way, timed out, or corrupted.
- If octets 4, 5, and 6 were sent, but 5 was lost, the receiver would only acknowledge up to 4, sending an Ack of 5.
- The sender would send 5 and wait to hear from the receiver where it should start up again.
- The receiver sends Ack 7, so the sender knows it can start sending again with octet 7.
- *There are **selective acknowledgements (SACK)** – not discussed here, which is a way of acknowledging selected pieces of the data stream.*

TCP Sequence and Acknowledgment Numbers

This is only if one octet was sent at a time, but what if multiple bytes are sent, which is usually the case?



Tech Note (FYI)

- Sender: The value in the sequence number is the first byte in the data stream.
- So, how does the receiver know how much data was sent, so it knows what value to send in the acknowledgement?
- Receiver: Using the sender's IP packet and TCP segment information, the value of the ACK is:

IP Length: (IP header) Total length - Header length

- TCP header length (TCP header): Header length

Length of data in TCP segment

ACK = Last Sequence Number acked + Length of data in TCP segment

- Check Sequence Number to check for missing segments and to sequence out-of-order segments.
- Remember that the ACK is for the sequence number of the byte you expect to receive. When you ACK 101, that says you've received all bytes through 100. This ignores SACK.

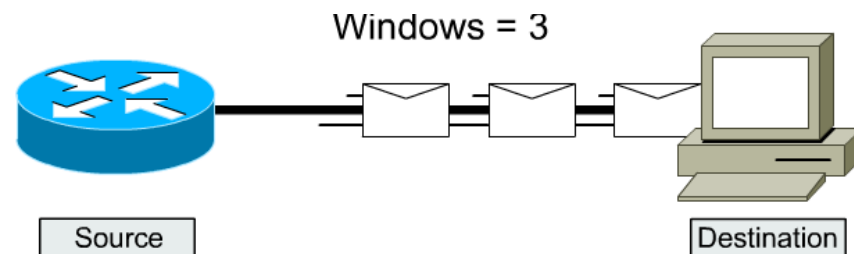
Sliding Windows

- Note: The following two slides on Sliding Windows contains corrections to the on-line curriculum followed by my slides on Sliding Windows.

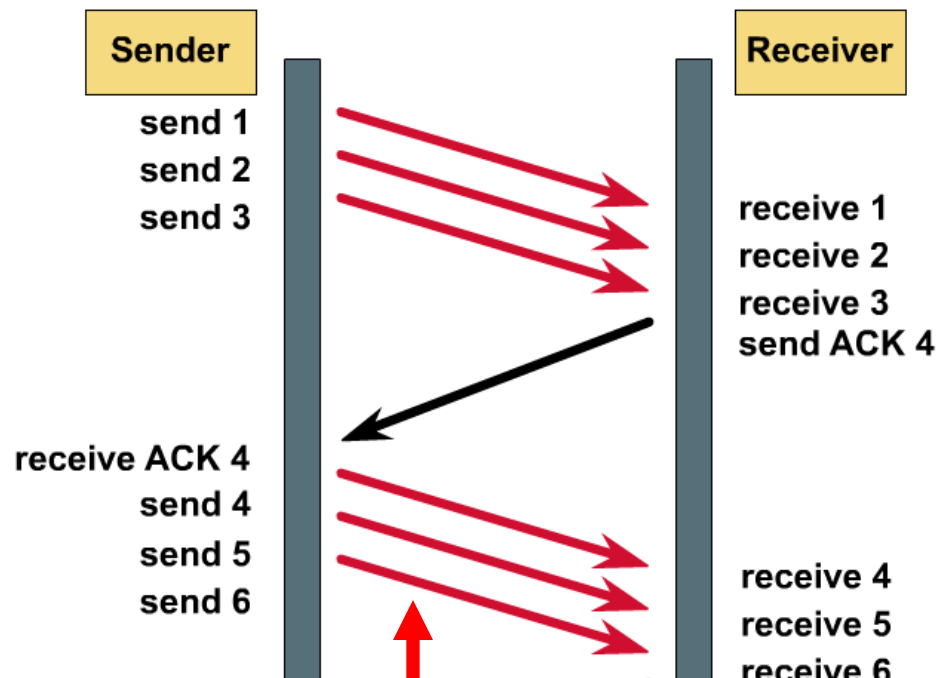
TCP Header

31

16-bit Source Port Number										16-bit Destination Port Number									
32-bit Sequence Number																			
32 bit Acknowledgement Number																			
4-bit Header Length				6-bit (Reserved)				U	A	P	R	S	F	16-bit Window Size					
								R	C	S	H	G	K						
16-bit TCP Checksum										16-bit Urgent Pointer									
Options (if any)																			
Data (if any)																			



TCP Sliding Window



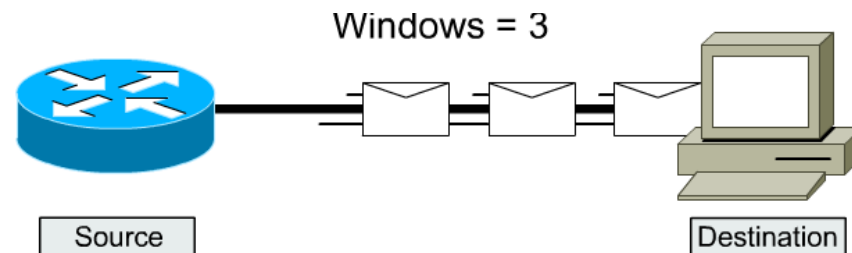
From Cisco Curriculum: This diagram is not an example of a sliding window, but of a window size of 3.

- TCP uses expectational acknowledgments, meaning that the acknowledgment number refers to the octet expected next.
- “The *sliding* part of *sliding window* refers to the fact that the window size is negotiated dynamically during the TCP session.” **(This is not exactly what a sliding window is! Coming soon!)**
- A sliding window results in more efficient host use of bandwidth because a larger window size allows more data to be transmitted pending acknowledgment.

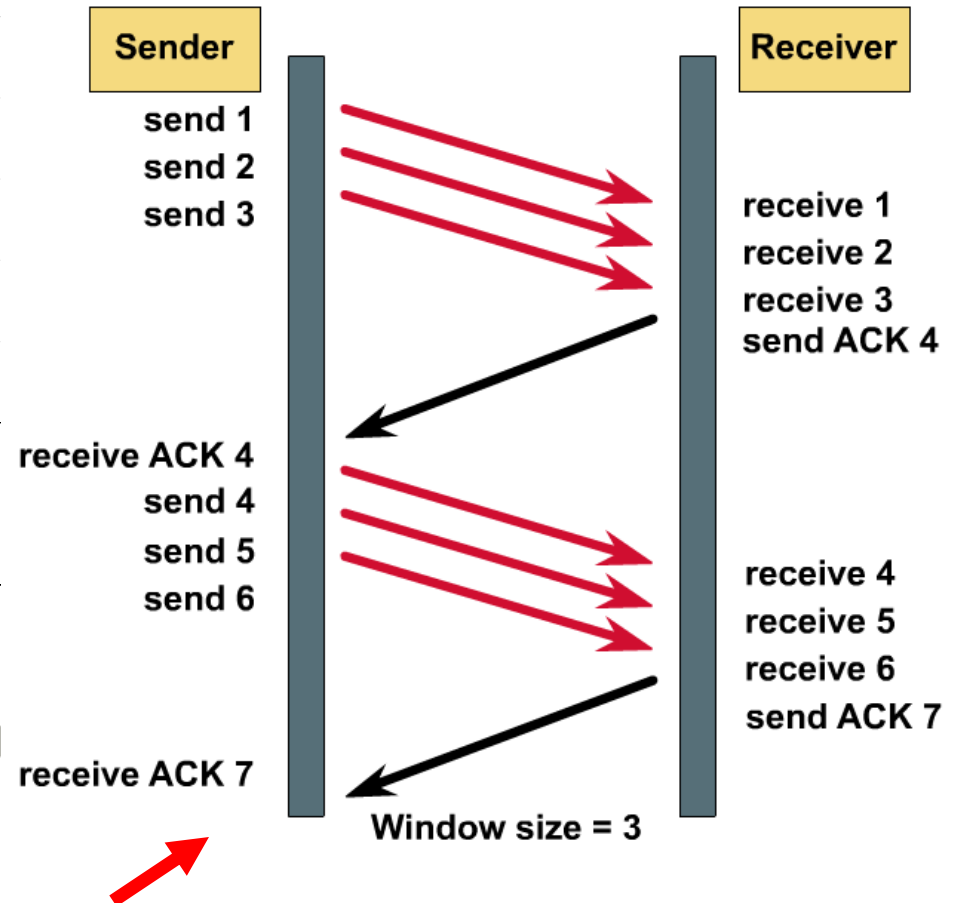
TCP Header

31

16-bit Source Port Number					16-bit Destination Port Number						
32-bit Sequence Number											
32 bit Acknowledgement Number											
4-bit Header Length		6-bit (Reserved)		U R G	A C K	P R S T	S S H	S Y N	F I N	16-bit Window Size	
16-bit TCP Checksum					16-bit Urgent Pointer						
Options (if any)											
Data (if any)											



~~TCP Sliding Window~~



This diagram is not an example of a sliding window, but of a window size of 3.

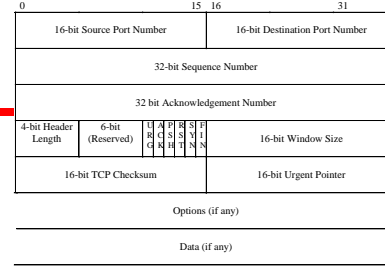
- From Cisco Curriculum: “A sliding window results in more efficient host use of bandwidth because a larger window size allows more data to be transmitted pending acknowledgment. “ **(A larger window size does this, not a sliding window.)**

Sliding Windows

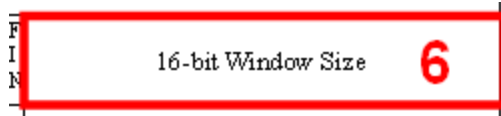
Initial Window size		Working Window size	
Usable Window		Octets sent	Usable Window
Can send ASAP		Not ACKed	Can send ASAP

Sliding Window Protocol

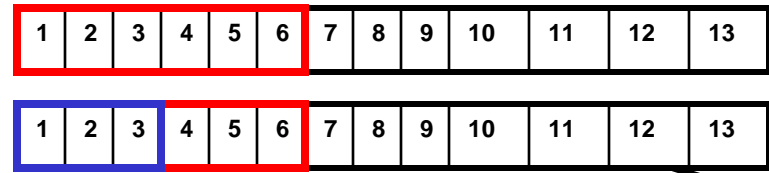
- Sliding window algorithms are a method of flow control for network data transfers using the receivers Window size.
- The sender computes its usable window, which is how much data it can immediately send.
- Over time, this sliding window moves to the rights, as the receiver acknowledges data.
- The receiver sends acknowledgements as its TCP receive buffer empties.
- The terms used to describe the movement of the left and right edges of this sliding window are:
 1. The left edge closes (moves to the right) when data is sent and acknowledged.
 2. The right edge opens (moves to the right) allowing more data to be sent. This happens when the receiver acknowledges a certain number of bytes received.
 3. The middle edge open (moves to the right) as data is sent, but not yet acknowledged.



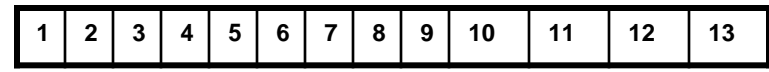
TCP Header



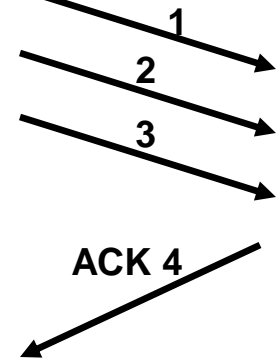
Host A - Sender



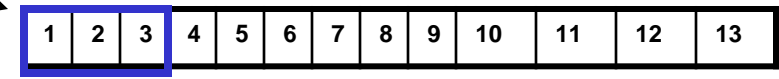
Host B - Receiver



Window size = 6

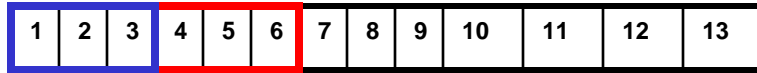


Octets received

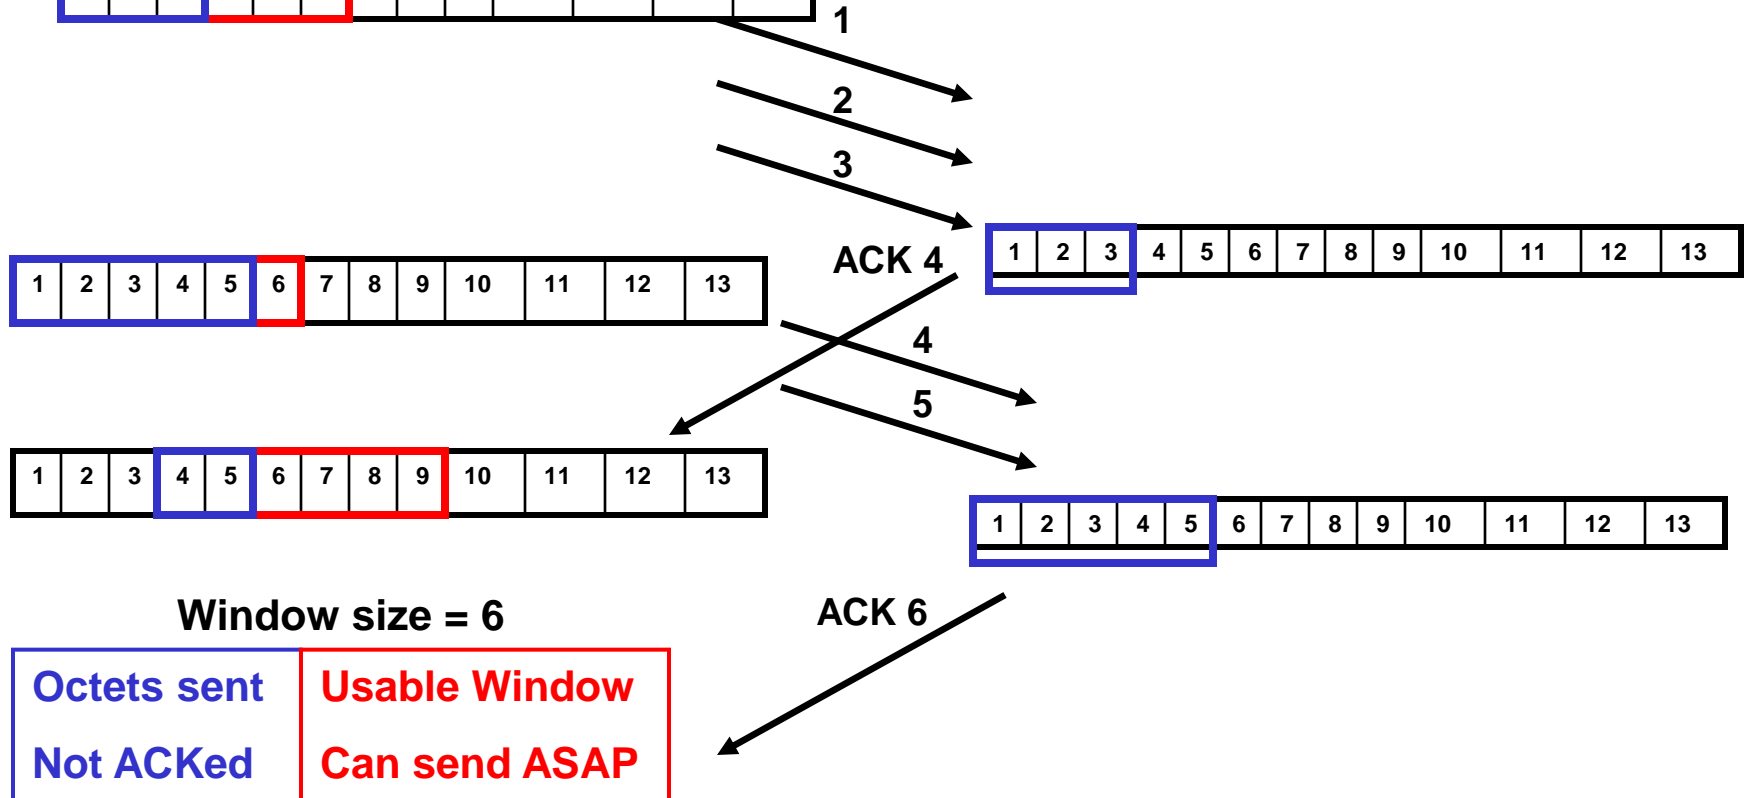


- Host B gives Host A a window size of 6 (octets).
- Host A begins by sending octets to Host B: octets 1, 2, and 3 and slides it's window over showing it has sent those 3 octets.
- Host A will not increase its usable window size by 3, until it receives an ACKnowledgement from Host B that it has received some or all of the octets.
- Host B, not waiting for all of the 6 octets to arrive, after receiving the third octet sends an expectational ACKnowledgement of “4” to Host A.

Host A - Sender

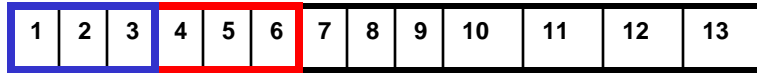


Host B - Receiver

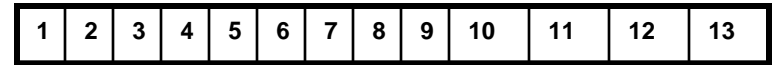


- Host A does not have to wait for an acknowledgement from Host B to keep sending data, not until the window size reaches the window size of 6, so it sends octets 4 and 5.
- Host A receives the acknowledgement of ACK 4 and can now **slide** its window over to equal 6 octets, 3 octets sent – not ACKed plus 3 octets which can be sent asap

Host A - Sender



Host B - Receiver



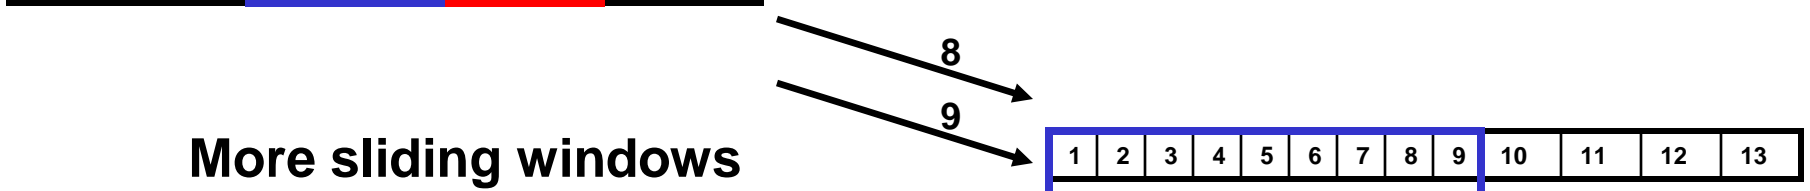
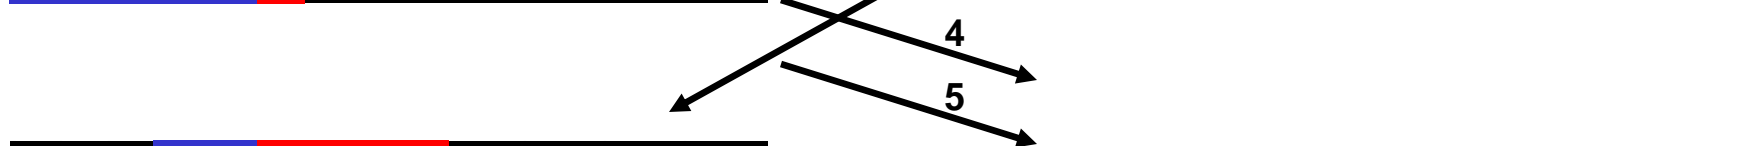
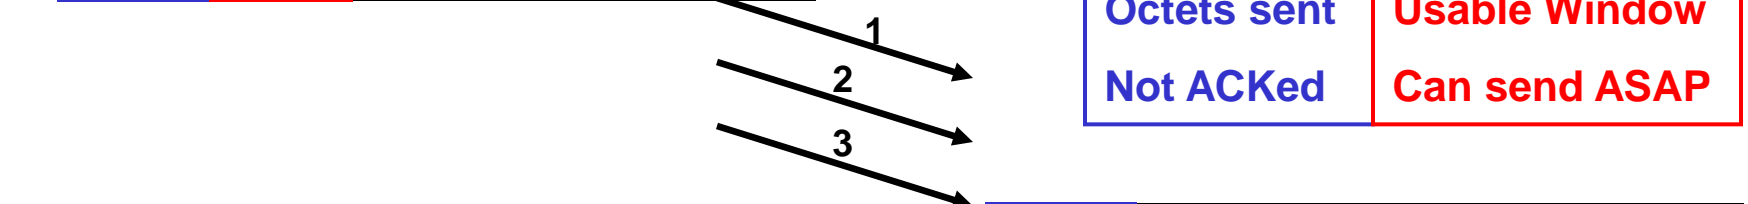
Window size = 6

Octets sent

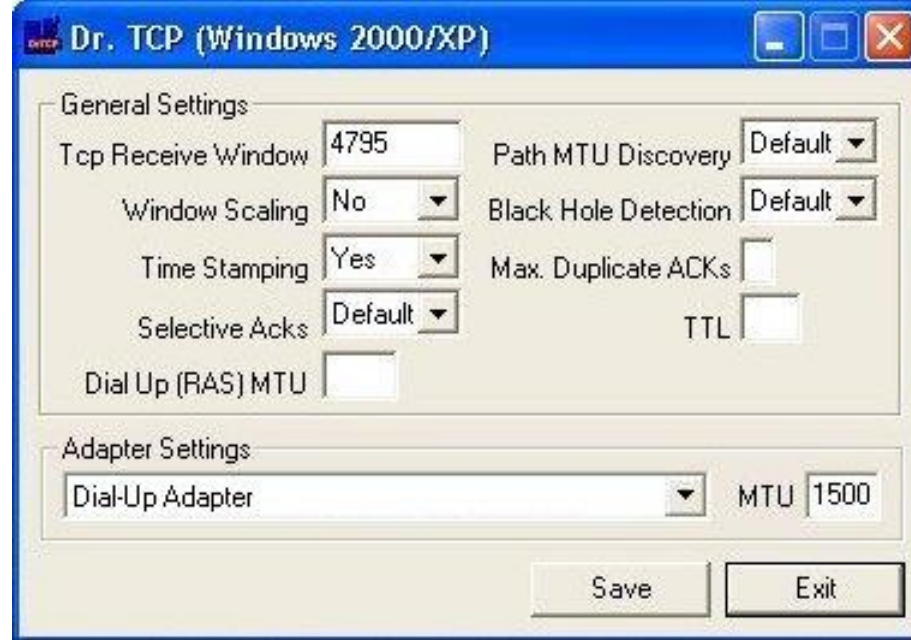
Usable Window

Not ACKed

Can send ASAP



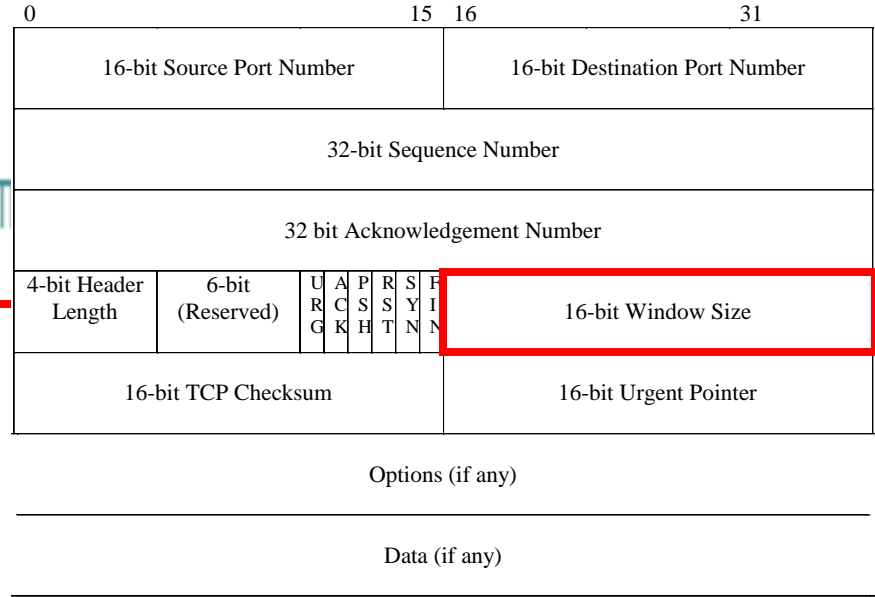
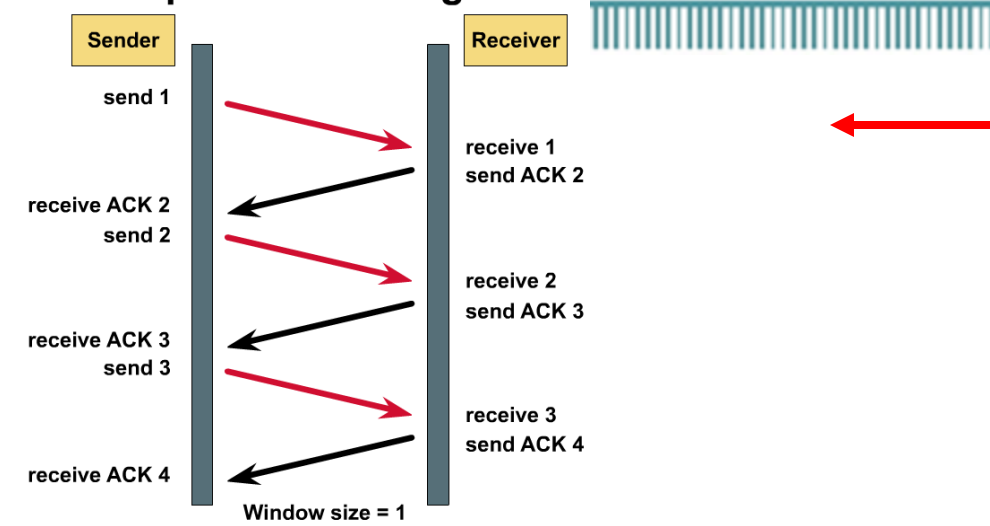
More sliding windows



- Default 8K for Windows, 32K for Linux,
- There are various unix/linux/microsoft programs that allow you to modify the default window size.
- I do not recommend that you mess around with this unless you know what you are doing.
- **“Disclaimer:** Modifying the registry can cause serious problems that may require you to reinstall your operating system. We cannot guarantee that problems resulting from modifications to the registry can be solved. Use the information provided at your own risk.”
- **NOTE:** I take no responsibility for this software or any others!

Receive and Send Windows

TCP Simple Acknowledgment



Receive Window

- The TCP Receive Window size is the amount of receive data (in bytes) that can be buffered by this host, at one time on a connection.
- The other (sending) host can send only that amount of data before getting an acknowledgment and window update from this (the receiving) host.

Send Window

- The TCP Receive Window size of the other host.
- How much data (in bytes) that can be sent by this host before receiving an acknowledgement from the other host.

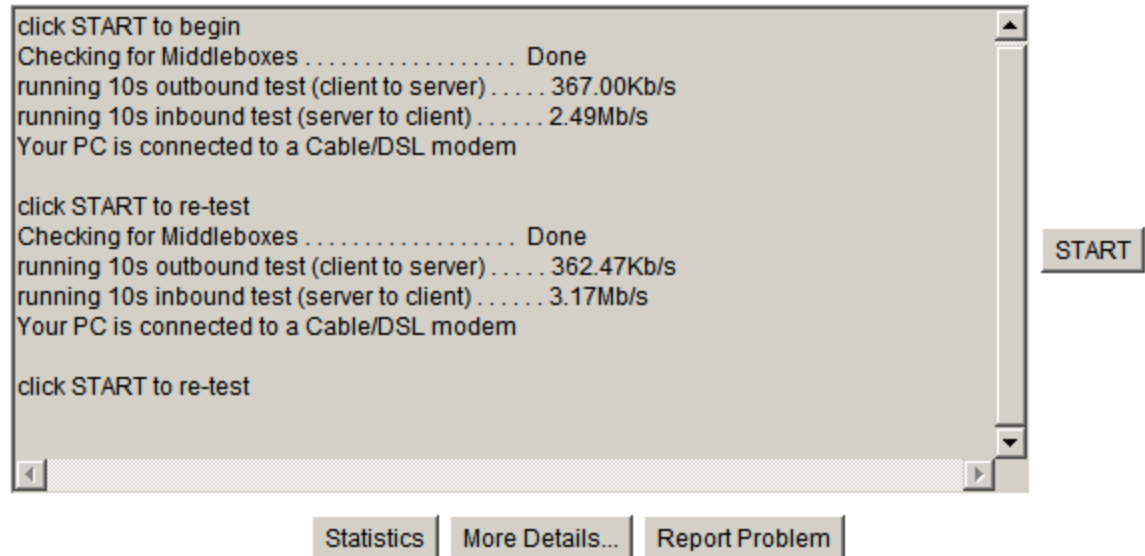
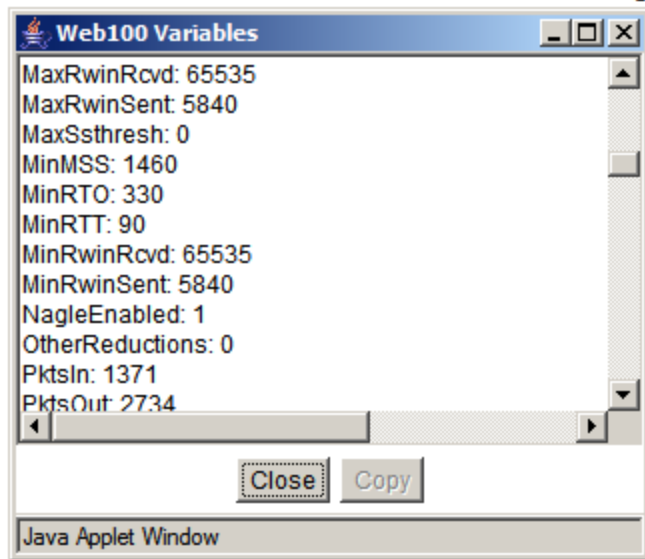
Sliding Windows – From TCPGuide.com

- Example: Server's window size was 360.
- This means the server is willing to take no more than 360 bytes at a time from the client.
- When the server receives data from the client, it places it into this buffer.
- The server must then do two distinct things with this data:
 1. **Acknowledgment:** The server must send an acknowledgment back to the client to indicate that the data was received.
 2. **Transfer:** The server must process the data, transferring it to the destination application process.
- The key point is that in the basic sliding windows system, data is acknowledged when received, but ***not necessarily*** immediately transferred out of the buffer.
- This means that is possible for the buffer to fill up with received data faster than the receiving TCP can empty it.
- When this occurs, the receiving device may need to adjust window size to prevent the buffer from being overloaded.

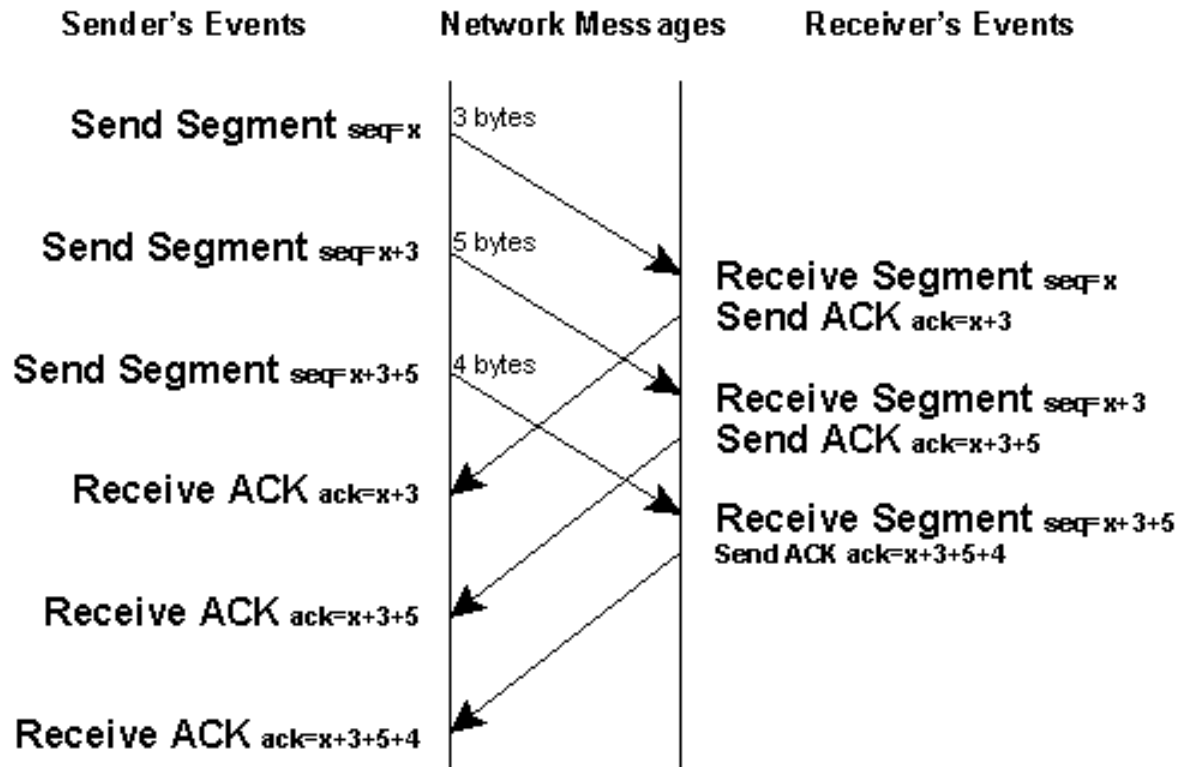
Bandwidth Testing and Other Statistics

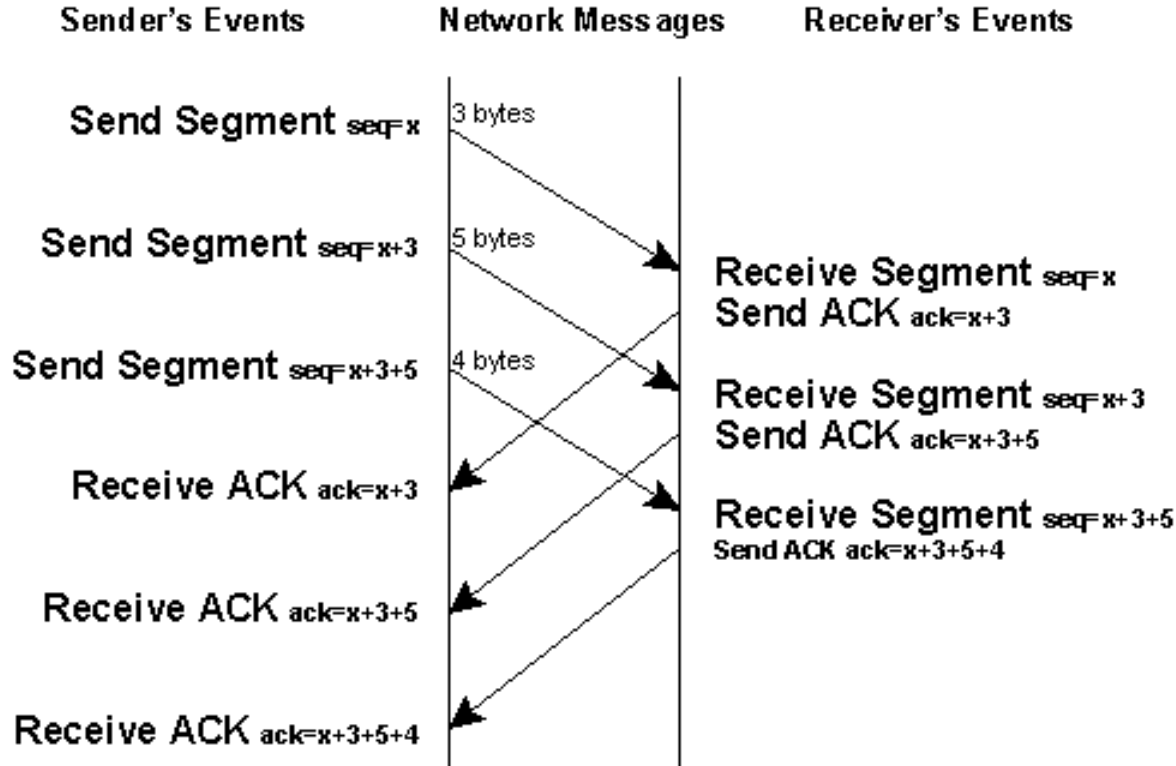
Using a browser go to this link and click on **start**:

- <http://miranda.ctd.anl.gov:7123/>



More on TCP Sequence Numbers and Acknowledgements

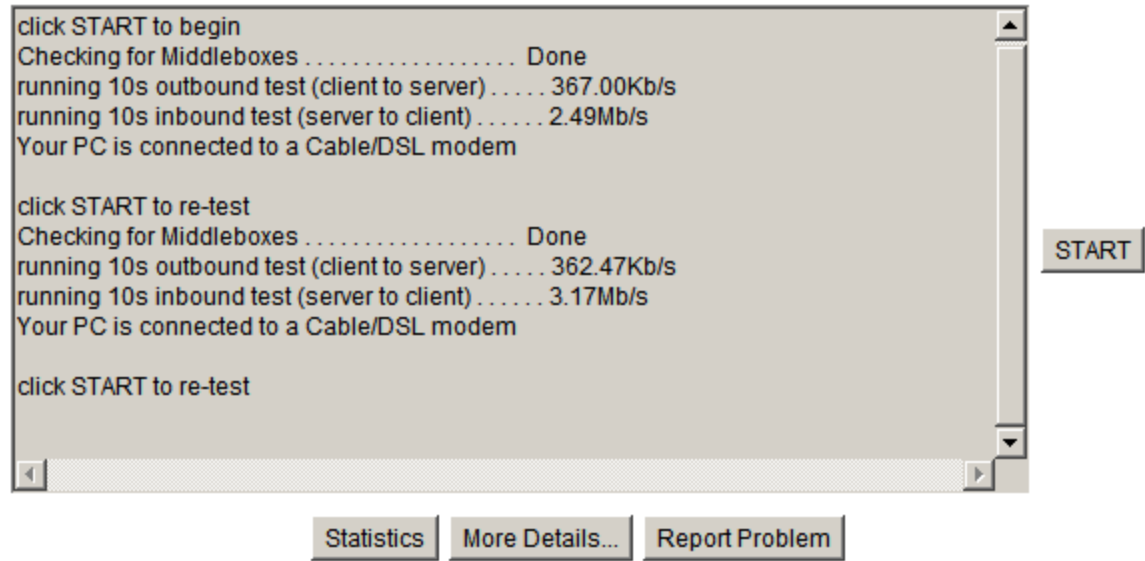
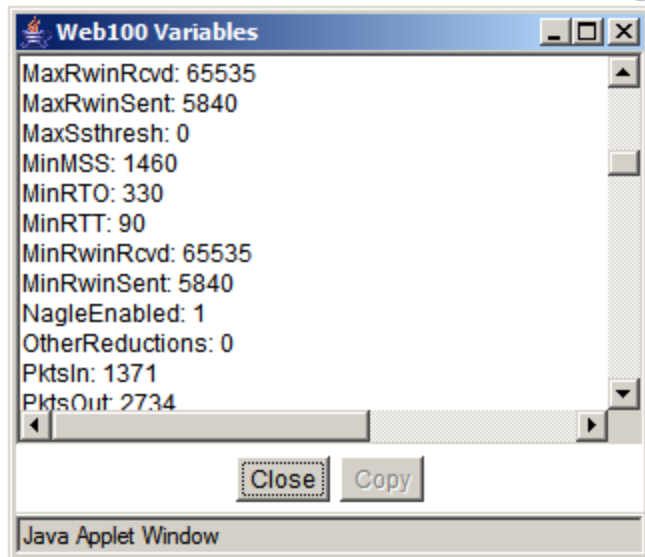




- The sequence and acknowledgment numbers are directional, which means that the communication occurs in both directions.
- The figure illustrates the communication going in one direction.
- The sequence and acknowledgments take place with the sender on the right.
- TCP provides **full-duplex service**, which means data can be flowing in each direction, independent of the other direction.
- Window sizes, sequence numbers and acknowledgment numbers are independent of each other's data flow.

We will do this in the lab:

- <http://miranda.ctd.anl.gov:7123/>



Viewing Receive Window Sizes in Ethereal

No. -	Time	Source	Destination	Protocol	Info
93	6.223575	146.137.222.101	192.168.1.100	TCP	3002 > 3158 [ACK] Seq=1 Ack=49153 win=21736 Len=0
94	6.223718	192.168.1.100	146.137.222.101	TCP	3158 > 3002 [ACK] Seq=57345 Ack=1 win=65535 Len=1460
95	6.223757	192.168.1.100	146.137.222.101	TCP	3158 > 3002 [ACK] Seq=58805 Ack=1 win=65535 Len=1460
96	6.223789	192.168.1.100	146.137.222.101	TCP	3158 > 3002 [ACK] Seq=60265 Ack=1 win=65535 Len=1460
97	6.223823	192.168.1.100	146.137.222.101	TCP	3158 > 3002 [ACK] Seq=61725 Ack=1 win=65535 Len=1460
98	6.223854	192.168.1.100	146.137.222.101	TCP	3158 > 3002 [ACK] Seq=63185 Ack=1 win=65535 Len=1460
99	6.223906	192.168.1.100	146.137.222.101	TCP	3158 > 3002 [PSH, ACK] Seq=64645 Ack=1 win=65535 Len=892
100	6.274716	146.137.222.101	192.168.1.100	TCP	3002 > 3158 [ACK] Seq=1 Ack=52073 win=21736 Len=0
Frame 94 (1514 bytes on wire, 1514 bytes captured)					
Ethernet II, Src: 192.168.1.100 (00:0a:e4:d4:4c:f3), Dst: 192.168.1.1 (00:0f:66:09:4e:0f)					
Internet Protocol, Src: 192.168.1.100 (192.168.1.100), Dst: 146.137.222.101 (146.137.222.101)					
Transmission Control Protocol, Src Port: 3158 (3158), Dst Port: 3002 (3002), Seq: 57345, Ack: 1, Len: 1460					
Source port: 3158 (3158)					
Destination port: 3002 (3002)					
Sequence number: 57345 (relative sequence number)					
[Next sequence number: 58805 (relative sequence number)]					
Acknowledgement number: 1 (relative ack number)					
Header length: 20 bytes					
Flags: 0x0010 (ACK)					
window size: 65535					
checksum: 0x2a22 [correct]					
Data (1460 bytes)					

Receive Window

- The TCP Receive Window size is the amount of receive data (in bytes) that can be buffered by this host, at one time on a connection.
- The other (sending) host can send only that amount of data before getting an acknowledgment and window update from this (the receiving) host.

Viewing the Send Window Sizes in Ethereal

No. -	Time	Source	Destination	Protocol	Info
3343	22.871400	192.168.1.100	146.137.222.101	TCP	3159 > 3003 [ACK] Seq=1 Ack=2668881 win=65535 Len=0
3344	22.875823	146.137.222.101	192.168.1.100	TCP	3003 > 3159 [ACK] Seq=2668881 Ack=1 win=5840 Len=1460
3345	22.876179	146.137.222.101	192.168.1.100	TCP	3003 > 3159 [ACK] Seq=2670341 Ack=1 win=5840 Len=1460
3346	22.876215	192.168.1.100	146.137.222.101	TCP	3159 > 3003 [ACK] Seq=1 Ack=2671801 win=65535 Len=0
3347	22.884436	146.137.222.101	192.168.1.100	TCP	3003 > 3159 [PSH, ACK] Seq=2671801 Ack=1 win=5840 Len=1460
3348	22.884843	146.137.222.101	192.168.1.100	TCP	3003 > 3159 [ACK] Seq=2673261 Ack=1 win=5840 Len=1460
3349	22.884894	192.168.1.100	146.137.222.101	TCP	3159 > 3003 [ACK] Seq=1 Ack=2674721 win=65535 Len=0
3350	22.885185	146.137.222.101	192.168.1.100	TCP	3003 > 3159 [ACK] Seq=2674721 Ack=1 win=5840 Len=1460
Frame 3344 (1514 bytes on wire, 1514 bytes captured)					
Ethernet II, Src: 192.168.1.1 (00:0f:66:09:4e:0f), Dst: 192.168.1.100 (00:0a:e4:d4:4c:f3)					
Internet Protocol, Src: 146.137.222.101 (146.137.222.101), Dst: 192.168.1.100 (192.168.1.100)					
Transmission Control Protocol, Src Port: 3003 (3003), Dst Port: 3159 (3159), Seq: 2668881, Ack: 1, Len: 1460					
Source port: 3003 (3003)					
Destination port: 3159 (3159)					
Sequence number: 2668881 (relative sequence number)					
[Next sequence number: 2670341 (relative sequence number)]					
Acknowledgement number: 1 (relative ack number)					
Header length: 20 bytes					
Flags: 0x0010 (ACK)					
window size: 5840					
checksum: 0xaec2 [correct]					
Data (1460 bytes)					

Send Window

- The TCP Receive Window size of the other host.
- How much data (in bytes) that can be sent by this host before receiving an acknowledgement from the other host.

Ch.11 – TCP/IP Transport and Application Layers



Cabrillo College

CIS 81 and CST 311

Rick Graziani

Spring 2006